
Statistic Preservation Does Not Identify Attack-Suppression Preservation: Limits of Composed Defenses in Collaborative Learning

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 **A preserved statistic is not preserved suppression, and no test that admits a**
2 **composition because one constituent already suppresses the attack can iden-**
3 **tify whether the downstream mechanism is responsible for the composition’s**
4 **success.** Collaborative learning composes defenses that were never designed to-
5 gether, so upstream processing can destroy what the downstream defense reads,
6 and the natural repair is to require the upstream stage to **preserve** it; for the broad
7 class of *positive per-client rescaling* compositions, where the invariance question
8 is exactly algebraic, that is not enough. The obstruction is structural: a criterion
9 gated on the outcome it predicts conditions on a descendant of both the mecha-
10 nism and the defense’s standalone strength, so **no screen of this form identifies**
11 **the mechanism’s contribution, at any sample size.** Ignoring this does not cost
12 power, it costs the sign: on the identical cell the outcome-gated ladder moves
13 -0.272 where the instrument moves $+0.098$. Preservation is a hierarchy, and we
14 break its last two links by intervening rather than comparing: Oracle intervention
15 pinning every adversary’s coefficient closes the attenuation channel, and disturb-
16 ing the statistic Krum reads then flips its decision in up to 80% of rounds while the
17 *support* of the admitted adversarial mass never changes in any measured round of
18 any arm, and suppression stays inside our pre-specified ± 0.15 equivalence margin.
19 We prove sufficient conditions for the weighting $\rightarrow$ aggregator class and replicate
20 the negative in a second invariance class and on a second dataset and architecture.
21 The implication for decentralized training is methodological: evaluate composed
22 mechanisms by intervening on adversarial contribution, not by inferring it from
23 downstream statistics.

24 1 Introduction

25 Federated learning [McMahan et al., 2017] is vulnerable to data poisoning [Chen et al., 2017, Bag-
26 dasaryan et al., 2020], and an arms race has produced defenses [Fung et al., 2020, Pillutla et al.,
27 2022, Yin et al., 2018, Blanchard et al., 2017] with accuracy/robustness trade-offs and none univer-
28 sally dominant (Appendix F). *Can we compose them for stronger robustness?* Composition does
29 not even commute: FoolsGold $\rightarrow$ CoordMedian suppresses a pixel backdoor to ASR 0.099 while
30 the reverse order reaches 0.732: a rank aggregator emits one vector rather than per-client updates,
31 so it cannot be an *upstream* stage and the pair collapses to FoolsGold alone, bit-identically. Naive
32 layering fails this way whenever upstream processing destroys what the downstream defense reads.

33 **Why the question is sharper without a trusted coordinator.** Such stacks compose independently
34 developed mechanisms among heterogeneous, partially trusted, often pseudonymous participants,




Figure 1: **The paper in one figure: the statistic moves, admission does not, suppression does not follow.** (a) Two paths run from an upstream rescaling T to suppression: the *chain* a preservation criterion reasons along, and a *confounded* path on which T merely attenuates whichever client carries the poison (Lem. 1) and reaches ASR **without passing through d_2 's statistic at all**. **Mode S cuts the second by construction**, pinning every adversary at $c=1$ (the one channel the dose closes, and not the only thing it moves). The admission count pools the four aggregators of §5; Krum admits no adversary at any rung, so its zero is a floor (App. C). (b) Mode S, Krum under model-scaling (§4): as ρ grows to 54.6 the *decision* flips in up to 80% of rounds while the *support* of the admitted adversarial mass does not, and max-committed ASR stays inside the pre-specified ± 0.15 equivalence margin ($|\Delta|=0.026$). Error bars in this panel are Student- t 95% intervals on the levels ($n=5$). (c) **Six cells, one contrast each, estimated two ways** (circle: outcome-gated ladder; diamond: Mode S, adversary pinned at $c=1$). **The designs disagree on 4 of 6 cells, and on coord_median/pixel they have opposite signs:** Δ ASR = -0.272 (paired 95% t interval $[-0.422, -0.122]$) against $+0.098$ ($[+0.018, +0.178]$), intervals excluding zero and not overlapping, so **closing the attenuation channel does not shrink the effect, it reverses its sign**. Above the dashed rule are the four cells a pre-registered mechanism was read off (*training data, not evidence for it*); below are its two out-of-sample cells, one of which **refutes it, so we withdraw it**.

so preservation is the obvious discipline to demand of them. Concretely: a contributor clips its own update, a scheduler reweights contributors by reputation, and whoever merges runs a robust aggregator: three parties, three mechanisms, nobody owning the stack. The only check at each seam is that the next stage still reads what it expects, so a preserved statistic is read as evidence the stack is robust. **That inference is what this paper refutes. The identification problem we isolate is defined independently of model scale**, so we study it in a controlled federated setting and claim no transfer beyond it.

Contributions. We propose no new defense, and **our central contribution is a negative identification result**: an outcome-gated comparison of defense pairs cannot determine whether preserving a downstream defense's mechanism is what keeps the composition secure. *Attack-suppression preservation* names the object throughout: suppression of the committed attack surviving composition, instantiated as ASR, never the statistic or the decision that follows; Fig. 1 is the paper in one picture. **The boundary.** Outcome-gated cross-defense screening is structurally unidentified: its eligibility gate is defined on the same outcome it predicts, so no cross-defense contrast *satisfying this outcome-gated design* identifies the mechanism's contribution and the admissible designs are within-defense (§3). **What is new is not the collider.** The identification failure is a positivity violation and that half is textbook; what is new is the quantity it applies to, a preservation hierarchy for federated ag-

gregation whose levels come apart at security-relevant channels, together with the algebra and the interventions that locate a pair inside it. **The demonstration.** Statistic preservation is accordingly not a valid causal proxy for attack-suppression preservation, established by intervention rather than correlation: on the identical cell the outcome-gated design estimates -0.272 where the intervention estimates $+0.098$, reversing the direction and not merely the precision (§4). **The consequence.** The screening rule follows from the boundary and is nothing more: evaluation prioritization, measurably broken under heterogeneity (App. J.2). **Also.** The preservation hierarchy (Def. 1); mechanism preservation as invariance under positive per-client rescalings (Prop. 1), with margins on the weighting→aggregator class (Thm. 1), supporting machinery, safe at *both* extremes of the weight ratio, so dispersion alone mispredicts; counterexamples across four aggregators of three structural kinds; and the negative replicated in a second invariance class and on a second dataset and architecture, while refuting the cross-arm admission ordering we had predicted. Scope is inventoried once, in Tables 1 and 2: what each claim does *not* license, and the evidence tiers (App. J.1).

2 Which Level a Composition Preserves, and How to Decide It Before Running

What does it mean for a composition to preserve a defense? Every per-client defense we consider acts as a positive scalar rescaling $T(u_i) = c_i u_i$, which gives the invariance question a sharp *algebraic* answer (Prop. 1), but writing C2 as an invariance of the statistic S leaves a gap our own experiments exposed. Three different things can be preserved: what the downstream defense *reads* (P1–P2), what it *decides* and what *survives* it (P3–P4), and whether the attack is still *suppressed* (P5). They come apart, so we name the levels and ask of each whether it determines the one below:

Level	What is preserved	Determines the level below?
(P1)	the statistic’s value	yes, by definition
(P2)	the ordering it induces	yes, by definition
(P3)	the decision d_2 makes	no , by counterexample (§4)
(P4)	the adversarial contribution admitted	no , by counterexample (§5)
(P5)	suppression of the committed attack	—

Both breaks are measured, at fixed d_2 and attack: Krum’s decision flips in up to 80% of rounds with the adversarial mass it admits unchanged (§4), and `cos_krum`’s decision and admitted mass are bit-identical while ASR falls 0.173 (§5); Fig. 2 draws the chain and both breaks.

The algebra fixes which level a given composition sits at: cosine similarity is exactly invariant, so C2 holds for every $\cdot \rightarrow$ FoolsGold composition; coordinate ordering survives only under a bounded-ratio condition; consensus distance and residual magnitude are not invariant at all, so $\cdot \rightarrow$ Reputation compositions are C2 failures; norm clipping does not blind FoolsGold, though it does blind Reputation.

Definition 1 (Preservation levels). *For an upstream transform T and a downstream defense d_2 reading statistic S :*

- (P1) **Value invariance:** $S(T(U)) = S(U)$.
- (P2) **Ordering invariance:** $S(T(U)) = \varphi(S(U))$ for some φ preserving the ordering S induces on clients. (this is C2).
- (P3) **Decision invariance:** the client or coordinate d_2 actually selects is unchanged.
- (P4) **Adversarial contribution (admission) invariance:** the adversarial input d_2 admits into its aggregate is unchanged. Admission is the pre-registered term throughout for how much adversarial input survives d_2 : a level, not one measurable quantity, instantiated by whatever each rule exposes: Krum’s and `cos_krum`’s selected-client indicator; `coord_median`’s selected-coordinate share, reputation’s adversarial weight share, which is why App. C’s influence column is aggregator-specific.
- (P5) **Attack-suppression preservation:** ASR stays below threshold.

We target (P4) $\nrightarrow$ (P3) in §4 and (P4) $\nrightarrow$ (P5) in §5. The conditions (C0)–(C3) the levels induce (**the rule being $C1 \wedge C2$**) are in App. A and B.

What the theory delivers: supporting machinery, not the claim. For a per-client weighting into a rank-based or geometric aggregator, Thm. 1 gives sufficient conditions for d_2 ’s ordering to

99 survive: it **certifies mechanism preservation, not security** and bounds no ASR. Its one prediction,
100 admitted adversarial mass $\Lambda_a \leq 12.6\%$ against a realized $\leq 1.8\%$, goes vacuous as the separation
101 margin approaches 1, which heterogeneity reaches, and does not follow from a (d_1, d_2) definition
102 pair alone, as a screen would need (App. I).

103 3 The Testability Boundary

104 *Why can we not simply compare many defense pairs?* The natural test keeps pairs whose constituents
105 already suppress the attack and asks whether the composition preserves it. That is exactly where
106 identification is lost. **The boundary in words** (Prop. 4, formal in App. J.4): let a screen certify
107 $(d_1 \rightarrow d_2)$ only if an *eligibility gate* holds: some constituent already suppresses the attack, measured
108 by the same outcome metric and threshold the screen predicts for the composition, and a mechanism
109 condition P holds, fixed by the two defenses' definitions alone. Four cells follow, and that design
110 never fills two of them:

	Mechanism condition	d_2 suppresses a alone	it does not
111	P holds (preserved)	observed: every certified pair	empty by the gate
	$\neg P$ (destroyed)	observed: predicted-FAIL pairs	empty by the gate

112 Reading down a column is the contrast that would identify P 's contribution, and **the only column**
113 **in which d_2 is not already doing the work is empty**, for every menu a screen of this form could
114 range over: excluded by design, not undersampled: more data cannot fill a cell the design never
115 permits. The gate cannot be dropped either, because without it C1 is vacuous. So wherever our
116 conditions certify a pair, a low composition ASR is attributable to d_2 's standalone strength rather
117 than to preservation. Equivalently, this is selection on a collider [Elwert and Winship, 2014, Pearl,
118 2009]: the gate is a descendant of both P and d_2 's standalone effectiveness, so conditioning on it
119 makes them co-vary. The hazard is standard; the gate reading as hygiene rather than as the step that
120 removes identification is not, and we made that mistake ourselves (§5).

121 **The admissible design, derived: the practical contribution.** Read forwards the boundary is
122 *constructive*, each clause forced by a term in Prop. 4: (1) hold d_2 and the attack fixed, so the gate's
123 input is the same quantity on both sides of every contrast; (2) intervene on the upstream transform
124 rather than selecting a different d_2 , whose variation is what empties the cell; (3) pin the adversarial
125 coefficient channel, since an upstream that attenuates the adversary moves the outcome by a route
126 unrelated to d_2 's statistic, the confound that inverted our own sign; (4) measure every channel:
127 aggregate displacement, decision, admission, influence. §4 is clause (3) as an instrument; the six-
128 step form is in App. J.3. A second boundary bounds what such a screen can *find*: **no emergent**
129 **composition can satisfy an inherited-suppression gate**, our own lowest-ASR composition being
130 the witness: $FG \rightarrow RFA$ reaches 0.093 where FoolsGold alone reaches 0.732, so at $|\mathcal{E}|/|\mathcal{S}|=1/5$ the
131 bound caps recall at 80%, above the 40% we measure (App. J.2).

132 4 Statistic Preservation Does Not Identify Attack-Suppression Preservation

133 *Does disturbing what a defense reads actually cost suppression?* The claim: it need not, and **on the**
134 **same cell the design that says otherwise disagrees in sign**. Our pre-registered pooled hypothesis
135 (H1), that ASR rises with disturbance of the statistic d_2 reads, is rejected, the effect *negative*: that
136 ladder moved adversarial influence and the statistic together (§5). We adjudicate C2's implication
137 that losing (P3) means losing (P5) by oracle intervention on two arms in two invariance classes, not
138 by a general law.

139 **The instrument, as a counterfactual.** The clean question asks what happens when the down-
140 stream statistic changes but the adversary cannot become weaker. **Mode S** enforces exactly that: it
141 pins every adversary's coefficient at $c=1.0$ and spreads only the benign clients over $\rho = e^{2\kappa}$, hold-
142 ing the adversarial coefficient share constant at 0.266667 to within float32 read-back noise, closing
143 Figure 1(a)'s attenuation path, so a rise in ASR here cannot be attenuation. Mode S and Mode A (§5)
144 are causal instruments built from adversary identity: they measure mechanism-level effects rather
145 than emulate deployable defenses (Table 1, App. J.5).

146 **Two readings of C2, frozen against each other.** *The statistic reading:* what matters is disturbance
 147 of d_2 's statistic, read off its *decisions*. *The admission reading:* what matters is the *adversarial input*
 148 d_2 *admits*. Mode S disturbs Krum's selection heavily while leaving the mass it admits untouched,
 149 so that arm adjudicates (Table 4, the pre-registration's H-statistic and H-admission): as ρ climbs
 150 to 54.6 we find the decision changing in up to 0.80 of rounds while the adversarial mass admitted
 151 changes in none of them, and ASR showing no evidence of the predicted increase, remaining in-
 152 side the pre-specified equivalence margin ($p_{\uparrow}=0.70$, $|\Delta|=0.026 < 0.15$). **The statistic reading is**
 153 **refuted:** statistic and decision disturbance need not change suppression once the adversarial coeffi-
 154 cient share is held fixed. *On this arm the zero is a floor:* Krum admits no adversary at any rung, so
 155 the dissociation is read across all four aggregators, each under its own committed attack:

Aggregator	Kind	Δ dec.	Δ adm.	$\Delta\Lambda_a$	Δ ASR
Krum	selector	0.733	0.000	0.000	-0.026
Reputation	weighted averager	1.000	0.000	0.020	+0.178
Cosine-Krum	cosine-invariant selector	0.000	0.000	0.000	-0.425
Coord. median	coordinate order statistic	0.482	0.000	0.033	+0.098

157 Δ adm. is the *support* of the admitted adversarial mass: it never moves in any of the 60 measured
 158 rounds of any arm, including the 48 per arm where reputation and coord.median carry adversarial
 159 mass at baseline. $\Delta\Lambda_a$ is the separate question of how much weight the adversary's direction carries.
 160 cos_krum's -0.425 is the full-ladder endpoint, below our accuracy floor; the witness we rest on is
 161 its first, gate-passing rung (0.173, §5); App. C adds EMNIST-byclass and the score-only control.

162 **What the instrument does not hold fixed, and the pre-registered control that closes it.** Mode S
 163 does not identify the effect of statistic disturbance in isolation: it identifies the effect of the upstream
 164 intervention while holding the adversarial coefficient share fixed, which is what lets statistic and
 165 decision disturbance be *dissociated* from suppression rather than shown to be inert. Mode S pins
 166 the *adversarial* coefficients and nothing else: the dose also moves Krum's pairwise distances, *which*
 167 benign client it selects, and the **magnitude and geometry of the update actually aggregated**. That
 168 third channel is measured: App. C's Δ agg. column, the aggregate the defense *emits*, reads **0.892**
 169 here (App. D). **Score-only Krum** closes it: Krum scores the transformed stack and so makes the
 170 same selection, but the *untransformed* selected update is aggregated: the decision channel preserved,
 171 the magnitude and geometry change removed. Rules were frozen at 35788d9 before any score-only
 172 ASR existed, $\Delta > +0.15$ meaning the published result was partly a magnitude artifact. We find
 173 $\Delta = -0.023$ (0.062 $\rightarrow$ 0.039) against -0.026 uncontrolled on the identical cell, at the same
 174 Δ dec. of 0.733, the refuting threshold 0.212 nowhere approached and every rung clearing the
 175 accuracy floor, so the central claim is instrumented against two independent confounds, attenuation
 176 and magnitude (App. E).

177 **A pre-registered replication that refutes the positive half, and the sign reversal.**
 178 coord.median under the pixel backdoor is *rank-discordant*: of the interpretable arms it has the
 179 smallest decision change (0.482) and the largest admission-level change ($\Delta\Lambda_a=0.033$), so the two
 180 readings predict opposite rank positions for its rise, and thresholds were frozen against the two pub-
 181 lished rises (confirm if $\Delta > +0.178$, refute if $\Delta < +0.150$). We find $\Delta = +0.098$, which **refutes**
 182 **the admission ordering:** the arm with the largest admission-level change rose *less* than one with a
 183 smaller change. The decision reading does not predict it either ($R^2=0.239$ pooled). What replicates
 184 is the negative on a coordinate-wise order statistic: a decision change of 0.482 costs no suppression
 185 by the frozen margin, at accuracy 0.698. Identity ASR here is 0.443, so equivalence is to the *identity*
 186 *rung*; the rise is positive in 5/5 seeds (post-hoc): small and consistent, not zero.

187 **The design choice, not the defense, sets the inferred direction.** On this identical cell (matched
 188 displacement, decision change and clean accuracy, Table 3), closing one channel reverses the sign,
 189 though the two are not the same intervention and no p -value attaches to the difference between them
 190 (Table 3):

Design	The adversarial coefficient	Δ ASR
Outcome-gated ladder (§5)	free to attenuate	-0.272 ($p_{\downarrow}=0.0015$)
Mode S intervention (§4)	pinned at $c=1$	+0.098

192 **The confounded design does not lose power, it gets the sign wrong.** Closing the attenuation chan-
 193 nel inverts the effect rather than sharpening it: a criterion calibrated against the confounded contrast

would have been calibrated against the confound, so §3’s boundary is not a technicality about power.
Downstream-statistic preservation is not, in general, sufficient to establish attack-suppression preservation (Prop. 2): no predicate of the statistic alone determines the admitted adversarial mass, and none of that mass alone determines suppression: the title claim, by counterexample *within this class* (Table 1).

5 What the Boundary Licenses, and What We Found

We validate on CIFAR-10 with a 4-layer CNN, $N=10$ clients, $f=0.2$ adversarial, two committed attacks (model-scaling, pixel backdoor), at $n=5$ seeds per cell; *ASR* means *max-committed ASR*; the rest is fixed in App. X. The core evidence is a seed-matched intervention on one cell, not a mean comparison across cells.

The confounded ladder that motivated the instrument, and where its sign disagrees. Our earlier suites tried to validate C2 by *choosing a different downstream defense*, the move Prop. 4 rules out. Varying only the *dose* at fixed d_2 and attack ($\kappa=0$ bit-identical to d_2 standalone) refutes that suite’s pooled prediction that disturbance predicts each arm’s rise: $\rho_s=+0.351$, one-sided $p=0.091$ against a frozen $p < 0.05$, and in three of four arms *ASR falls* as the dose rises (80 runs, one arm confirming decisively; Table 5). **The fault is ours and is measured:** at fixed $\text{mean}(c)$ a large spread mostly *dilutes* whichever client carries the poison, so this is attenuation, not statistic disturbance (App. M, R).

The dissociation holds across four aggregators and replicates on a second dataset. Applying the instrument to that ladder, most of the one confirmation was attenuation: with the adversary’s share pinned, reputation’s rise at $\rho=54.6$ shrinks from $+0.762$ to $+0.178$, and what remains tracks a small *admission* change, not the large decision change (App. J.5). The class-(a) control refutes us the other way: *cos_krum*’s selection is bit-identical at every rung yet *ASR* is 0.584 at the identity (it never suppresses) and *falls* 0.173 at the first accuracy-gate-passing rung (one-sided $p=0.044$, 5/8 seeds); the full-ladder fall sits below our accuracy floor, so **the gate-passing rung is the witness**. **Mode A** (payload share at fixed dispersion) separates Thm. 1’s two regimes: confirmed for *coord_median*, refuted for reputation. A *pre-registered* replication on EMNIST-byclass/simple_cnn moves the decision on 0.643–0.714 of rounds with the adversary never admitted and suppression inside the frozen ± 0.15 margin ($\Delta = -0.017$ beside CIFAR-10’s -0.026): **replicated** on a second dataset and architecture, never “generalizes” (App. C, K).

And what a screen then buys. Only the 13/20 prospective sub-tier risks falsification, limited by our menu: FLTrust suppresses both attacks alone (0.048), so its pairs are out of scope or near-identities, and two further attempts to validate the *screen* were defeated by the same confound. **A third closed a census.** The region Prop. 1(a) settles *exactly* is 3×3 ; its 2 unspent cells and 2 measured-C2-fail controls were labelled at 96f8856, and **all four landed as registered**. It holds **one** low-ASR pair, $\text{rep} \rightarrow \text{cos_krum}$ at 0.263, which neither constituent suppresses alone, **and the criterion selects it**: C2 alone admits 9 (11%, 90 runs), $C1 \wedge C2$ admits 1 (100%/100%, 10). **Where invariance is exact the trade reverses**, on one accuracy-gated cell at $n=3$ whose C1 input is itself seed-count-contingent (33% ungated; App. K.2). **Scale:** the PASS predictions hold at $N=100/\text{ResNet18}$, far from converged accuracy. Under *one* adaptive stress test $\text{FG} \rightarrow \text{RFA}$ rises $2.4 \times (n=30)$: **inherited robustness is not adaptive robustness** (App. K, O, S).

6 Conclusion

The finding. Downstream-statistic preservation is not sufficient to identify preserved attack suppression. An intervention changes a defense’s statistic and decision substantially without changing the adversarial mass it admits or the suppression that follows; conversely a bit-identical statistic and decision coexist with a change in suppression, each direction with a measured witness.

Why. An outcome-gated cross-defense comparison conditions on the same outcome it seeks to explain, so it cannot separate mechanism preservation from standalone effectiveness: a structural identification boundary (Prop. 4) no amount of additional data within that design removes. Within-

defense intervention fixes d_2 and the attack instead: on the identical cell **the confounded design does not lose power, it gets the sign wrong.**

The implication. Evaluate composed defenses by measuring the adversarial contribution that survives the downstream mechanism, not by inferring security from the statistic it consumes. That object is (P4), which Mode S reads only by knowing adversary identity, so we identify the channel rather than screen for it (App. J.3). The open question: **can (P4) be estimated without knowing who is malicious?**

References

- E. Bagdasaryan, A. Veit, Y. Hua, D. Estrin, and V. Shmatikov. How to backdoor federated learning. In *International Conference on Artificial Intelligence and Statistics*. PMLR, 2020.
- A. N. Bhagoji, S. Chakraborty, P. Mittal, and S. Calo. Analyzing federated learning through an adversarial lens. In *International Conference on Machine Learning*. PMLR, 2019.
- P. Blanchard, E. M. El Mhamdi, R. Guerraoui, and J. Stainer. Machine learning with adversaries: Byzantine tolerant gradient descent. In *Advances in Neural Information Processing Systems*, 2017.
- X. Cao, M. Fang, J. Liu, and N. Z. Gong. FLTrust: Byzantine-robust federated learning via trust bootstrapping. In *Network and Distributed System Security Symposium (NDSS)*, 2021.
- X. Chen, C. Liu, B. Li, K. Lu, and D. Song. Targeted backdoor attacks on deep learning systems using data poisoning. *arXiv preprint arXiv:1712.05526*, 2017.
- F. Croce, M. Andriushchenko, V. Sehwag, E. Debenedetti, N. Flammarion, M. Chiang, P. Mittal, and M. Hein. RobustBench: A standardized adversarial robustness benchmark. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2021.
- F. Elwert and C. Winship. Endogenous selection bias: The problem of conditioning on a collider variable. *Annual Review of Sociology*, 40:31–53, 2014.
- C. Fung, C. J. M. Yoon, and I. Beschastnikh. The limitations of federated learning in sybil settings. In *23rd International Symposium on Research in Attacks, Intrusions and Defenses (RAID)*, pages 301–316, 2020.
- P. Kairouz, H. B. McMahan, B. Avent, A. Bellet, M. Bennis, A. N. Bhagoji, K. Bonawitz, Z. Charles, G. Cormode, R. Cummings, et al. Advances and open problems in federated learning. *Foundations and Trends in Machine Learning*, 14(1–2):1–210, 2021.
- H. B. McMahan, E. Moore, D. Ramage, S. Hampson, and B. A. y. Arcas. Communication-efficient learning of deep networks from decentralized data. In *Proceedings of the 20th International Conference on Artificial Intelligence and Statistics (AISTATS)*, volume 54, pages 1273–1282. PMLR, 2017.
- T. D. Nguyen, P. Rieger, H. Chen, H. Yalame, H. Möllering, H. Fereidooni, S. Marchal, M. Miettinen, A. Mirhoseini, S. Zeitouni, F. Koushanfar, A.-R. Sadeghi, and T. Schneider. FLAME: Taming backdoors in federated learning. In *USENIX Security Symposium*, 2022.
- J. Pearl. *Causality: Models, Reasoning, and Inference*. Cambridge University Press, 2nd edition, 2009.
- K. Pillutla, S. M. Kakade, and Z. Harchaoui. Robust aggregation for federated learning. *IEEE Transactions on Signal Processing*, 70:1142–1154, 2022.
- V. Shejwalkar and A. Houmansadr. Manipulating the byzantine: Optimizing model poisoning attacks and defenses for federated learning. In *Network and Distributed System Security Symposium (NDSS)*, 2021.
- F. Tramèr, A. Kurakin, N. Papernot, I. Goodfellow, D. Boneh, and P. McDaniel. Ensemble adversarial training: Attacks and defenses. In *International Conference on Learning Representations (ICLR)*, 2018.

- 290 H. Wang, K. Sreenivasan, S. Rajput, H. Vishwakarma, S. Agarwal, J.-y. Sohn, K. Lee, and D. Pa-
291 pailiopoulos. Attack of the tails: Yes, you really can backdoor federated learning. In *Advances in*
292 *Neural Information Processing Systems (NeurIPS)*, volume 33, 2020.
- 293 C. Xie, K. Huang, P.-Y. Chen, and B. Li. DBA: Distributed backdoor attacks against federated
294 learning. In *International Conference on Learning Representations*, 2020.
- 295 D. Yin, Y. Chen, K. Ramchandran, and P. L. Bartlett. Byzantine-robust distributed learning: To-
296 wards optimal statistical rates. In *Proceedings of the 35th International Conference on Machine*
297 *Learning*, 2018.
- 298 Z. Zhang, A. Panda, L. Song, Y. Yang, M. W. Mahoney, P. Mittal, K. Ramchandran, and J. E.
299 Gonzalez. Neurotoxin: Durable backdoors in federated learning. In *International Conference on*
300 *Machine Learning (ICML)*, 2022.

301 Table 1 sits ahead of the appendices because it governs all of them: it is the one place each scope
302 condition is stated, and no later section restates them at each use.

Table 1: **Scope of claims.** The one place each scope condition is stated. **Claims do not transfer upward:** the formal guarantee does not certify the general framing, and the empirical figures certify neither level above them.

kind	content	status, and what it does <i>not</i> license
General	An evaluation gated on the outcome it predicts cannot identify the mechanism’s contribution; admissible designs are within-defense (Prop. 4, Prop. 5).	Proved, and not specific to FL, to backdoors, or to our conditions. Licenses no claim that a composition <i>is</i> secure, nor that mechanism preservation implies suppression.
Formal	Mechanism preservation as <i>invariance</i> under positive per-client rescalings $T(u_i) = c_i u_i$, $c_i > 0$, into specific rank and geometric aggregators (Thm. 1, Lem. 1, Cor. 1).	Proved on that class only: a <i>laboratory</i> for composition, not a model of FL defenses in general. A d_1 that alters participation, carries cross-round state, mixes clients, or aggregates first has no such T and is covered by nothing here (C0, applied retrospectively to arms frozen before it). Bounds no ASR , and Regime A’s condition is not operationally checkable (S is not a stable per-attack constant).
Empirical	§4’s interventions and §5’s 42-pair menu: CIFAR-10, CifarCNN/ResNet18, $N \in \{10, 100\}$, $f \in [0.05, 0.6]$, Dirichlet $\alpha=0.5$, star topology , three attack families; one arm on EMNIST-byclass/simple_cnn.	Measured at $n=5$ seeds per cell ($n=8$ for <code>cos_krum</code> , $n=3$ on EMNIST-byclass), fixed before the runs, so intervals are wide and no null is read as flatness without a pre-registered equivalence margin. Dataset transfer is tested on the flagship arm alone: “replicated on a second dataset and architecture,” never “generalizes”. For the screen, dataset transfer on EMNIST-byclass is out of scope rather than merely untested: no single defense measured there suppresses the committed pixel backdoor alone (minimum 0.609), so C1’s per-attack precondition fails over the menu measured (§K.3). “Refuted” always means a pre-registered <i>directional</i> hypothesis was rejected against a rule frozen for it, never that a mathematical claim was disproved.
Both instruments are oracles	Mode S and Mode A build c_i from adversary identity, as does Δ_c .	They identify a causal channel and are not screens ; computing (P4) without adversary identity is our open problem, and admission is a <i>direction to measure</i> , not a validated predictor. C1 and all built on it assume a committed adversary: what is bounded is <i>inherited</i> robustness under a fixed attack, never adaptive robustness .

303 A The composition model and the invariance classification

304 **The two witnesses under Definition 1, and which rung each rests on. (P1)⇒(P2)⇒(P3) hold**
305 **by definition: logical implications, not empirical claims;** their content is that they stop at (P3), and
306 §2’s table states both breaks below it. *Holding* decision *and* admission invariance does not deliver
307 suppression: `cos_krum` falls 0.173 by the first rung alone through the magnitude channel, at $\rho=2.72$,
308 where clean accuracy moves only $0.599 \rightarrow 0.567$ and our accuracy gate passes (Appendix N).
309 Across the whole ladder it falls 0.425, but that endpoint sits below our accuracy floor, so **the gate-**
310 **passing rung is the witness we rest on (§5).** *Losing* decision invariance does not cost suppression:
311 Krum’s selection flips in up to 80% of rounds under an intervention pinning the adversarial mass, and
312 its ASR stays inside the pre-specified equivalence margin ($|\Delta|=0.026 < 0.15$) (§4). So a criterion
313 stated at (P1)–(P2) constrains security in *neither* direction, **(P4), where the security content sits, is**
314 **not sufficient for (P5) either, and the hierarchy answers “which level does a security claim have**
315 **to be stated at?”, not “which level predicts ASR?”**

316 **Definition 2** (Defense composition). *Given defenses d_1 and d_2 , composition ($d_1 \rightarrow d_2$) operates*
317 *as: d_1 produces per-client transformed updates $U' = \{T_{d_1}(u_i)\}_{i=1}^K$; d_2 aggregates U' using its*
318 *standard algorithm. For FoolsGold→CoordMedian, FG outputs $\{w_i u_i\}$ and CoordMedian takes*
319 *the unweighted coordinate-wise median of the K transformed values, not a “weighted median”.*
320 *Clients with $w_i = 0$ are **not removed** but reach d_2 as zero vectors, so at least one point of U' is*
321 *always the origin, which is what Lemma 1 is about.*

322 Each defense suppresses malicious updates through a *discriminative property* (a structural feature
323 separating adversarial from benign contributions) and the property is defense-class-specific: *simi-*
324 *larity spread* (FoolsGold), *coordinate ordering* (TrimmedMean, CoordMedian), *residual magnitude*
325 (RFA), *consensus distance* (Reputation, in ℓ_2 from the coordinate-wise median) and *norm magni-*
326 *tude* (NormClip). Mechanism preservation is then governed by whether d_1 ’s transform leaves d_2 ’s
327 property intact, and because every per-client defense we consider acts as a positive scalar rescaling,
328 that question has a sharp answer:

Proposition 1 (Invariance classes under positive per-client rescaling). *Let $T(u_i) = c_i u_i$ with $c_i > 0$, strictly: the form taken by NormClip ($c_i = \min(1, \tau/\|u_i\|)$), Reputation ($c_i \propto e^{-d_i/s}$) and RFA-reweighting ($c_i \propto 1/\max(\|u_i - \mu'\|, \varepsilon)$), each positive by construction. **FoolsGold is deliberately not in that list:** its shift-and-clamp gives the least-trusted client weight exactly 0 in every round (2 of 5 clients here), so FoolsGold-as- d_1 falls outside the hypothesis and is governed by Lemma 1 rather than by (a)–(c). Then:*

- (a) **Cosine similarity is exactly invariant:** $\cos(c_i u_i, c_j u_j) = \cos(u_i, u_j)$. Hence FoolsGold’s weight vector is unchanged, $w(T(U)) = w(U)$, and **C2 holds for every $\cdot \rightarrow$ FoolsGold composition.**
- (b) **Coordinate ordering is not invariant** in general; it survives under the bounded-ratio condition of Theorem 1(1).
- (c) **Consensus distance and residual magnitude are not invariant:** heterogeneous c_i can invert their ordering, so $\cdot \rightarrow$ Reputation compositions are C2 failures and $\cdot \rightarrow$ RFA requires the separation condition of Theorem 1(2).

So mechanism preservation is a question about the *invariance class* of the downstream signal, and the classification corrects a natural but wrong intuition: norm clipping does not blind FoolsGold, though it does blind Reputation, whose signal is consensus distance rather than similarity.

The framework in general form. The structure does not depend on rescaling. Write d_2 ’s discriminative property as a statistic $S(U)$ and let T be d_1 ’s per-client transform; then $(d_1 \rightarrow d_2)$ is **mechanism-preserving** when $S(T(U)) = \varphi(S(U))$ for some φ preserving the ordering d_2 acts on. C2 is exactly that requirement, C3 asks in addition that φ not compress the gap below d_2 ’s operating threshold, and Proposition 1 is the case $T(u_i) = c_i u_i$ (Appendix I), so the conditions are algebraic rather than experimental. It also states the real limit: **a d_1 whose transform is not per-client (mixing updates, or carrying cross-round state) has no T of this form, and nothing here applies.**

The full causal chain. Definition 1 is a prefix of a longer one: statistic $\rightarrow$ decision $\rightarrow$ admission $\rightarrow$ aggregate adversarial influence $\rightarrow$ per-round update $\rightarrow$ persistence $\rightarrow$ ASR. **We measure the four emphasized links, not update geometry or persistence dynamics:** an acknowledged gap, not an assumed-benign one. A selection flipping from one benign client to another disturbs the statistic maximally and the mechanism not at all, which is why (P3) and (P4) come apart. **“Mechanism” always means the downstream defense’s discriminative operation, never its security behaviour,** that implication being what this paper refutes.

B The signal-preservation conditions in full

Definition 3 (Signal-Preservation Conditions). *Composition $(d_1 \rightarrow d_2)$ is predicted to suppress attack a when: (C0) d_2 ’s mechanism is load-bearing, $\text{ASR}(d_1, a) \geq 0.5$ for some committed a : a condition on the criterion’s domain that **never certifies a pair**, and one we formulated only after §5 made it bind, so it is applied **retrospectively** and is **not part of what the 42 pairs validate**; (C1) some constituent suppresses every committed attack alone, $\min_{d \in \{d_1, d_2\}} \text{ASR}(d, a) < 0.5$ for all a , defined through the same threshold that constitutes the outcome, hence an **empirical prerequisite and not a mechanistic property**; (C2) d_1 ’s transform does not destroy d_2 ’s discriminative property, decided from definitions alone by Proposition 1; (C3) it does not compress the quantitative margin below d_2 ’s operating threshold. **The prioritization rule is $C1 \wedge C2$, C3 demoted to a margin diagnostic (Appendix U); the pre-registered labels of §5 were frozen with C3 in the gate and are not relabelled. Theorem 1 bounds no ASR: that mechanism preservation plus C1 predicts low ASR is a separate empirical hypothesis, tested in §5, and neither stage alone establishes low ASR, nor does their conjunction (stage structure, Appendix B).***

Definition 3 in the body states the four conditions and every disclosure attached to them; what follows is the same criterion with its *stage structure*, which is worth separating because conflating the stages is the main source of confusion about what the framework claims.

Composition $(d_1 \rightarrow d_2)$ suppresses attack a when the conditions below hold. The criterion is a **two-stage screen**, and we label the stages because conflating them is the main source of confusion about what the framework claims: stage 1 is empirical screening of the *constituents*, stage 2 is mechanistic screening of the *composition*.

381 **Stage 0: applicability (C0).** Whether the framework has anything to say about this pair at all. This
 382 is a condition on the criterion’s *domain*, and Lemma 2 is stated over it together with C1.

383 **(C0) The downstream mechanism is load-bearing.** $\text{ASR}(d_1, a) \geq 0.5$ for at least one committed
 384 attack a . C2 and C3 concern preservation of d_2 ’s mechanism, so they carry information only when
 385 that mechanism does work the composition needs: if d_1 alone suppresses everything, the composi-
 386 tion inherits that suppression and a correct prediction is correct for a reason the framework did not
 387 supply. The asymmetry is deliberate: a strong d_2 is the intended case. **Failing C0 removes a pair**
 388 **from the criterion’s domain; it never certifies one.** *Provenance:* C0 was vacuous on the menu
 389 the framework was developed on (strongest single defense 0.519) and was formulated only after the
 390 prospective suite of §5 made it bind, so it is applied **retrospectively** and is **not part of what the 42**
 391 **pairs validate** (Appendix Q).

392 **Stage 1: empirical constituent screening (C1).** Read off standalone evaluations. This screens; it
 393 does not explain, and it is not a mechanistic claim.

394 **(C1) Individual effectiveness.** For *every* committed attack a in the evaluated set,
 395 $\min_{d \in \{d_1, d_2\}} \text{ASR}(d, a) < 0.5$. The quantifier matters: our headline metric is the *maximum*
 396 over attacks, so requiring a suppressing constituent for only one attack would license compositions
 397 that fail on the other. C1 is defined through the same ASR threshold that constitutes the outcome,
 398 so it is an **empirical prerequisite, not a mechanistic property**: the one condition that cannot be
 399 checked without running something.

400 **Stage 2: mechanistic composition screening (C2, with C3 as a diagnostic).** Determined by d_1 ’s
 401 and d_2 ’s definitions alone, no experiment required, and what Theorem 1 proves. **The evaluation-**
 402 **prioritization criterion is $C1 \wedge C2$,** with C3 retained as a **margin diagnostic** rather than a gate for
 403 the reasons scored in Appendix U. *The pre-registered labels of §5 were frozen with C3 in the gate*
 404 *and are not relabelled*, and the demotion changes no reported number.

405 **(C2) Discriminative-property preservation.** d_1 ’s transform does not destroy d_2 ’s discriminative
 406 property, as determined by Proposition 1: automatic for cosine-similarity downstream defenses, condi-
 407 tional for coordinate ordering, and violated for consensus distance under heterogeneous rescaling.

408 **(C3) Margin preservation.** d_1 does not compress the quantitative margin below d_2 ’s operating
 409 threshold. For rank aggregators: extremeness gap survives reweighting. For geometric aggrega-
 410 tors: separation condition holds on transformed points.

411 **What is proved and what is hypothesised.** Stage 1 is empirical and outcome-defined, stage 2
 412 mechanistic, and **neither alone establishes low ASR, nor does their conjunction**: Theorem 1 es-
 413 tablishes that $C2 \wedge C3$ preserve d_2 ’s discriminative mechanism on the reweighted inputs and **bounds**
 414 **no ASR**. That mechanism preservation plus C1 predicts *low ASR* is a separate empirical hypothesis,
 415 tested in §5.

416 C The four channels, all four aggregators

417 Endpoint change ($\kappa=0$ against $\kappa=2$) in each measured channel under Mode S, for ev-
 418 ery downstream defense in the paper. **The first column, Δ agg., is the relative**
 419 **displacement of the aggregate the defense emits**, $\|\text{agg}(T(U)) - \text{agg}(U)\|/\|\text{agg}(U)\|$:
 420 what enters training, and **not** the value of a scoring statistic. The measured chan-
 421 nels are therefore aggregate displacement, decision, admission and influence. **Recomputed,**
 422 **not transcribed:** every cell is emitted by `experiments/build_channel_table.py` from
 423 `admission_measurement.json`, `femnist_admission.json`, `targeted_dose/summary.json`,
 424 `dose_replication/summary.json` and `dose_femnist/summary.json` under `results/`, which
 425 the script only reads; it also asserts that the identity rung displaces nothing (0.00e+00 for all
 426 six rows). **The sixth row is the score-only control** (Appendix E): the first row’s arm with the
 427 magnitude channel closed, so its aggregate column is the re-selection component alone (from
 428 `displacement_decomposition.json`) and its decision, admission and influence columns are
 429 the first row’s own numbers, score-only leaving the selection untouched; its ASR comes from
 430 `score_only/summary.json`, which the emitter refuses to merge into the ASR sources it shares
 431 cell keys with, so it cannot shadow the first row’s published number. The influence column is
 432 **aggregator-specific by necessity** (a selector has no weights) and its instantiation is: `adversarial`

433 share of the Weiszfeld/reputation weight mass for the averager, adversarial share of selected co-
 434 ordinates for the order statistic, and the indicator that an adversarial update is selected at all for
 435 the two selectors. Earlier drafts labelled this arm “FEMNIST” in the tables while the prose said
 436 EMNIST-byclass: the loader reads torchvision EMNIST(split="byclass") under a femnist key
 437 and partitions it by the label-Dirichlet split used throughout, *not* LEAF’s by-writer split, so the ta-
 438 bles now read EMNIST-byclass while the frozen artifact and pre-registration filenames keep the
 439 femnist key. The EMNIST-byclass row is the second-dataset replication of §5; its channels were
 440 measured prospectively, before its ASR existed.

Aggregator	Kind	Δ agg.	Δ dec.	Δ adm.	$\Delta\Lambda_a$	Δ ASR
Krum	selector (pairwise distance)	0.892	0.733	0.000	0.000	−0.026
Reputation	weighted averager (cross-round state)	0.659	1.000	0.000	0.020	+0.178
441 Cosine-Krum	selector (cosine distance)	1.297	0.000	0.000	0.000	−0.425
Coord. median	coordinate-wise order statistic	0.670	0.482	0.000	0.033	+0.098
Krum (EMNIST-byclass)	selector (pairwise distance)	0.842	0.643	0.000	0.000	−0.017
Krum, score-only	selector, magnitude channel closed	0.986	0.733	0.000	0.000	−0.023

442 **What this table does and does not show.** It shows that the dissociation between the decision
 443 channel and the admission channel is not a property of one aggregator: it holds for a selector, for a
 444 cosine-invariant selector, for a coordinate-wise order statistic and for a weighted averager, on two
 445 datasets. It does **not** order the arms ($\Delta\Lambda_a$ and Δ ASR disagree in rank, which is the refutation of §4
 446 stated in one row-pair) and it is *not* an out-of-sample test: every row is a Mode-S ladder reported
 447 elsewhere in this paper.

448 D Decomposing the emitted displacement: re-selection against rescaling

449 The Δ agg. column above mixes two things the causal claim needs apart. For a selector, $\text{agg}(U) =$
 450 u_{s_0} and $\text{agg}(T(U)) = c_{s_1} u_{s_1}$, with s_0 the client selected on the raw stack and s_1 the client selected
 451 on the transformed one, so exactly

$$c_{s_1} u_{s_1} - u_{s_0} = \underbrace{(u_{s_1} - u_{s_0})}_{\text{re-selection}} + \underbrace{(c_{s_1} - 1) u_{s_1}}_{\text{rescaling}},$$

452 both components reported relative to $\|u_{s_0}\|$. They are **not orthogonal** (the measured angle runs
 453 from $\cos = -0.99$ to -0.43 , mean ≈ -0.65), so their norms need not sum to the total, and the
 454 table prints the angle rather than leaving that discrepancy unexplained. What is asserted instead
 455 is the exact vector identity: worst residual $2.66\text{e-}07$ over all 240 (arm, family, rung, seed, round)
 456 cells, float32 noise. **This decomposes the published number, not a second similar one:** every
 457 recomputed total agrees with the corresponding Δ agg. value in `admission_measurement.json`
 458 to $0.00\text{e}+00$, on the same seeds, rounds and configuration. Selectors only: a weighted averager
 459 and a coordinate-wise order statistic emit no single selected client, so the split is undefined for
 460 them. `experiments/measure_displacement_decomposition.py` trains no ASR, writes a new
 461 file, and leaves every frozen artifact byte-identical.

Design	Selector	κ	total	re-selection	rescaling	cos	sel. chg.
Confounded ladder	krum	0	0.000	0.000	0.000	—	0.000
		0.5	0.637	0.685	0.333	−0.65	0.600
		1	0.731	0.851	0.685	−0.66	0.733
		2	0.880	1.675	1.674	−0.72	0.867
Instrument (Mode S)	krum	0	0.000	0.000	0.000	—	0.000
		0.5	0.663	0.646	0.308	−0.64	0.533
		1	0.847	1.075	0.666	−0.69	0.800
		2	0.892	0.986	0.973	−0.66	0.733
Confounded ladder	cos_krum	0	0.000	0.000	0.000	—	0.000
		0.5	0.270	0.000	0.270	—	0.000
		1	0.515	0.000	0.515	—	0.000
		2	0.834	0.000	0.834	—	0.000
Instrument (Mode S)	cos_krum	0	0.000	0.000	0.000	—	0.000
		0.5	0.395	0.000	0.395	—	0.000
		1	0.754	0.000	0.754	—	0.000
		2	1.297	0.000	1.297	—	0.000

What it says, and what it does not. Mode S into Krum at $\kappa=2$: the published 0.892 is 0.986 of re-selection and 0.973 of rescaling, so the dose moves *which* benign client is aggregated and *how large* it arrives, in comparable measure. Of the 11/15 rounds whose selection changes there, **none** lands on an adversary: the quantitative form of the claim that Krum’s decision changes are benign-to-benign, and the reason the support of the admitted adversarial mass changes in **none of the 60 measured rounds** while the decision channel moves by 0.733, on this arm no adversary is admitted at *any* rung, so that count is a floor, and the arms carrying adversarial mass at baseline are reputation and `coord_median` (48 of 60 rounds each, App. C). `cos_krum` is the informative contrast: its statistic is exactly invariant to positive rescaling, so 0/15 selections change at every rung of *both* designs and its entire displacement (1.297 under the instrument, 0.834 under the ladder) is rescaling. Its fall therefore cannot travel through the decision channel at all, which corroborates §5’s attribution of that fall to the magnitude channel from a quantity measured independently of any ASR. **This is a decomposition of a displacement, not of an ASR effect:** it says what the dose moves in the update entering training and makes no claim about suppression, and it is an analysis of runs already reported rather than a new experiment.

E The score-only control in full

§4 reports the primary outcome: $\Delta = -0.023$ (mean ASR $0.062 \rightarrow 0.039$, seeds 42–46) against -0.026 for the uncontrolled arm on the identical cell, inside the frozen ± 0.15 equivalence margin, with the accuracy floor cleared at every rung (0.565/0.629/0.624/0.630). The verdict wording is the pre-registration’s own for that branch: *the negative survives with the magnitude channel closed*.

What the arm is. Krum scores on the transformed stack (statistic and decision disturbed exactly as in the uncontrolled arm, since the scoring input is the same) and aggregates the *untransformed* selected update, so what enters training is bit-for-bit what a client produced. It is the reviewer-suggested control in its own terms: the selected update’s magnitude and direction are held fixed while the selection statistic is perturbed independently. Implementation is a defaulted `score_only=False` kwarg, so no existing call site changes and every frozen computation stays bit-identical; the suite’s own harness check verifies that on the published arms. The $\kappa=0$ rung is *imported* rather than re-run, because at the identity the transform returns the stack unwrapped and score-only Krum *is* Krum alone: an exact claim, verified bit-identically at one seed before the arm started.

Three disclosures, from the freeze rather than from hindsight. (i) The secondary Jonckheere–Terpstra test was specified to detect a monotone *rise* (the outcome that would have refuted flatness with attenuation already closed by construction) and the observed trend is a *fall* ($J=58.5$, $z=-1.110$, $p_{\uparrow}=0.867$, permutation $p_{\uparrow}=0.874$). We report it and do **not** score it as support for flatness; the pre-registration says in advance that at $n=5$ a null here is not evidence of flatness. (ii) The branches had unequal headroom. With the identity rung at 0.062, refuting us required mean $\text{ASR}(\kappa=2) > 0.212$ against a ceiling of 1 and was fully reachable, whereas the indeterminate branch

($\Delta < -0.15$) was *arithmetically unreachable*: ASR cannot fall below 0, so the largest possible fall on this cell is 0.062. (iii) **Selectors only.** A weighted averager and a coordinate-wise order statistic emit no single selected client, so “the selected update” is undefined for them and the control does not cover them: the `coord_median` arm carrying the +0.098 rise is therefore **not** magnitude-controlled, and neither is reputation’s +0.178. Score-only Krum is also an instrument and not a defense: it aggregates an update its own scoring stage never saw. And closing the magnitude channel still does not hold the model trajectory fixed, because which client is selected continues to change, by design, that being the channel left open.

The `cos_krum` corollary, and exactly what it settles. Section 5 of the same pre-registration asserted, before the run, that score-only `cos_krum` must be *bit-identical* across rungs: its statistic is exactly invariant under positive rescaling, so the selection cannot move, and with the emitted update untransformed the whole trajectory must repeat. It does (accuracy 0.528400 and ASR 0.003111 at all four rungs, $|\Delta| = 0.00e+00$ throughout) and the $\kappa=0$ value reproduces the published cell at that seed exactly, so the two pipelines agree at 0.000000 rather than approximately. **What the check does not do is confirm the arm’s published fall.** It runs one seed, as frozen, and that seed is unrepresentative of the arm: under the uncontrolled dose its own ladder *rises* +0.224 where the eight-seed mean *falls* −0.425. Closing the magnitude channel removes all 0.224 of that seed’s movement, which corroborates the *channel*, for `cos_krum` the emitted displacement is entirely rescaling (Appendix D), so no other route is open, but the attribution of the arm’s fall rests on that decomposition rather than on this check, and settling it empirically would need the arm’s other seven seeds. Expanding a pre-registered seed count after seeing the data is what pre-registration exists to prevent, so we report the invariance as confirmed and the arm-level attribution as **corroborated mechanically, not established at the arm mean.**

F Related work

The contrast in one sentence. Prior work asks whether a defense withstands an attack, and the composition literature treats composition quality as an *outcome to predict*. We ask a different question: the **identifiability of mechanistic explanations** for that outcome, under the evaluation designs commonly used to compare defenses. Concretely, whether an evaluation criterion can causally attribute a composition’s security to preservation of the downstream mechanism.

Byzantine-robust aggregation, studied in isolation. Krum [Blanchard et al., 2017], Trimmed Mean [Yin et al., 2018], RFA [Pillutla et al., 2022], FoolsGold [Fung et al., 2020], FLTrust [Cao et al., 2021] and FLAME [Nguyen et al., 2022] are each evaluated as a single stage; we propose no new defense and reason about *ordered pairs* of existing ones. Ensemble adversarial training [Tramèr et al., 2018] *assumes* upstream processing preserves downstream signals; we make that assumption explicit, checkable, and (as it turns out) **the wrong abstraction for security**.

Adaptive attacks are why an inherited-robustness guarantee is not a security guarantee. Optimized model poisoning defeats Krum, trimmed mean and FoolsGold once the attack is chosen *against* the defense [Shejwalkar and Houmansadr, 2021, Bhagoji et al., 2019], as do distributed and durable backdoors [Xie et al., 2020, Wang et al., 2020, Zhang et al., 2022]. C1 quantifies over a *committed* attack, so this literature bounds what any inherited-suppression criterion (ours included) can promise; our single whitebox stress test (§6) measures that gap rather than defending the criterion.

Evaluation methodology, which is where this paper sits. The adjacent precedent is adversarial robustness, where reported numbers proved unreliable until evaluations were required to be adaptive [Croce et al., 2021]. Our point is methodological in the same way but about **identification** rather than attack strength: composition studies gate candidates on standalone effectiveness [Kairouz et al., 2021] and then read the composition’s outcome as evidence about the mechanism, which Proposition 4 shows cannot work at *any* sample size. What is new is the level hierarchy at which the preservation intuition fails and the design class that can test it (Appendix J.4).

The attack side of the same literature engineers *persistence* rather than magnitude: model-scaling [Bagdasaryan et al., 2020] and Neurotoxin [Zhang et al., 2022] both aim at triggers that survive rounds of clean aggregation, which is the quantity we measure as γ (§1) and the reason temporal defense switching does not help in our regime (Appendix O). Nothing in this paper proposes a new aggregator, a new attack, or a new defense; the object of study is the *composition* of defenses

that already exist, and the contribution is an identification argument about how such compositions can and cannot be evaluated.

G Notation

symbol	meaning
c_i	upstream coefficient applied to client i 's update, $T(u_i) = c_i u_i$, $c_i > 0$
A	the set of adversarial participants in a round
ρ	statistic-level dispersion, $\max_i c_i / \min_i c_i$: <i>how far apart</i> the coefficients are
Δ_c	adversary-correlated dispersion, $\mathbb{E}[c_i \mid i \in A] - \mathbb{E}[c_i \mid i \notin A]$: <i>who</i> receives them
r	adversary-to-benign weight ratio, $w_a / \max_b w_b$; the Mode-A dial, $r = e^\nu$
S	the attack's coordinate extremeness (Theorem 1(1))
m	RFA separation margin, $\Delta'_{\text{sep}}/2(R'_B + \delta')$
Λ_a	aggregate adversarial influence, instantiated per aggregator: Weiszfeld/reputation
	weight share, selected-coordinate share, or the selection indicator
κ	the Mode-S / ladder dispersion dial, $\rho = e^{2\kappa}$
γ	attack <i>persistence</i> , $\text{ASR}(t+1)/\text{ASR}(t)$ after a clean round (≈ 0.66); never a weight

H Deployment guidance in full

Deployment guidance. (0) Check C0: if one candidate defense already suppresses every attack you have, deploy it alone: composition has nothing to add. (1) Check C1: does at least one defense suppress known attacks? (2) Check C2 by invariance class (Proposition 1): a cosine-similarity downstream defense has its signal preserved behind any positive rescaling, a consensus-distance one does not. (3) For weighting→aggregator compositions Theorem 1 gives verifiable sufficient conditions for C2∧C3. Gate on C2; read C3 as a *quantitative diagnostic*: a 5–10 round pilot returns ρ and the separation margin m , and it is $S > \rho$ (rank aggregators) or $m \gg 1$ (RFA) that matters, not whether d_1 happens to be heterogeneous, the binary reading that produced both of our false negatives. (4) Best candidate on our evidence, for further evaluation rather than deployment: **FG→CM** (accuracy 0.68, max ASR 0.107). FG→RFA is marginally lower on the point estimate (0.093 at $n=30$) but statistically indistinguishable, and carries the *highest* adaptive ASR (0.157) and an 18-point accuracy cost, so it is not the more secure choice on our evidence. (5) In our persistent-backdoor setting fixed-order composition substantially outperforms randomized ordering; we have not tested regimes where randomization might pay.

I Proofs

What the two arms jointly establish. Both halves are witnessed by runs already reported, so the following formalizes measurements rather than adding a claim.

Proposition 2 (Statistic-level non-identifiability). *Within the positive per-client rescaling class there exist transform families T_1, T_2 and downstream defenses d_2 such that (a) T_1 and T_2 induce arbitrarily different statistics and decisions at d_2 while inducing the same admitted adversarial mass and the same suppression; and (b) T_1 and T_2 induce the same admitted adversarial mass (indeed the same statistic and the same decision) while inducing different suppression. Consequently no predicate of the downstream statistic alone determines the admitted adversarial mass, and no predicate of the admitted mass alone determines suppression.*

Both parts are proved by exhibited witnesses (Appendix I): Mode S into Krum under model-scaling for (a), `cos_krum` for (b). **The corollary is the paper's title claim:** no metric defined *solely* on preservation of the downstream statistic can identify attack-suppression equivalence, (a) exhibiting statistic disturbance without suppression change and (b) the converse. **Scope:** a non-existence claim *within the positive-rescaling class on these defenses and attacks*, by exhibited counterexample: **not** that no metric of any kind can order security.

Proposition 3 (ρ and Δ_c are independently controllable). *Within the same class, the statistic-level dispersion $\rho = \max_i c_i / \min_i c_i$ and the adversary-correlated dispersion $\Delta_c = \mathbb{E}[c_i \mid i \in A] - \mathbb{E}[c_i \mid i \notin A]$ are not functionally related: there is a family with $\Delta_c \equiv 0$ exactly while ρ ranges over $[1, 54.6]$, and a family with the statistic's dispersion held fixed while Δ_c ranges over $[-1.138, +2.469]$. Hence neither quantity determines the other.*

593 This makes §2’s warning precise: **a transform reported by ρ alone leaves Δ_c entirely free**. It does
 594 *not* follow that Δ_c is a predictor; its magnitude fails to order suppression either (§5). The instrument
 595 is elaborate because §3 leaves nothing simpler. §5 answers what remains: *what happened when we*
 596 *tried it the obvious way first?*

597 I.1 Formal Guarantee: Weighting $\rightarrow$ Aggregator Class

598 For the practically important class containing our certified compositions (a per-client *weighting*
 599 defense d_1 followed by a *rank-based* or *geometric* aggregator d_2) C2 and C3 can be formally estab-
 600 lished. **Read what follows as a *boundary* theorem and deliberately not a *security* theorem**: it
 601 says exactly how much an invariance argument about composition can prove, which is preservation
 602 of d_2 ’s discriminative mechanism at levels (P1)–(P3). It bounds no ASR, and the counterexamples
 603 in the box above are counterexamples to reading it as though it did.

604 **Theorem 1** (Bounded reweighting preserves the downstream discriminative mechanism). *Let d_1*
 605 *apply per-client scalar weights $w_i \geq 0$ with ratio $\rho := \max_i w_i / \min_{i:w_i>0} w_i$ (over strictly positive*
 606 *weights), operating on n_b benign and n_a adversarial updates in a round of size K . Assume $w_a > 0$*
 607 *for the adversarial clients; the $w_a = 0$ case is Lemma 1. Then:*

608 1. **Rank aggregators.** *The composition passes $\{w_i u_i\}$ to an unweighted order statistic over*
 609 *K transformed values per coordinate. The mechanism is preserved when:*

- 610 (i) Count majority: $n_a < K/2$.
- 611 (ii) Signed ordering (C2): *on suppression-relevant coordinate k , the adversarial trans-*
 612 *formed value lies strictly beyond every benign one: $w_a u_{a,k} > \max_b w_b u_{b,k}$ (upper*
 613 *tail) or $w_a u_{a,k} < \min_b w_b u_{b,k}$ (lower tail).*

614 Sufficient magnitude condition for (ii): *if $\text{sign}(u_{a,k})$ is common to all adversaries and*
 615 *$|u_{a,k}| > \rho \cdot \max_b |u_{b,k}|$, then (ii) holds. This branch is one of two mechanism-preserving*
 616 *regimes, the other (strong attenuation, $w_a \rightarrow 0$) being Lemma 1.*

617 2. **Geometric median** (RFA): *the mechanism is preserved when separation dominates dis-*
 618 *placement: the minimum adversarial–benign separation among the transformed updates*
 619 *strictly exceeds twice the benign spread plus the displacement the adversaries induce in*
 620 *the geometric median (Appendix I defines the four quantities and proves it), under which*
 621 *every adversarial residual exceeds every benign one, so RFA assigns adversarial updates*
 622 *strictly lower inverse-residual weight.*

623 *Proofs are in Appendix I. Both parts are sufficient conditions for d_2 to preserve the downstream*
 624 *suppression ordering, not to exclude adversarial contributions outright; **neither is sufficient for low***
 625 *ASR, and part (2) constrains the aggregate adversarial coefficient only through Corollary 1.*

626 Theorem 1(2) orders residuals *per client*, which alone does not bound the adversary’s *total* share
 627 of the aggregate. A realized *separation margin* $m \geq 1$, the factor by which that condition is met,
 628 defined with the bound in Corollary 1 (Appendix I), does bound it, but only quantitatively: as $m \rightarrow 1$
 629 the bound degrades to the adversarial count fraction, so **the separation condition on its own is not**
 630 **enough**. This is the one place the theory predicts a number and the measurement lands inside it: at
 631 the measured $m \geq 2.82$ in every round with $w_a > 0$ (mean 4.71 over those rounds) the bound is
 632 $\Lambda_a \leq 12.6\%$, against a realized $\Lambda_a \leq 1.8\%$. The $m \rightarrow 1$ boundary is where §3’s heterogeneity
 633 failure sits, and it is reached ($m=1.01$ at $\alpha=0.1$).

634 Λ_a is level (P4) in measurable form. (P4) is the first level whose content is the adversary’s
 635 influence on the aggregate, and for a geometric aggregator Λ_a is that level made measurable:
 636 Corollary 1 bounds it, Lemma 1 drives it to a benign-only convex combination, and §4’s instru-
 637 ments intervene on it directly. **So the level the evidence points to is one the theory can bound,**
 638 **just not from a (d_1, d_2) definition pair alone**, which is what a screen would need. Unqualified,
 639 “admission” means *admitted adversarial mass*, whether an adversarial update enters d_2 ’s aggregate
 640 at all; the three aggregator-specific quantities Definition 1 sets out are never used interchangeably
 641 (Appendix G).

642 The remaining case dominates in practice: FoolsGold drives the adversarial weight to zero. This
 643 does *not* remove the adversary from RFA’s input (Definition 2): RFA receives a zero vector, a point
 644 of the geometric-median problem like any other. The correct statement is stronger than removal:

645 **Lemma 1** (Zero-weight annihilation). *Our implementation runs the ε -smoothed Weiszfeld iteration,*
646 *whose fixed point $\mu' = \sum_i \lambda_i u'_i$, $\lambda_i \propto 1/\max(\|u'_i - \mu'\|, \varepsilon)$, is a convex combination for every*
647 *configuration, so no non-degeneracy assumption is needed (Appendix I). If $w_a = 0$ for every ad-*
648 *versarial client then $u'_a = 0$ for each, and writing $\Lambda_a := \sum_{a \text{ adv}} \lambda_a$ the adversarial terms vanish*
649 *identically:*

$$\mu' = \sum_{b \text{ benign}} \lambda_b u'_b = (1 - \Lambda_a) \bar{u}, \quad \bar{u} := \sum_b \frac{\lambda_b}{1 - \Lambda_a} u'_b,$$

650 *a convex combination of benign transformed updates. The aggregate is a **contraction of a benign-***
651 ***only weighted average**: u_a does not appear in μ' at all, and the origin point's only effect is a smaller*
652 *effective step: **a utility cost, not a security failure**.*

653 The same argument covers CoordMedian. The two cases are exhaustive over every measured round
654 (FG zeroes all adversaries in 32% of adversary-present rounds, the separation condition holding in
655 the other 68%), and the mechanism survives for *opposite* reasons: under Lemma 1 the adversaries
656 hold most of the Weiszfeld mass ($\Lambda_a = 0.71$) but carry no payload; under Corollary 1 they carry
657 payload but at most 12.6% of it. The contraction $1 - \Lambda_a = 0.29$ is composition's accuracy cost:
658 FG→RFA reaches accuracy 0.50 and FG→CM 0.66 against 0.78 for the aggregators alone, paid
659 even when the zeroed clients are benign.

660 The proofs below are short; they are collected here rather than inline because what each result *rules*
661 *in or out* is the content, not the derivation. Both witnesses of Proposition 2 are runs reported in §4
662 and §5: Mode S into Krum for part (a), `cos_krum` for part (b).

663 *Proof of Proposition 4.* The load-bearing condition excludes d_1 from attaining the minimum in (i),
664 so certification forces $M(d_2, a) < \theta$; both cells with $M(d_2, a) \geq \theta$ are unoccupied whatever the
665 value of P . A contrast varying d_2 therefore compares pairs differing in P only inside the selected
666 region $M(d_2, a) < \theta$, where P and $M(d_2, a)$ are both functions of the same choice of d_2 ; the
667 missing cell is the one in which they vary independently, and no reweighting of the occupied cells
668 recovers it. The emptiness follows from the gate rather than from $\mathcal{D}$. $\square$

669 *Proof of Lemma 2.* C1 at a^* gives $\min_{d \in \{d_1, d_2\}} \text{ASR}(d, a^*) < 0.5$. By C0 the minimum is not
670 attained at d_1 , so it is attained at d_2 . $\square$

671 *Proof of Proposition 5.* Emergence is the negation of the gate, so each member of $\mathcal{E}$ is rejected while
672 belonging to $\mathcal{S}$: a false negative. The bound is the resulting recall. $\square$

673 *Proof of Proposition 2.* Both halves are exhibited. (a) Mode S into Krum under model-scaling, $\kappa=0$
674 against $\kappa=2$: ρ moves from 1 to 54.6 and Krum's decision changes in up to 0.80 of rounds, while
675 the support of the admitted adversarial mass changes in none of the 60 measured rounds (a floor
676 on this arm: no adversary is admitted at any rung), the adversarial coefficient share is constant at
677 0.266667, and suppression shows **no evidence of a practically meaningful change under the pre-**
678 **specified ± 0.15 equivalence margin** ($|\Delta| = 0.026$). (b) `cos_krum` under the pixel backdoor: by
679 Proposition 1(a) its statistic is invariant, and its decision and admission are measured bit-identical
680 (0.000) at every rung, yet ASR falls 0.173 by the first rung (§5). $\square$

681 *Proof of Proposition 3.* By construction: Mode S pins every adversarial coefficient at $c=1$ and
682 spreads only benign clients, giving $\Delta_c \equiv 0$ at every ρ ; Mode A holds the benign weights uni-
683 form and sweeps only the adversary-to-benign ratio, moving Δ_c at fixed dispersion. Realized values
684 are tabulated in Appendix R, computed with no ASR involved. $\square$

685 *Proof of Proposition 1.* (a) $c_i u_i / \|c_i u_i\| = u_i / \|u_i\|$ for $c_i > 0$, so the matrix of pairwise cosines is
686 unchanged; FoolsGold's weights are a function of that matrix alone (via each client's normalized
687 maximum similarity), hence unchanged. (c) With $u_1=(1, 0)$, $u_2=(0, 1)$, $u_3=(3, 3)$ the coordinate-
688 wise median is $(1, 1)$ and u_3 is farthest; under $c = (1, 1, 0.1)$ the median becomes $(0.3, 0.3)$ and u_3
689 becomes the *closest* point, reversing the ordering Reputation depends on. $\square$

690 *Proof of Theorem 1.* (1) CoordMedian operates on $\{w_i u_{i,k}\}_{i=1}^K$ per coordinate. If the adversar-
691 ial values lie strictly beyond all benign values on one side (ii) and $n_a < K/2$, then at most
692 $n_a < \lceil K/2 \rceil$ of the extreme order statistics are adversarial, so the median is a benign value. For

the magnitude condition: $w_a > 0$ gives $\rho \geq \max_b w_b/w_a$, so $|u_{a,k}| > \rho \max_b |u_{b,k}|$ implies $|w_a u_{a,k}| > w_a \rho \max_b |u_{b,k}| \geq (\max_b w_b) \max_b |u_{b,k}| \geq \max_b |w_b u_{b,k}|$; with a common sign this places $w_a u_{a,k}$ strictly beyond every benign value in that direction. Model-scaling sets $u_a = S\tilde{u}_a$ with $S=10$, which amplifies the adversary’s own direction and so satisfies the sign assumption by construction; with measured $\rho \leq 2.1$ the margin is $S/\rho \geq 4.8$. TrimmedMean additionally needs the trimming budget to cover the adversarial count.

(2) Every benign residual satisfies $\|u'_b - \mu'\| \leq R'_B + \delta'$. For adversarial u'_a : $\|u'_a - \mu'_B\| \geq \Delta'_{\text{sep}} - R'_B$, so $\|u'_a - \mu'\| \geq \Delta'_{\text{sep}} - R'_B - \delta'$. Requiring this to exceed $R'_B + \delta'$ gives (i). $\square$

Theorem 1(2) orders residuals *per client*, which alone does not bound the adversary’s *total* share of the aggregate: several adversaries could each be individually down-weighted yet collectively hold substantial mass. The separation margin controls this, quantitatively:

Corollary 1 (Bounded aggregate adversarial influence). *Let $m := \Delta'_{\text{sep}}/[2(R'_B + \delta')]$ be the realized separation margin, so Theorem 1(2) reads $m > 1$. Benign residuals are at most $D_b := R'_B + \delta'$ and adversarial residuals at least $(2m - 1)D_b$, so bounding $\sum_j 1/d_j$ below by the benign terms alone gives*

$$\Lambda_a = \sum_a \lambda_a \leq \frac{n_a}{n_b(2m-1) + n_a}.$$

The aggregate is thus a convex combination whose adversarial share is at most Λ_a : writing $\mu' = (1 - \Lambda_a)\bar{u}_B + \Lambda_a\bar{u}_A$ for the benign and adversarial parts, the displacement from the benign-only combination $\bar{u}_B$ is at most $\Lambda_a\|\bar{u}_A - \bar{u}_B\|$.

The bound is only as good as the margin: as $m \rightarrow 1$ it degrades to $n_a/(n_b + n_a)$, approaching 40% of the mass at $n_a=2, n_b=3$. We measure $m \geq 2.82$ in every round with $w_a > 0$ (mean 4.71 over those rounds), capping the adversarial share at 12.6% for $n_a=2$; the realized Λ_a never exceeds 1.8%, and the bound holds in all 17 measured rounds.

J Body material, at full length

The body compresses the following; nothing here is a new claim, a new number or a new run. Figure 2 summarizes the preservation chain and the two directions in which it breaks.

J.1 Contributions, organization and scope in full

§1 states these compressed. Nothing here is a new claim.

J.1.1 The three evidence tiers

Table 2: The three tiers of evidence, of very unequal weight and never pooled. §1 states them in prose; every empirical claim in the paper is tagged with one of the three words where it is made.

tier	design	what it licenses
Causal	within-defense intervention at fixed d_2 , attack and C1 input: Mode S, Mode A (§4)	statements about the <i>mechanism</i> , scoped to these cells
Replication	EMNIST-byclass, second dataset <i>and</i> architecture ($\Delta=-0.017$), and <code>coord_median</code> , a second invariance class (+0.098)	“replicated on a second dataset and architecture”, never “generalizes”
Screening	the outcome-gated 42-pair suite, 38/42	evaluation prioritization only ; by Proposition 4 it cannot identify the mechanism’s contribution at <i>any</i> sample size

J.1.2 The five contributions

1. **The preservation hierarchy, and the two-directional break above.** Because the per-client defenses in use are all positive scalar rescalings, mechanism preservation is a decidable *invariance* question (Proposition 1), and on the weighting→aggregator class we prove it with explicit

725 margins (Theorem 1), a **boundary** theorem, not a security theorem. Definition 1 gives the five
 726 levels the box states, and §4 measures both breaks. One corollary is worth stating alone: weight
 727 heterogeneity is safe at *both* extremes and dangerous in the middle, so **the weight ratio ρ alone**
 728 **mispredicts** (Lemma 1, §R.1).

729 2. **An identification boundary for outcome-gated screening, and the design class it leaves**
 730 **admissible.** Proposition 4 proves the general claim above; its FL instance is that $C0 \wedge C1$
 731 force the downstream defense of every certified pair to suppress the attack alone, **mechanism-**
 732 **effectiveness confounding** (§3), and its evidential cost is real: a matched ablation built to break
 733 the confound had its premise refuted by its own pre-freeze baselines, so **the confound re-**
 734 **materialized inside the experiment designed to remove it.** A second boundary limits coverage
 735 rather than validation: **emergent** suppression (present in the pair, absent in both constituents)
 736 cannot satisfy an inherited-suppression gate, capping such a screen’s recall (Proposition 5), and
 737 our own best composition is its witness.

738 3. **Three central empirical findings, all causal interventions, not a validation of any screen.** (i)
 739 *The negative, instrumented and replicated (Causal, then Replication).* Mode S pins every ad-
 740 versary’s coefficient, closing Lemma 1’s attenuation channel *by construction*; a **pre-registered**
 741 **replication** in a *different* invariance class and on a *second dataset and architecture* reproduces it
 742 while **refuting the cross-arm admission ordering we had predicted** ($R^2=0.239$ over 26 cells),
 743 so **admission is a direction to measure, not a quantitative predictor.** (ii) *A sign reversal under*
 744 *a change of intervention (Causal).* The same cell yields -0.272 ($p_{\downarrow}=0.0015$) with adversarial
 745 influence free to co-vary and $+0.098$ with it pinned: the confounded design does not lose power,
 746 **it gets the sign wrong**; hence we withdraw the pooled form of C2’s causal claim (80 runs;
 747 $\rho_s = +0.351$, $p = 0.091$; §5). (iii) *Heterogeneity closes the theorem’s own margin.* At Dirichlet
 748 $\alpha=0.1$ RFA’s all-round mean separation margin collapses from $m=3.24$ to $m=1.01$, the limit
 749 at which Corollary 1’s bound goes vacuous, and two of three certified compositions cross the
 750 threshold (§3). A fourth predictor falls with them: **adversary-correlated dispersion Δ_c** is in-
 751 dependently controllable from ρ (Proposition 3) and locates the two real anchors our synthetic
 752 curve misses, but **orders suppression no better.** Statistic disturbance, admission magnitude and
 753 coefficient targeting are each refuted *as a predictor*: what survives is a *preservation level*, not a
 754 statistic of the coefficient vector.

755 4. **The screening rule is a consequence of (2), not a claim of its own** (*Screening* throughout).
 756 What survives the negative is a conservative *evaluation-prioritization* rule: read C2 off the down-
 757 stream signal’s invariance class, then evaluate the survivors first. We report it as an application
 758 (§3), never as a validated security predictor; it certifies only *inherited* robustness under a *com-*
 759 *mitted* attack, at recall 40%, and outputs an evaluation *ordering*, never a deployment decision.

760 5. **Secondary results.** Fixed-order composition beats randomized ordering $9\times$, because preser-
 761 vation conditions are order-sensitive; and **one** whitebox composition-aware *adaptive* stress test
 762 raises $FG \rightarrow RFA$ to 0.157 ± 0.073 at $n=30$, $2.4\times$ its seed-matched 0.064 on the pixel arm it is
 763 built from, **an inherited-robustness guarantee is not adaptive robustness** (§6, Appendix O).

764 J.1.3 How the paper is organized, as an argument

765 Composition creates an **identification problem**. §2 defines the object to be identified: mechanism
 766 preservation as invariance, the five levels, the conditions; §3 characterizes the **boundary** on iden-
 767 tifying it and *derives* the class of designs that can; §4 runs the experiment that class dictates; §5
 768 reports what the boundary **falsifies**, our own pooled hypothesis included; §3 states what survives.

769 J.1.4 Why this class, and the scope discipline

770 Everything we *prove* is instantiated for positive per-client rescalings, deliberately: the class is a
 771 **laboratory for studying composition, not a general model of it**, it covers the per-client defenses in
 772 current use (norm clipping, FoolsGold, RFA reweighting, FLTrust) and makes invariance decidable
 773 from definitions. **Every claim here is general, formal or empirical, and they do not transfer**
 774 **upward.** Table 1, first in the appendix, is the one place each scope condition is stated (what falls
 775 outside the class, what the theorems do and do not bound, the sample sizes, what “refuted” means,
 776 why our instruments are not deployable) and we do not restate them at each use.

777 J.1.5 Threat model and scope

778 We consider persistent backdoor attacks where an adversary controlling fraction f of clients embeds
779 triggers that survive across rounds, and define *persistence* γ as the fraction of backdoor ASR retained
780 per round of clean aggregation ($\gamma \approx 0.66$ for CIFAR-10/CNN). Attacks include model-scaling [Bag-
781 dasaryan et al., 2020], pixel-pattern backdoors [Chen et al., 2017], distributed backdoor attacks [Xie
782 et al., 2020], and adaptive composition-aware adversaries [Shejwalkar and Houmansadr, 2021]. The
783 framework constrains the aggregation pipeline, not the communication graph, and is evaluated only
784 in **star-topology FL**. **A scope condition on the criterion itself, not a limitation of our runs:** C1
785 and every prediction built on it are defined against a *committed* adversary, so what is bounded is
786 *inherited* robustness under a fixed attack, **never adaptive robustness** (Table 1).

787 J.2 The prioritization criterion, and where it breaks

788 **This section is a consequence of the boundary, and what survives is an *evaluation-prioritization***
789 ***criterion*, not a security screen.** C1 $\wedge$ C2 answers one question: given a menu, which ordered com-
790 positions are worth running first. **We disclaim predictive validity explicitly;** it does not predict
791 whether a composition is secure, its positive direction is unvalidated, and the cross-arm admission
792 ordering a predictive reading would need was *refuted* by our own pre-registered replication (§4). It
793 also costs us our own best composition: under the per-attack quantifier **FG $\rightarrow$ RFA no longer sat-**
794 **isfies C1**, so the lowest-ASR composition we found is a **false negative** and Proposition 5’s witness
795 (uncertifiable by *any* inherited-suppression screen rather than missed by ours) while the certified
796 in-distribution set (FG $\rightarrow$ CM 0.107, Rep $\rightarrow$ CM 0.463 at $n=30$) is correct and the miss costs noth-
797 ing detectable (0.093 against 0.107, Welch $p=0.52$). Out of sample it licenses *positive* predictions
798 only, against uninformative FAIL controls, so **this is not an out-of-sample classification result:** the
799 *Screening*-tier 38/42 at recall 2/5 rests on the in-distribution tiers. The defensible claim is *conser-*
800 *vative prioritization*; scoring condition by condition, the confusion matrix, the false negatives and
801 what the rule can and cannot do are in Appendices S–T, U and Table 9 (App. V).

802 **The procedure, and where it breaks.** (0) If one defense already suppresses every attack you
803 have, stop (C0); (1) require a constituent that suppresses the committed attack (C1); (2) read C2 off
804 the downstream signal’s *invariance class*, which §3 says you cannot test directly; (3) read C3 as a
805 margin from a short pilot, never as a gate. Then **evaluate the survivors first, and evaluate them:**
806 none of this certifies security. **Where the rule breaks is measured rather than argued: under**
807 **extreme heterogeneity (Dirichlet $\alpha=0.1$) two of three certified compositions cross the threshold**
808 **and RFA’s all-round mean separation margin collapses from $m=3.24$ to $m=1.01$:** the $m \rightarrow 1$
809 limit at which Corollary 1’s bound says nothing. Theorem 1’s guarantee is therefore not merely
810 conservative under heterogeneity but *empty* there, and the theorem earns its keep by predicting its
811 own failure region: a short pilot measuring $m \approx 1$ declines to deploy. **Post-hoc, not pre-registered**
812 (Appendix O.1).

813 J.3 What is proved, what is supported, and the six-step procedure

814 **What is proved, what is supported, and what is not. Proved:** the invariance classes (Prop. 1);
815 sufficient conditions for mechanism preservation on the weighting $\rightarrow$ aggregator class (Thm. 1,
816 Lem. 1); the identification boundary and the design class it leaves admissible (Prop. 4, Lem. 2);
817 the recall bound on inherited-suppression screens (Prop. 5). **Empirically supported** (*Causal*, then
818 *Replication*): that substantial disturbance of the downstream statistic need not produce a corre-
819 sponding change in suppression when the adversarial coefficient share is held fixed, on the Krum
820 arm under Mode S and, by pre-registered replication, in a *different* invariance class and on a *second*
821 *dataset and architecture*; and that the criterion behaves as a conservative prioritizer on 42 pairs. **Not**
822 **established:** its positive predictive direction ($R^2=0.239$, cross-arm ordering refuted), and general-
823 ization beyond the scope Table 1 fixes, $\gamma \approx 0.66$ being architecture-conditional. The criterion **trades**
824 **discovery for reliability:** it rejects working compositions (recall 40%, one of them ours), cannot in
825 principle do otherwise (Prop. 5), and draws negative-direction evidence only from the in-distribution
826 tiers. FG $\rightarrow$ CM is an evaluation *ordering*, **not a deployment recommendation** (Appendix H).

827 **What to do instead: six steps for any composition claim.** (1) Name the downstream mechanism
828 and its invariance class. (2) State *which* preservation level you claim, (P1)–(P5). (3) Check the

upstream transform lies in the class your argument covers: C0 here; one altering participation, carrying cross-round state or mixing clients is covered by no invariance argument. (4) Hold d_2 , the attack, and therefore the gate’s input fixed; **do not select candidates on the outcome**. (5) Intervene on the upstream transform and **measure the channel**: aggregate displacement, decision, admission, influence, not the outcome-conditional contrast. (6) Report the level at which the chain broke. Steps (4)–(5) are what Proposition 4 licenses, and following them turned our -0.272 into $+0.098$; **this, not the criterion, is what we ask readers to carry away**. One boundary it does not cross: under one whitebox composition-aware *adaptive* stress test $FG \rightarrow RFA$ rises to $ASR = 0.157 \pm 0.073$ at $n=30$, $2.4\times$ its seed-matched 0.064 on the pixel arm it is built from, so **an inherited-robustness guarantee is not adaptive robustness**: the condition the framework most lacks (Appendix O).

J.4 The identification boundary, formally

This section states §3’s two propositions and their FL instance at full length; the body gives the boundary in words and the design it dictates.

Proposition 4 (Why constituent-effect contrasts cannot identify composition effects under inherited-effect gating). *Let a screen certify $(d_1 \rightarrow d_2)$ against attack a only if both (i) an **eligibility gate** $E(d_1, d_2, a) := \min_{d \in \{d_1, d_2\}} M(d, a) < \theta$ holds, where M and θ are the same outcome metric and threshold the screen predicts for the composition, and (ii) a **mechanism condition** $P(d_1, d_2)$ holds, determined by the two defenses’ definitions alone. Suppose the screen is applied only where the downstream mechanism is load-bearing, $M(d_1, a) \geq \theta$. Consider estimating P ’s contribution to $M(d_1 \rightarrow d_2, a)$ from contrasts across a defense menu $\mathcal{D}$. Then every certified pair satisfies $M(d_2, a) < \theta$, so of the four cells of $\{P, \neg P\} \times \{M(d_2, a) < \theta, M(d_2, a) \geq \theta\}$ the certified set occupies at most two, and **the cell needed to separate P from d_2 ’s standalone effectiveness is empty by the gate, for every $\mathcal{D}$. No contrast varying d_2 identifies P ’s contribution. Identification requires holding d_2 , a and therefore E fixed while intervening on the upstream transform: **the admissible designs are within-defense**.***

The structure is a collider: the gate E is a function of constituent effectiveness, effectiveness affects the composition’s outcome, and the mechanism condition P is fixed by the same pair of definitions, so conditioning on E makes P and effectiveness co-vary across every certified candidate. The design table is the whole argument, and **the missing counterfactual is a cell, not a sample size**:

	$M(d_2, a) < \theta$ (d_2 suppresses a alone)	$M(d_2, a) \geq \theta$ (it does not)
P holds (mechanism preserved)	observed: every certified pair	empty by the gate
$\neg P$ (mechanism destroyed)	observed: predicted-FAIL pairs	empty by the gate

Reading down a column contrasts P against $\neg P$ at fixed effectiveness (the comparison that would identify P ’s contribution) and the right-hand column, the only one in which d_2 is not already doing the work, is empty for every menu $\mathcal{D}$. **The counterfactual needed is excluded by design, not undersampled**. The emptiness follows from the gate, not the menu, and nothing in the argument is specific to FL, to backdoors, or to our conditions (Appendix I). Its instance for *our* criterion is one line:

Lemma 2 (C0 and C1 place the effective constituent downstream). *Let $(d_1 \rightarrow d_2)$ satisfy C0 and C1, and let a^* be a committed attack witnessing C0, so $ASR(d_1, a^*) \geq 0.5$. Then $ASR(d_2, a^*) < 0.5$: the downstream defense suppresses a^* on its own.*

Instantiated: in every in-scope pair our conditions certify, d_2 already suppresses the attack d_1 fails, so a certified composition’s low ASR is always attributable to d_2 ’s standalone strength (C1) rather than to d_1 having preserved d_2 ’s mechanism (C2, C3): confounded **by the conditions themselves**, not by an unlucky menu. We name this **mechanism–effectiveness confounding**.

This is selection on the outcome; where it hides is the point. The hazard is standard (conditioning on a descendant of both the putative cause and the outcome [Elwert and Winship, 2014, Pearl, 2009]) and **we claim no novelty for it**. What is not standard is that in composition studies the gate is presented not as conditioning but as a *reasonableness* filter (“we study defenses that work”), so it reads as hygiene rather than as the step that removes identification. Two consequences are routinely assumed away: **a better estimator cannot fix it**, no reweighting of occupied cells recovering an unoccupied one, and **the gate cannot be dropped**, because without it C1 is vacuous.

879 **Both worked instances are ours, and one of them inverted a sign.** We name no other paper;
880 our own two designs are the instances. The C1-gated study of §5 is the template at the level of the
881 *criterion*; the dose ladder of §5 is a second instance one level down, its gated quantity being the
882 dose’s incidental attenuation of whichever client carries the poison. **We ran the outcome-gated**
883 **design first, and on the identical cell it returned the opposite sign** (Table 3): the hazard cost us
884 the sign of our own effect, in the direction that flattered us.

885 **The admissible design, derived: this, not the criterion, is the practical contribution.** Read
886 forwards the boundary is *constructive*: each clause below is forced by a term in Proposition 4, not
887 chosen for convenience. **(1) Hold d_2 and the attack fixed**, so the gate’s input $M(d_2, a)$ is numeri-
888 cally the same quantity on both sides of every contrast and cannot co-vary with the mechanism. **(2)**
889 **Intervene on the upstream transform** instead of selecting a different d_2 off a menu, because it is
890 varying d_2 that empties the identifying cell. **(3) Pin the adversarial coefficient channel**, because
891 an upstream that attenuates the adversary moves the outcome by a route with nothing to do with d_2 ’s
892 statistic: the confound that inverted our own sign above. **(4) Measure every channel: aggregate**
893 **displacement, decision, admission, influence, rather than the outcome-conditional contrast**, the
894 outcome being exactly the variable the gate conditioned on. §4 is clause (3) built as an instrument;
895 Appendix J.3 states all four as a six-step procedure for any composition claim.

896 **A second boundary: what such a screen cannot find.** The gate that limits what a screen can
897 be *validated* on also limits what it can *cover*, which converts the most obvious objection to our
898 criterion, that its recall is poor, into a scope result.

899 **Proposition 5** (Inherited-suppression screens cannot certify emergent synergy). *Call $(d_1 \rightarrow d_2)$ ’s*
900 *suppression of a **emergent** if $M(d_1 \rightarrow d_2, a) < \theta$ while $\min_{d \in \{d_1, d_2\}} M(d, a) \geq \theta$. No emergent*
901 *composition satisfies the eligibility gate of Proposition 4. Hence, writing S for the compositions*
902 *with $M < \theta$ and $\mathcal{E} \subseteq S$ for the emergent ones, any such screen’s recall obeys $\text{recall} \leq 1 - |\mathcal{E}|/|S|$*
903 *on every menu and every attack.*

904 **This is a scope theorem, not a defect found by tuning**, and it is specific neither to C1 nor to
905 FL: any screen requiring a constituent to suppress the attack independently inherits it. Our own
906 lowest-ASR composition is the witness (FG→RFA reaches 0.093 where FoolsGold alone reaches
907 0.732 and RFA alone 0.845), so at $|\mathcal{E}|/|S| = 1/5$ the bound caps any such screen’s recall at 80%,
908 above the 40% we measure (§3). Screening for inherited robustness and searching for synergistic
909 compositions are different problems; only the first is ours.

910 **What the boundary rules out, our own previously-decisive test included.** These proofs are
911 short and their value is what they rule in and out. They rule *out* the experiment we had previously
912 named as decisive: a d_2 whose signal is invariant under positive rescaling but *weak in isolation*
913 cannot satisfy C1 in scope, so no such certified pair exists to test. §4–§5 run four members of
914 the class they rule in. **The strongest disturbs the statistic while pinning the adversarial mass,**
915 **isolating statistic disturbance from adversarial coefficient attenuation, so we built it (§4).**

916 J.5 The targeted suite in full

917 The body compresses these paragraphs; nothing here is a new claim or a new run.

918 **Reputation’s confirmation, in full.** With the adversary’s share pinned, reputation’s rise at $\rho=54.6$
919 shrinks from +0.762 to +0.178 (one-sided Welch $t=-5.22$, $p=0.0017$); what remains tracks a
920 small *admission* change (0.000 → 0.020) rather than the large decision change (0.000 → 1.000).

921 **The arc that produced this section.** Our pre-registered pooled hypothesis (H1) (ASR should rise
922 with measured disturbance of the statistic the downstream defense reads) was **rejected**, the effect
923 coming out *negative* (-0.272 , $p_{\downarrow}=0.0015$ on the flagship cell), because that ladder moved adver-
924 sarial influence and the statistic *together* and so never tested H1 at all (§5). This section builds the
925 instrument that fixes it, pinning the adversarial coefficient share while the statistic moves, so statistic
926 disturbance is isolated from adversarial coefficient attenuation. C2 is written at level (P2); if losing
927 (P3) meant losing (P5) the criterion would transfer straight to security, and the experiment below
928 adjudicates that implication by **oracle intervention on two arms in two invariance classes** (*Causal*
929 tier), not by a law about defense compositions in general.

930 **An instrument that closes the attenuation channel by construction.** A positive rescaling
 931 reaches d_2 through two channels at once: it disturbs the *statistic*, and it *attenuates* whichever client
 932 carries the poison (Lemma 1). Separating them needs an upstream that moves the first while pinning
 933 the second. **Mode S** does this: it fixes every adversary’s coefficient at $c=1.0$ exactly and spreads
 934 only the benign clients over $\rho = e^{2\kappa}$, so the **adversarial coefficient share is constant at 0.266667**
 935 **across all four rungs** (spread 3.5×10^{-7}), and **a rise in ASR here cannot be attenuation**. Design
 936 and thresholds were frozen before any run, the share verified with no ASR involved (Appendix X).
 937 **Mode A** sweeps the payload instead (§5). **Both are oracle instruments, not defenses and not**
 938 **models:** they read adversary identity, which no deployable defense can, and they do **not** quantita-
 939 tively reproduce any real defense’s behaviour, at matched disturbance two real upstream anchors sit
 940 $+0.767$ and $+0.789$ above the synthetic curve (§5), the measured size of that gap. Their purpose is
 941 causal identification, not defense evaluation.

942 **Two readings of C2, frozen against each other.** *H-statistic:* what matters is disturbance of d_2 ’s
 943 statistic, read off its *decisions*. *H-admission:* what matters is the *adversarial input* d_2 *admits*. On
 944 the Krum arm these predict opposite outcomes, since Mode S disturbs Krum’s selection heavily
 945 while, the adversary’s weight being pinned, leaving the adversarial mass it admits untouched. That
 946 arm adjudicates (Table 4, App. N): as ρ climbs to 54.6, Krum’s **decision changes in up to 0.80**
 947 **of rounds** while its **admission changes in none of the 60 measured rounds**, and ASR shows
 948 **no evidence of the predicted increase, remaining inside the pre-specified equivalence margin**
 949 ($p_{\uparrow}=0.70$; $|\Delta|=0.026 < 0.15$). **H-statistic is refuted:** on this arm **substantial statistic and deci-**
 950 **stion disturbance does not produce a corresponding change in suppression once the adversarial**
 951 **coefficient share is held fixed, and the evidence points to admission as the relevant channel in**
 952 **these experiments.**

953 **What the instrument does *not* hold fixed, named in full, and what it pins exactly.** Mode S pins
 954 the *adversarial* coefficients; it does not pin the benign ones and was never claimed to. So the dose
 955 also moves (a) Krum’s pairwise distances, hence the statistic itself; (b) *which* benign client Krum
 956 selects, in $0.533/0.800/0.733$ of rounds at $\rho=2.72/7.39/54.6$; (c) the **magnitude and geometry of**
 957 **the update actually aggregated**, since the selected client enters scaled by its own coefficient; and
 958 (d) the model trajectory, which diverges across rungs in consequence. Channel (c) is **already mea-**
 959 **sured and printed:** the Δ agg. column of Appendix C is $\|\text{agg}(T(U)) - \text{agg}(U)\|/\|\text{agg}(U)\|$, the
 960 relative displacement of the aggregate the defense *emits*, and reads **0.892** here. Appendix D splits
 961 that 0.892 into 0.986 of *re-selection* and 0.973 of *rescaling*, with none of the re-selections landing
 962 on an adversary. So the claim is sharper: **that displacement, with decision changes in up to 80%**
 963 **of rounds, leaves suppression inside the pre-specified equivalence margin, provided the adver-**
 964 **sarial coefficient share does not move. What is pinned is pinned exactly:** $c_{\text{adv}}=1.0$ at every rung
 965 with maximum deviation $0.00\text{e}+00$ and the share constant to $3.5\text{e}-07$ on CIFAR-10, $3.9\text{e}-07$ on
 966 EMNIST-byclass, float32 read-back noise, not slack in the instrument (`measure_admission.py`).

967 **Closing the magnitude channel as well: the pre-registered score-only control.** Channel (c)
 968 is the one confound Mode S leaves open, so we closed it too. **Score-only Krum** scores on the
 969 transformed stack exactly as before and aggregates the *untransformed* selected update, so the dose
 970 changes *which* client contributes and nothing about *what* it contributes: Δ dec. is the same 0.733
 971 while Δ agg. is no longer 0.892 but the re-selection component alone, 0.986 relative to $\|u_{s_0}\|$.
 972 Rules were frozen at 35788d9 before any score-only ASR existed: primary Δ against the same
 973 ± 0.15 margin, where $\Delta > +0.15$ would have said **the published result was partly an artifact of**
 974 **the magnitude channel** and narrowed the central claim. **The outcome is $\Delta = -0.023$ (mean ASR**
 975 **$0.062 \rightarrow 0.039$) against -0.026 uncontrolled on the identical cell, so closing the magnitude chan-**
 976 **nel moves the effect by 0.003;** the refuting threshold 0.212 is nowhere approached and every rung
 977 clears the accuracy floor. **The central claim is thus instrumented against two independent con-**
 978 **found, attenuation and magnitude (row six of Appendix C).** The secondary trend test, the headroom
 979 asymmetry, the selectors-only scope and the `cos_krum` corollary are in Appendix E.

980 **A pre-registered replication, on a different defense class, that refutes the positive half.** The
 981 two existing Mode-S arms move together on both predictors, so neither adjudicates. `coord_median`
 982 under the pixel backdoor is *rank-discordant*: of the interpretable arms it has the **smallest decision**
 983 **change (0.482) and the largest admission-level change ($\Delta\Lambda_a=0.033$),** so the two readings predict
 984 opposite rank positions for its rise. Thresholds were frozen against the two published rises: confirm

if $\Delta > +0.178$, refute if $\Delta < +0.150$ (Appendix X). **The outcome is $\Delta = +0.098$, so the admission ordering is refuted:** the arm with the largest admission-level change rose *less* than one with a smaller change. Nor does the decision reading predict it, having called the *smallest* rise when $+0.098$ is the middle one; the added cells leave the pooled model unconfirmed ($R^2=0.239$, the share coefficient’s sign failing the pre-specified conjunction). **What replicates is the negative**, on a coordinate-wise order statistic rather than Krum’s selector: a decision change of 0.482 costs no suppression by the frozen margin, at accuracy 0.698. Two qualifications: this cell’s identity ASR is 0.443, so equivalence is to the *identity rung*, not a low absolute ASR; and the rise is positive in 5/5 seeds (post-hoc), so the cost is **small and consistent, not zero**. **The instrument earns its keep:** the confounded ladder on this identical cell *falls* 0.272 ($p_{\downarrow}=0.0015$), so closing the attenuation channel **inverts the sign of the effect**.

What the two arms jointly establish. Both halves are witnessed by runs already reported, so what follows formalizes measurements rather than adding a claim. **Statistic-level non-identifiability** (Proposition 2, proved from the exhibited witnesses in Appendix I): within the positive per-client rescaling class there are transform families inducing arbitrarily different statistics and *decisions* at the *same* admitted adversarial mass and the same suppression (Mode S into Krum under model-scaling) and families inducing the same statistic, decision *and* admitted mass at *different* suppression: `cos_krum`. Consequently **no predicate of the downstream statistic alone determines the admitted adversarial mass, and no predicate of the admitted mass alone determines suppression**. **That corollary is the paper’s title claim:** no metric defined *solely* on preservation of the downstream statistic can identify attack-suppression equivalence. **Scope:** a non-existence claim *within this class on these defenses and attacks*, by exhibited counterexample: **not** that no metric of any kind can order security. A second proposition sharpens §2’s warning: statistic-level dispersion ρ and adversary-correlated dispersion Δ_c are **independently controllable** ($\Delta_c \equiv 0$ while ρ ranges over $[1, 54.6]$; Δ_c over $[-1.138, +2.469]$ at fixed dispersion), so **a transform reported by ρ alone leaves Δ_c entirely free** (Proposition 3), which does *not* make Δ_c a predictor, its magnitude failing to order suppression either (§5). The instrument is elaborate because §3 leaves nothing simpler. §5 answers what remains: *what happened when we tried it the obvious way first?*

1013 K What the boundary licenses: the suites in full

1014 This section holds §5’s three suites at full length: the two attempts the identification boundary
1015 defeated, the confounded dose ladder that motivated Mode S, and the remaining targeted arms.
1016 Nothing here is rescored; the body states the outcomes.

1017 K.1 Two attempts that the confound defeated

1018 Both were built to validate the *screen*; neither could, so they are evidence for Proposition 4 rather
1019 than for the framework.

1020 **Two pre-registered suites, and the two ways the boundary asserted itself.** (i) *A prospective*
1021 *suite on three unused defenses: 13/20, and no positive test.* Everything scored in §3 was developed
1022 on the menu it is scored on, so we ran a prospective test on FLTrust [Cao et al., 2021], Krum and
1023 Multi-Krum [Blanchard et al., 2017], for which the theory predicts *opposite* outcomes. All 20 labels
1024 were git-committed before the output directory existed and both attacks clear the FedAvg sanity gate,
1025 so the negatives are informative (11/13 in scope, 5/7 on the falsifiable arm; Appendix Q). **Three**
1026 **disclosures dominate it:** FLTrust alone suppresses both attacks (0.048), so its 7 pairs are **out of**
1027 **scope** by C0 (applied *retrospectively*); its 6 in-scope predicted-LOW pairs are **near-identities**, four
1028 bit-identical to FLTrust alone, so **this suite supplies no falsifiable prospective positive test** (one
1029 arrived only with the census of §K.2, and it is a single cell); and MultiKrum at $K=5$ selects all five
1030 participants and so is bit-identical to FedAvg, which together with the pass-through d_1 rows leaves
1031 8 of the 20 rows reducing to a single defense and the falsifiable arm with 2 pairs composing two
1032 mechanisms rather than 7. (ii) *A matched metric-swap ablation, whose premise its own baselines*
1033 *refuted.* To break the confound by construction we held the suppression *principle* fixed and varied
1034 only the invariance class, mirroring Reputation and Krum scoring from *directions* alone. Measuring
1035 the C1 input before freezing **refuted the design’s premise:** `cos_reputation` collapses (0.017 $\rightarrow$
1036 0.982) and `cos_krum` reaches 0.078 only at 0.146 accuracy, failing the C1 gate applied everywhere

1037 else (**magnitude information is what lets this family handle a magnitude attack**), so C1 bound
 1038 for all eight pairs, every label was HIGH, and the suite bears on specificity only (8/8; App. L, P).
 1039 **The confound re-materialized inside the experiment built to remove it.**

1040 **The one contrast that escapes, and the design it dictates.** One comparison evades the boundary:
 1041 hold d_2 and the attack fixed and vary only the upstream, so C1 is identical by construction. On
 1042 *Krum*/model-scaling, NormClip disturbs its selection 0/9 and reaches 0.062 while RFA disturbs it
 1043 6/9 and reaches 0.779 (Welch $p=0.012$, $n=5$), with no overlap across the six cells where standalone
 1044 d_2 genuinely suppresses. **This is post-hoc** (Appendix L) but has what nothing above it does: C1
 1045 held exactly fixed, disturbance measured on the mechanism rather than read off the outcome.

1046 K.2 A third attempt that closed a census, pre-registered

1047 **The region where mechanism preservation is settled on both sides is small enough to enu-**
 1048 **merate, so we censused it rather than sampling it.** It requires a d_1 emitting strictly posi-
 1049 tive per-client coefficients $T(u_i) = c_i u_i$, so Proposition 1(a) applies, against a d_2 reading an
 1050 exactly scale-invariant statistic: over this project’s defenses and the metric-swap mirrors that is
 1051 $\{\text{norm_clip}, \text{rfa}, \text{reputation}\} \times \{\text{foolsgold}, \text{cos_krum}, \text{cos_reputation}\}$, exactly 9 ordered
 1052 pairs, 7 already measured at the identical protocol and 2 unspent. Enumeration is what re-scoped
 1053 a planned 20-pair wave: 118 **ordered pairs of this project are already spent**, and once a non-
 1054 degenerate d_1 and C0 are required only 6 candidates remain anywhere. Pre-registered at 96f8856
 1055 with per-pair C0–C3 verdicts, a frozen scoring rule and a two-sided resolution table; 24 runs, 4.87
 1056 h.

$d_1 \downarrow d_2 \rightarrow$	foolsgold	cos_krum	cos_reputation
norm_clip	0.887 HIGH	0.508 HIGH [†]	0.935 HIGH [†]
rfa	0.914 HIGH	0.566 HIGH	0.861 HIGH
reputation	0.807 HIGH	0.263 LOW*	0.787 HIGH*

Max-committed mean ASR over seeds 42, 43, 44; LOW iff < 0.5 .

* the two unspent cells, labels frozen before the output directory existed.

[†] norm_clip’s clip is inactive on the pixel arm (max norm $3.043 < \tau=5$), so those

pixel cells are bit-identical to d_2 alone; both verdicts rest on the scaling arm, where it is not.

1058 foolsgold is excluded from the d_1 column by measurement, and that measurement changed
 1059 the design before any prediction was frozen. It normalizes with `weights - weights.min()`,
 1060 forcing at least one coefficient to *exactly* zero every round, which puts it outside Proposition 1(a)’s
 1061 strict $c_i > 0$ hypothesis and under Lemma 1 instead: measured through the shipped transform, 2/5
 1062 exact zeros every round and both downstream cosine statistics moved in 9 of 9. Synthetically at
 1063 1000 trials the effect is entirely the zero’s: strictly positive coefficients change cos_krum’s selection
 1064 0/1000 times, one exact zero 471/1000. reputation emits 0/5 zeros at a positive-part ratio up
 1065 to 18839.8, and under that rescaling in the same rounds the cosine statistics moved 0/9 while
 1066 krum’s selection moved 8/9 and reputation’s own ordering 9/9: **a contrast between statistics**
 1067 **on identical inputs, not between conditions.** So foolsgold $\rightarrow$ cos_* were entered as *measured-*
 1068 *C2-fail controls*. **Had C2 been assumed from the definitions rather than measured, this wave**
 1069 **would have reported two mislabelled pairs.**

1070 4/4 as registered, including the pair flagged uncertain. Pair 1 (reputation $\rightarrow$ cos_krum, cer-
 1071 tified) LOW at 0.263; pair 2 (reputation $\rightarrow$ cos_reputation, the C1-fail $\wedge$ C2-true cell that dis-
 1072 criminates the strategies) HIGH at 0.787; the two controls HIGH. The near-identity check is clean
 1073 on all 8 arms (0 of 3 seeds identical to d_2 alone), so these are compositions. Pair 1 was the live
 1074 risk and the prereg said so: its LOW rests on two *different* attacks suppressed by two *different* con-
 1075 stituents, reputation holding model-scaling (0.017 at accuracy 0.777) and cos_krum the pixel
 1076 backdoor (0.300 at 0.598), while on the scaling arm the two nearest measured analogues (0.508,
 1077 0.566) pointed the other way. **So the census’s one low-ASR pair is suppressed by neither con-**
 1078 **stituent alone, and it is the pair C1 $\wedge$ C2 selects:** over the closed census statistic preservation admits
 1079 all 9 at 11% precision for 90 runs, C1 $\wedge$ C2 admits 1 at 100% precision and 100% recall for 10. In
 1080 this region C2 alone is not a screen at all, since everything here satisfies it by construction, and at the
 1081 frozen seed count B dominates A on precision, recall and cost at once: **the opposite of the trade in**
 1082 **§5.** Disclosure (v) below is the qualification that phrase needs.

1083 **Five disclosures.** (i) B’s 100% is the accuracy-gated reading: `cos_krum` alone reaches 0.078 on
1084 model-scaling at accuracy 0.146, barely above CIFAR-10’s 0.10 chance, so it is a destroyed model;
1085 read at the bare 0.5 threshold C1 also admits `norm_clip` $\rightarrow$ `cos_krum` and `rfa` $\rightarrow$ `cos_krum`, both
1086 HIGH, and B selects 3 for 33%. The analyzer emits both readings. (ii) The one positive is $n=3$
1087 and wide: pair 1’s pixel arm is 0.000/0.107/0.681 per seed, so one seed sits above the threshold
1088 and the verdict is the frozen mean rule; we add no seeds after seeing that. (iii) The controls are
1089 weaker than intended: every `foolsgold-as-d1` scaling arm sits at chance accuracy (0.100), and one
1090 control is HIGH only through the accuracy gate (raw ASR 0.086 at accuracy 0.255), so they show
1091 the C2 probe is not inert but are not independent confirmations. (iv) 7 of the 9 cells are retrospective;
1092 only the two `reputation` cells were blind, and the closure is prospective at the margin rather than
1093 throughout. The wave tripped two frozen sanity flags, both reported and neither adjusted: pair 2’s
1094 scaling arm at accuracy 0.790 against the FedAvg reference 0.667, and that same control’s scaling
1095 ASR identically 0.0 across all 3 seeds at 0.100 accuracy. (v) The C1 input that certifies pair 1 is the
1096 lowest-seed reading of a high-variance cell: `cos_krum` alone measures 0.300 on the pixel arm as the
1097 mean over the frozen seeds 42–44 (0.003/0.269/0.630), and the identical cell, same protocol and
1098 identity upstream, is measured twice more in this paper at 0.490 over 5 seeds and 0.584 over 8, the
1099 latter being the value §J.5 reports as `cos_krum` never suppressing. Six of the eight seeds are above
1100 0.5, and at $n=8$ C1 fails for pair 1, B certifies nothing here, and the census’s one low-ASR pair
1101 becomes another false negative; 0.490 is still below the threshold, so C1 survives at $n=5$. We report
1102 this instead of rescoring, because the pre-registration fixes the seed set at three and forbids adding
1103 seeds afterwards; what it costs is the *dominance* claim above, which holds at the frozen seeds and is
1104 not established beyond them.

1105 **K.3 Why the screen has no second dataset: C1’s pixel-arm input is absent on** 1106 **EMNIST-byclass**

1107 **A second-dataset composition wave was designed, costed at ~ 55 h, and then not run, be-**
1108 **cause the quantity it needed does not exist.** C1 is defined on *standalone* suppression, so its
1109 inputs are measured before any per-pair label is frozen: on CIFAR-10 a pair clears it because the
1110 two constituents cover *different* attacks, as in the census above where `reputation` holds model-
1111 scaling (0.017 at accuracy 0.777) and `cos_krum` the pixel backdoor (0.300 at 0.598). On EMNIST-
1112 byclass/simple_cnn the pixel side is missing. Re-measured through the identical runner at the iden-
1113 tical protocol ($N=10$, $K=5$, 50 rounds, $f=0.2$, Dirichlet 0.5, seeds 42/43/44), `coord_median`
1114 reaches **0.609** at clean accuracy 0.836 (per seed 0.717/0.493/0.618), and over the **seven** de-
1115 fenses with an EMNIST-byclass pixel cell that is the *minimum*: `krum` 0.660, then trimmed mean
1116 0.988, `norm_clip` 0.996, `fedavg` 0.997, `multi_krum` 0.997, `rfa` 0.998 (six *distinct* aggregators:
1117 `multi_krum` reduces to `fedavg` at $K=5$, which is why those two agree). **So no pair drawn from**
1118 **them satisfies C1 there:** the screen certifies nothing over them, every such pair is predicted HIGH,
1119 and the strategy comparison of §K.2 cannot be made on this dataset from these inputs.

1120 **The failure is the pixel arm’s, and we do not escape it by dropping the other.** On the model-
1121 scaling arm `krum` does suppress there (0.027 at clean accuracy 0.767), so a scaling-only protocol
1122 would supply a C1 input and manufacture the comparison; we decline it for the reason the prospec-
1123 tive suite declines the mirror move, that the committed protocol requires *each* attack and choosing
1124 the arm that turns out to have an input is attack-shopping regardless of the resulting labels.

1125 **Two things this is not.** It is *not a failed replication*: no prediction was made and refuted, and the
1126 question a screen answers does not arise where nothing suppresses the attack under test, so it is a
1127 scope limit on C1’s precondition. And it is *not unsatisfiability over the whole menu*: `foolsgold`,
1128 `reputation` and the cosine mirrors have no EMNIST-byclass pixel cell, so the coverage is ex-
1129 actly those cells and we state the gap instead of generalizing past it. We report the obstacle
1130 rather than lower f , subset the 62 classes or lengthen training until a positive appeared, each of
1131 which is a new design axis chosen after seeing the obstacle. Finally, the `coord_median` figure
1132 came back *bit-identical* to the corresponding cell of the earlier payoff matrix (0.6091787658... at
1133 0.8363207047...), which by the standard we apply to ourselves in §J.5 establishes protocol *identity*
1134 and not corroboration: the same computation run twice cannot corroborate itself. What it does buy is
1135 that the matrix’s other cells are protocol-comparable, so the six borrowed rows above did not cost an-
1136 other 21 h (`measure_femnist_c1_inputs.py`, `results/femnist_c1_inputs/summary.json`).

1137 K.4 H1, and the confounded design that motivated the instrument

1138 The prospective suite and the metric-swap ablation both tried to validate C2 by *choosing a different*
 1139 *downstream defense*: the move Proposition 4 rules out. The admissible alternative varies only the
 1140 *dose* at fixed d_2 and attack, so the C1 input is the same quantity at every rung and the confound is
 1141 **impossible by construction** rather than matched away.

1142 **The instrument, the abscissa, and the predictions frozen against them.** The synthetic d_1 has
 1143 one parameter, the *dispersion* of a positive per-client rescaling: $c_j \propto \exp(\kappa(2j/(K-1) - 1))$ at
 1144 position j of a (seed, round)-keyed permutation, normalized so $\text{mean}_j(c_j) = 1$. It is Proposition 1's
 1145 form, the **dose is** $\rho = e^{2\kappa}$, $\kappa=0$ is bit-identical to d_2 standalone, and normalizing the mean while
 1146 keying the permutation on (seed, round) holds aggregate scale fixed and stops the dose parking ad-
 1147 versaries at a fixed end: neither of which, as it turns out, pins Δ_c . Frozen grid $\kappa \in \{0, 0.5, 1.0, 2.0\}$,
 1148 four cells across three invariance classes. What matters to C2 is not ρ but what the rescaling does to
 1149 the statistic d_2 reads, so we measured that first, with no ASR involved, establishing that *there is a*
 1150 *dose* (0 at identity, rising to 0.87/0.93 for class (c)) and that *class-(a) invariance holds in code*. *Pri-*
 1151 *mary, pooled*: across all 16 (arm, κ) cells the ASR rise over that arm's own $\kappa=0$ rung is positively
 1152 rank-correlated with measured disturbance (Spearman, one-sided, $\alpha=0.05$). *Secondary, per arm*:
 1153 monotone rise for class (c), **FLAT** (the frozen label, scored as equivalence) for class (a), shallower
 1154 for class (b); one-sided Jonckheere–Terpstra at $n=5$, “flat” scored as *equivalence* ($|\Delta| < 0.15$, every
 1155 rung below 0.5), a margin fixed in advance. **The ladder was falsifiable both ways**: `cos_krum`'s
 1156 aggregate is still $c_j u_j$ for the selected j , so suppression can be lost through the magnitude channel,
 1157 and a class-(c) arm inside the margin would refute the central fragility claim. **No frozen prediction**
 1158 **uses S** : measured, S is **not a stable per-attack constant** (medians 0.57/0.92, range to 12.59), a
 1159 limit on the theorem's checkability. Design, grid, arms, shapes and refutation criteria were frozen in
 1160 advance (Appendix M).

1161 **Outcome: one arm confirmed, two refuted, one indeterminate; the pooled test fails.** All 80
 1162 runs completed and were scored per seed against the frozen rules (Table 5, App. N). The **primary**
 1163 pooled prediction is **refuted**: Spearman $\rho_s = +0.351$ over 16 cells, one-sided $p = 0.091$, against
 1164 a frozen $\rho_s > 0$ at $p < 0.05$; **in three of four arms ASR falls as the dose rises**. Arm 2 confirms
 1165 decisively: reputation $0.017 \rightarrow 0.779$ ($p = 1.7 \times 10^{-6}$) at nearly constant accuracy, the paper's
 1166 only prospective positive for C2, and it establishes the *destruction* half only. Arm 1 refutes more
 1167 than its own shape: Krum, whose statistic the fragility claim is stated over, shows no rise (inside
 1168 the equivalence margin) despite measured disturbance 0.867, and does not reproduce our within-
 1169 defense contrast ($0.062 \rightarrow 0.779$ under RFA at a *lower* disturbance), so selection-flip rate cannot be
 1170 its mechanism. Arm 3 is indeterminate and uninterpretable anyway (`cos_krum/pixel` falls $0.490 \rightarrow$
 1171 0.086 , the last rung below the accuracy gate) and arm 4 is refuted outright (`coord_median/pixel`
 1172 falling $0.443 \rightarrow 0.170$ at accuracy ≥ 0.69). Both fell where a rise was frozen, which is why neither
 1173 is scored in our favour.

1174 **Two faults, both ours, and now measured rather than conjectured.** *The instrument.* Normal-
 1175 izing $\text{mean}(c)=1$ fixes the aggregate's scale but not any client's *relative* weight, and because the
 1176 dose ignores adversary status a $55 \times$ spread mostly *dilutes* whichever client carries the poison: on
 1177 the same rounds the ladder's Δ_c drifts monotonically negative with the dose ($0.000 \rightarrow -0.140 \rightarrow$
 1178 $-0.240 \rightarrow -0.336$), so three arms measure attenuation rather than C2 disturbance. **The adversary-**
 1179 **targeted dose is the fix**: Mode S is the Δ_c -**pinned** instrument (Δ_c exactly 0 at every rung, ρ up to
 1180 54.6) and Mode A the Δ_c -**swept** one. *The abscissa.* At matched disturbance two of six *real* upstream
 1181 anchors miss the synthetic curve badly (RFA $\rightarrow$ Krum $+0.767$, NormClip $\rightarrow$ Reputation $+0.789$), and
 1182 not through a range mismatch; Δ_c locates the discrepancy, as Proposition 3 says it may, those two
 1183 being targeted far beyond any ladder rung ($-1.21, -0.72$ against -0.34) at ρ of only 14.8 and 4.19
 1184 against the ladder's 54.6: **matched on the statistic, badly unmatched on the targeting. But this**
 1185 **rescues no predictor**: Δ_c 's magnitude does not order suppression either, and FoolsGold's is *pos-*
 1186 *itive*: **post-hoc and exploratory**, computed with no ASR and rescoring nothing (Appendices N,
 1187 R). Two checks did pass: the $\kappa=0$ rung reproduces every standalone baseline exactly per seed, and
 1188 `cos_krum`'s selection never moved (0/120).

1189 K.5 Mode A, the four aggregators, and the pre-registered second dataset

1190 **Mode A: the two regimes, and where they hold.** Mode A (payload-only) closes the other chan-
 1191 nel: it keeps the benign weights uniform and sets the adversary-to-benign weight ratio to $r = e^\nu$
 1192 at fixed mean, so the adversarial share sweeps $0.049 \rightarrow 0.710$ with the statistic’s dispersion held
 1193 fixed: the mirror of Mode S. That gives Theorem 1 a single falsifiable shape: ASR single-peaked
 1194 at $r=1$, since an over-attenuated adversary carries no payload (Regime B) and an over-amplified
 1195 one is discarded by a rank aggregator (Regime A). It is **confirmed for coord_median**, the rank
 1196 aggregator the regimes are stated for: ASR rises $0.124 \rightarrow 0.315 \rightarrow 0.443$ into $r=1$ and falls
 1197 $0.443 \rightarrow 0.296 \rightarrow 0.098$ past it, both sides significant ($p_\downarrow=3.0\times 10^{-4}$, 6.4×10^{-4}) at accuracy
 1198 ≥ 0.74 (Figure 3). It is **refuted for reputation**, which is monotone, maximal at the most attenuated
 1199 rung (0.894 at $r=0.135$), because an attenuated adversary reads as *consistent* to a reputation score.
 1200 Reputation is not a rank aggregator, so the shape was never licensed for it.

1201 **The same channels, across four aggregators of three structural kinds.** Across selectors on pair-
 1202 wise and on cosine distance (krum, cos_krum), a coordinate-wise order statistic (coord_median)
 1203 and a weighted averager with cross-round state (reputation), the negative is **not Krum-specific**: the
 1204 decision channel moves by 0.482–1.000 on three of the four while **the support of the admitted ad-**
 1205 **versarial mass changes in none of the 60 measured rounds of any of the four**, and on reputation
 1206 and coord_median that mass is *present* at baseline in 48 of those 60, so their counts are not floors
 1207 the way Krum’s is, and the only two arms whose ASR *rises* are the only two with nonzero $\Delta\Lambda_a$. **An**
 1208 **analysis of runs already reported, not a new experiment, and it orders nothing**: the larger $\Delta\Lambda_a$
 1209 (0.033) belongs to the arm with the *smaller* rise (Appendix C).

1210 **A second dataset and architecture, pre-registered.** The same cell on EMNIST-
 1211 byclass/simple_cnn in place of CIFAR-10/cifar_cnn, frozen in git before any EMNIST-byclass
 1212 ASR existed. Premise measured first, no ASR anywhere: the decision moves on 0.643–0.714 of
 1213 rounds, the support of the admitted mass unchanged in all 56 measured rounds (again a floor: no
 1214 adversary admitted at any rung). **Suppression shows no evidence of a practically meaningful**
 1215 **change**: $\Delta = -0.017$ against the frozen ± 0.15 margin, beside CIFAR-10’s -0.026 , at accuracy
 1216 $0.767 \rightarrow 0.835$, so the gate never binds. **Replicated on a second dataset and architecture**: the
 1217 frozen wording, never “generalizes”. Two limits from the freeze: the Jonckheere–Terpstra was
 1218 specified to catch a monotone *rise* and the trend is a *fall* ($z=-1.563$), so it is not read as support for
 1219 flatness; and with the identity rung at 0.027 the refuting branch was reachable (> 0.177) while the
 1220 indeterminate one was not, so equivalence is to the identity rung, not a low absolute ASR.

1221 It addresses neither recall nor discovery, and nothing in §5 is rescored.

1222 L Verification of the invariance mirrors

1223 All eight metric-swap labels were frozen in advance, and these mirrors are mechanism-
 1224 isolating ablations on the same footing as the instruments of §4. We check the statis-
 1225 tic directly rather than inferring invariance from an outcome (verify_cos_invariance.py,
 1226 results/cos_invariance_check.json), comparing each d_2 ’s internal state on $\{u_i\}$ against
 1227 $\{c_i u_i\}$ for the c_i NormClip and RFA actually produce. Over 18 live rounds (3 seeds $\times$ 2 up-
 1228 streams $\times$ 3 rounds) cos_reputation’s weight vector matches to $\leq 2.5 \times 10^{-7}$ with its ordering
 1229 never changing (0/18) and cos_krum selects the same client in 18/18; the Euclidean originals do
 1230 not (Reputation’s weights move by up to 0.222, ordering changes 12/18; Krum’s selection changes
 1231 6/18). Disturbance tracks the realized weight ratio ρ as Proposition 1 predicts: under NormClip
 1232 ($\rho \in [1.00, 4.90]$) Reputation changes in 3/9 and Krum in 0/9, under RFA ($\rho \in [2.11, 17.38]$)
 1233 in 9/9 and 6/9. On 1000 synthetic draws the separation is cleaner still (0/1000 vs. 992/1000
 1234 for the Reputations, 0/1000 vs. 618/1000 for the Krums), and the mirrors match the shipped im-
 1235 plementations they stand in for to $\approx 10^{-8}$ relative aggregate error. **The post-hoc within-defense**
 1236 **contrast of §5 re-uses this frozen suite**, so it is not pre-registered, and two seeds were added after
 1237 seeing the $n=3$ result: a top-up that cut both ways, degrading the invariant control’s Δ from ≈ 0
 1238 to $+0.19$ while leaving the disturbed/undisturbed split without overlap. One asymmetry matters:
 1239 cos_reputation’s *weights* are exactly invariant but its aggregate is not, since the weighted aver-
 1240 age is over the rescaled $c_i u_i$: a magnitude contraction of the Lemma 1 kind, so the prediction is
 1241 two-sided, with ASR staying low while clean accuracy pays a contraction cost.

M Dose–response caveats recorded before the freeze

The pre-freeze disturbance measurement. The abscissa was measured over 5 seeds $\times$ 3 rounds per cell with no ASR involved, which is why it precedes the freeze: the class-(a) mirror’s cosine scores move by at most 1.1×10^{-4} in 0/120 cells, i.e. the invariance holds in the shipped code and not merely on paper, while the class-(c) arms reach 0.87 and 0.93. **Attack assignment.** The class-(c) arms are on model-scaling and the class-(a)/(b) arms on the pixel backdoor: not a free choice, since on model-scaling *no* invariant d_2 suppresses at usable accuracy (§5). “Same attack” is required *within* each ladder, which is where the causal claim lives; across arms we compare curve shapes and slopes, never levels. `coord.median/pixel` sits only 0.057 below threshold, so its binary above/below-0.5 reading is weak either way. The dose is uncorrelated with adversary status by construction but not adversary-*adversarial*; a dose that deliberately up-weights adversaries would be strictly harder and is not run. A synthetic transform is not a real defense, so we overlay the two *real* upstreams already measured on these cells (NormClip, RFA) as anchors at their own measured disturbance; landing far off the synthetic curve would be a finding about the instrument and reported as one. Design, grid, arms, shapes, margin and refutation criteria are frozen in `pre_registration_dose_response.md` at commit e711a95, timestamped before `results/dose_response/` existed and reproduced verbatim in the runner’s arm table. **What it does not address:** recall, and discovery: `fg→rfa` remains a false negative of our own rule.

N Dose–response ladder: per-seed detail, intervals, and the anchor residuals

Table 3: **The confounded ladder against the Mode-S instrument, on the one arm that has both.** Same downstream defense, same attack, same seeds, same nominal dose, and the two designs share their $\kappa=0$ cell *numerically* (mean and worst per-seed gap 0.00e+00, asserted at 10^{-12}). At the top rung they are matched on aggregate displacement, decision change and mean clean accuracy, so **the ladder’s fall is not an accuracy collapse**; they differ chiefly in the **adversarial coefficient channel**, and the ASR effect **changes sign**. **The one non-match, disclosed:** aggregate adversarial influence Λ_a moves in *both* designs, 0.081 against 0.033 ($2.4\times$), so neither column is a clean statistic-only intervention, which is why the contrast is “differ chiefly in the coefficient channel”, never “only in it”. *Vocabulary:* this row is the $\Delta\Lambda_a$ column of the channel table (the adversarial share of selected coordinates, the same field in the same artifact) and **not** its Δ adm. column, which counts rounds where the *support* of the adversarial mass changes and reads 0.000 on every row: what moves here is how much weight the adversary’s direction carries, not whether it enters at all. A two-column contrast on one arm: it locates where the two designs differ and attaches no p -value to the difference between the columns. Every cell recomputed from per-seed and per-round rows by `build_comparability_table.py`.

	Confounded ladder dose_kappa	Instrument (Mode S) doseS_kappa
downstream d_2	<code>coord.median</code>	<code>coord.median</code>
attack	committed pixel	committed pixel
seeds	42–46	42–46
nominal dose	κ 0.0 $\rightarrow$ 2.0, ρ 1 $\rightarrow$ 54.6	κ 0.0 $\rightarrow$ 2.0, ρ 1 $\rightarrow$ 54.6
benign coefficients	mean 1 over <i>all</i> clients	mean 1 over <i>benign</i> clients
adversarial coefficients	drawn from the same spread	pinned $c = 1.0$ exactly
adv. coefficient share	0.267 $\rightarrow$ 0.187	0.266667, constant
Δ_c	+0.000 $\rightarrow$ -0.336	0 by construction (read back 9e–08)
aggregate displacement	0.000 $\rightarrow$ 0.722	0.000 $\rightarrow$ 0.670
decision change	0.000 $\rightarrow$ 0.487	0.000 $\rightarrow$ 0.482
adversarial influence $\Delta\Lambda_a$	0.000 $\rightarrow$ 0.081	0.000 $\rightarrow$ 0.033
mean clean accuracy	0.767 $\rightarrow$ 0.690	0.767 $\rightarrow$ 0.698
ASR	0.443 $\rightarrow$ 0.170	0.443 $\rightarrow$ 0.541
ASR effect	-0.272 ($p_{\downarrow} = 0.0015$)	+0.098 ($p_{\uparrow} = 0.069$)

Table 6 is every cell of the ladder; Figure 4 plots it on the measured axis and Figure 5 on the nominal one. Both are produced by `analyze_dose_response.py` and `dose_response_figure.py` from `results/dose_response/`; nothing below is transcribed from a log.

Table 4: Mode S (statistic-only), $n=5$ per arm ($n=8$ for `cos_krum`): mean ASR with the two measured channels. **The decision channel moves far more than ASR throughout. The admission row is the admitted adversarial mass Λ_a , not the support of that mass, which is the Δ adm. of §4 and is 0.000 in every measured round of every arm.** [†] below the 0.35 accuracy gate. [‡] the replication arm, rank-discordant by design, frozen separately: $\Delta=+0.098$ is below reputation’s $+0.178$, so the cross-arm admission *ordering* is refuted while the negative replicates. Verdicts otherwise scored against readings frozen in advance; per-seed values in Appendix N.

arm	channel	$\rho=1$	2.72	7.39	54.6	test	verdict
krum / scaling	ASR	0.062	0.016	0.023	0.036	$p_{\uparrow}=0.70$	H-admission confirmed, H-stat. refuted
	decision	0.000	0.533	0.800	0.733		
	admission	0.000	0.000	0.000	0.000		
reputation / scaling	ASR	0.017	0.019	0.025	0.195	rise $+0.178$ vs. ladder $+0.762$ ($p=0.0017$)	H-admission confirmed (mostly atten.)
	decision	0.000	0.933	1.000	1.000		
	admission	0.000	0.002	0.008	0.020		
cos_krum / pixel	ASR	0.584	0.411	0.295	0.159 [†]	$p_{\downarrow}=8 \times 10^{-4}$ $ \Delta =0.425$	refuted (Prop. 1 practical)
	decision	0.000	0.000	0.000	0.000		
	admission	0.000	0.000	0.000	0.000		
coord_median / pixel [‡]	ASR	0.443	0.442	0.490	0.541	rise $+0.098$ $< +0.150$ ($p_{\uparrow}=0.069$)	ordering refuted ; negative repl.
	decision	0.000	0.232	0.358	0.482		
	admission	0.000	0.005	0.013	0.033		

Table 5: The ladder as measured, $n=5$, per-attack mean ASR (mean clean accuracy). *Predicted* shapes were frozen before any run (Appendix X); $p_{\uparrow}$ is the one-sided Jonckheere–Terpstra probability of an increasing trend, the only direction any frozen rule tests. An arm meeting neither its confirmation nor its refutation criterion is reported *indeterminate* rather than scored in our favour. [†] below the 0.35 accuracy gate: a collapsed model, not suppression. Intervals, per-seed values, anchor residuals and both plotted axes: Appendix N.

arm (d_2 / attack)	predicted shape	mean ASR (clean acc.)				JT $p_{\uparrow}$	verdict
		$\kappa=0$	0.5	1.0	2.0		
1 krum / scaling	monotone rise	0.062 (0.57)	0.011 (0.65)	0.014 (0.68)	0.054 (0.62)	0.63	refuted confirmed
2 reputation / scaling	monotone rise	0.017 (0.78)	0.022 (0.78)	0.487 (0.77)	0.779 (0.70)	1.7×10^{-6}	
3 cos_krum / pixel	FLAT (control)	0.490 (0.61)	0.450 (0.57)	0.153 (0.47)	0.086 (0.22) [†]	0.999	indeterminate
4 coord_median / pixel	rise, shallower	0.443 (0.77)	0.428 (0.76)	0.327 (0.75)	0.170 (0.69)	0.998	refuted

The replication arm’s post-hoc seed-matched difference. The frozen primary for the `coord_median` replication (§4) is the *unpaired* $\Delta = +0.098$, and it is not superseded by what follows. Between-seed variance on this cell is large (identity ASR spread 0.287–0.677), which the unpaired test cannot see through, so we also report the seed-matched difference: positive in 5/5 seeds, paired $t = 3.39$, one-sided $p = 0.014$. **This is post-hoc and was not pre-registered**; we report it because it is the honest description of a small consistent effect, and we label it post-hoc wherever it appears. It does not change the frozen verdict: the admission *ordering* is refuted either way.

The primary test, and two post-hoc sensitivities that cannot change its verdict. The frozen primary test is a one-sided Spearman correlation of *rise* against *measured disturbance* pooled over all 16 cells, confirmed only at $\rho_s > 0$ with $p < 0.05$: we obtain $\rho_s = +0.351$, $p = 0.091$, so it is **refuted**. Two sensitivities are recorded because they are informative about *why*, and are labelled post-hoc because neither was frozen. (i) The four $\kappa=0$ cells are constructed (0, 0) points (zero disturbance and zero rise by definition) and dropping them gives $\rho_s = +0.780$, $p = 0.001$ ($n=12$); so the pooled correlation is carried by the arm-2 rise and destroyed by the two arms whose ASR falls, not by the identity anchors. (ii) Against the *nominal* dose κ (i.e. $\log \rho$) rather than measured disturbance, $\rho_s = -0.354$, $p = 0.911$: the wrong sign, which is the same fact as three arms falling. Neither is scored.

Per-arm trend tests. One-sided Jonckheere–Terpstra across the four ordered rungs, 5 seeds each, $\alpha = 0.05$; the normal-approximation p agrees with a 20,000-draw permutation p to three decimals throughout. Arm 1: $J = 70$, $z = -0.34$, $p_{\uparrow} = 0.632$. Arm 2: $J = 144$, $z = +4.64$, $p_{\uparrow} = 1.7 \times 10^{-6}$. Arm 3: $J = 29$, $z = -3.10$, $p_{\uparrow} = 0.999$ ($p_{\downarrow} = 9.8 \times 10^{-4}$); dropping the gated $\kappa=2.0$




Figure 2: **The preservation chain, and where it breaks.** (P1) $\Rightarrow$ (P2) $\Rightarrow$ (P3) holds by definition; **below (P3) the chain breaks, in both directions**, and both witnesses come from the oracle intervention that pins every adversary’s coefficient and so closes Lemma 1’s attenuation channel by construction (§4). The pooled statistic-level prediction we had frozen is separately **rejected** by a pre-registered ladder (§5), and a pre-registered replication in a different invariance class reproduces the negative but **refutes the cross-arm admission ordering**, so **admission is more causally proximal than the statistic, and not a validated predictor**. Every number in the figure is recomputed from the scoring scripts’ own JSON by `story_chain.figure.py`.

Table 6: All 16 rungs, $n=5$. *disturb* is the independently measured disturbance of that arm’s own statistic at that rung, computed before the freeze with no ASR involved; *rise* is ASR relative to the arm’s own $\kappa=0$ rung and is the primary test’s ordinate. CIs are Student- t at $n=5$ and are wide by design: the ladder was powered to detect a monotone trend, not to resolve single cells. † fails the 0.35 clean-accuracy gate.

arm	κ	ρ	dist.	ASR	95% CI	rise	acc	per-seed ASR (42–46)				
krum	0	1.00	0.000	0.062	[0.019, 0.105]	+0.000	0.565	0.115	0.021	0.047	0.057	0.069
scaling	0.5	2.72	0.600	0.011	[−0.001, 0.023]	−0.051	0.652	0.000	0.010	0.011	0.027	0.007
class (c)	1.0	7.39	0.733	0.014	[0.003, 0.024]	−0.048	0.682	0.003	0.015	0.021	0.021	0.008
	2.0	54.60	0.867	0.054	[0.002, 0.106]	−0.008	0.619	0.005	0.039	0.077	0.113	0.037
reputation	0	1.00	0.000	0.017	[0.012, 0.022]	+0.000	0.777	0.014	0.012	0.017	0.023	0.019
scaling	0.5	2.72	0.867	0.022	[0.014, 0.031]	+0.005	0.779	0.016	0.018	0.024	0.033	0.020
class (c)	1.0	7.39	0.867	0.487	[0.349, 0.624]	+0.470	0.765	0.439	0.547	0.578	0.557	0.313
	2.0	54.60	0.933	0.779	[0.673, 0.885]	+0.762	0.701	0.710	0.843	0.730	0.896	0.714
cos_krum	0	1.00	0.000	0.490	[0.058, 0.922]	+0.000	0.607	0.003	0.269	0.630	0.864	0.686
pixel	0.5	2.72	0.000	0.450	[0.130, 0.769]	−0.041	0.569	0.764	0.236	0.165	0.648	0.435
class (a)	1.0	7.39	0.000	0.153	[−0.127, 0.433]	−0.337	0.466	0.132	0.090	0.000	0.543	0.001
	2.0	54.60	0.000	0.086	[−0.100, 0.271]	−0.405	0.220 †	0.083	0.000	0.000	0.345	0.000
coord_median	0	1.00	0.000	0.443	[0.263, 0.622]	+0.000	0.767	0.395	0.287	0.449	0.677	0.405
pixel	0.5	2.72	0.251	0.428	[0.241, 0.615]	−0.015	0.764	0.454	0.215	0.430	0.638	0.402
class (b)	1.0	7.39	0.378	0.327	[0.181, 0.472]	−0.116	0.749	0.351	0.139	0.352	0.463	0.328
	2.0	54.60	0.487	0.170	[0.113, 0.228]	−0.272	0.690	0.179	0.090	0.204	0.192	0.187

1286 rung leaves $z = -2.06$, $p_\downarrow = 0.020$, so the decrease is not an artifact of the collapsed cell alone.
 1287 Arm 4: $J = 31$, $z = -2.96$, $p_\uparrow = 0.998$ ($p_\downarrow = 1.5 \times 10^{-3}$). No frozen rule tests the decreasing
 1288 direction; $p_\downarrow$ is descriptive.

1289 **Clean accuracy against dose separates two ways of losing ASR.** Descriptive, not pre-registered.
 1290 Accuracy trends by the same test: arm 1 $z = +1.35$, arm 2 $z = -3.30$, arm 3 $z = -3.50$, arm
 1291 4 $z = -4.11$. Arm 3 pays $0.607 \rightarrow 0.220$: a 0.39 collapse through the Lemma 1 contraction
 1292 channel, and the reason its $\kappa=2.0$ cell is uninterpretable rather than a success. Arm 4 pays only
 1293 $0.767 \rightarrow 0.690$ while ASR more than halves, so there the dose genuinely *helps* coordinate median:




Figure 3: **The two regimes, exhibited: the payload dose separates them.** *The picture of regimes A and B of §2, produced by the instrument of §5.* With the statistic’s dispersion fixed and only the adversary/benign weight ratio r swept, coord_median (a rank aggregator, the class the regimes are stated for) is single-peaked at $r=1$ as Theorem 1 predicts, while reputation is monotone: off-spec, and the theorem is not claimed for it. Bars are Student- t 95% intervals, recomputed per seed by analyze_targeted_dose.py; the Mode-S companion is Table 4.




Figure 4: **The measured axis: ASR against disturbance of each arm’s own statistic**, an abscissa fixed before the freeze ($n=5$, Student- t 95% intervals; the figure script imports the arm and anchor construction from the scoring script, so the two cannot diverge). Stars and crosses are the two *real* upstream transforms measured on these same cells: at matched disturbance they sit +0.77 and +0.79 ASR *above* the synthetic curve, though their realized weight ratios (mean 3.61 and 11.93, maximum 4.90 and 17.38) lie *inside* the ladder’s span; §5 attributes the gap to Δ_c , post-hoc. The class-(a) arm’s four rungs coincide at disturbance exactly 0 and its ASR falls anyway: Lemma 1’s magnitude channel, not C2 disturbance. Hollow markers fail the 0.35 accuracy gate.

1294 a refutation of our predicted shape, but a benign mechanism. Arm 2 is the cleanest cell in the ladder,
1295 losing 0.076 accuracy while ASR rises 46×.

1296 **Anchor residuals, and why the instrument is the finding.** Residual is real-transform ASR
1297 ($n=5$, from metric_swap_topup) minus the synthetic rung at the nearest measured distur-
1298 bance. NormClip→Krum: disturbance 0/9, ASR 0.062 vs. rung 0.062, residual +0.001.
1299 RFA→Krum: 6/9 = 0.667, ASR 0.779 vs. rung 0.011 at disturbance 0.600, residual
1300 +0.767. NormClip→Reputation: 3/9 = 0.333, ASR 0.806 vs. rung 0.017, residual +0.789.




Figure 5: **The nominal axis:** the same 16 cells against the dose $\rho = e^{2\kappa}$, coloured by Proposition 1 class. Of the two class-(c) arms the fragility claim is stated over, only reputation rises; the other two arms *fall*. This is the axis the frozen per-arm shapes were stated on, and the pooled prediction was stated on the measured axis of Figure 4. Hollow markers fail the 0.35 accuracy gate.

1301 RFA→Reputation: 9/9, 0.895 vs. 0.779, residual +0.116. Both *cos_krum* anchors sit at distur-
1302 bance 0 and land (+0.000, −0.044). *coord_median* was not in the metric-swap menu, so arm 4 has
1303 *no* real-transform anchor and its ecological validity is untested. The two large residuals are not a
1304 range mismatch: realized ρ is mean 3.61 / max 4.90 for NormClip and mean 11.93 / max 17.38 for
1305 RFA, both inside the ladder’s [1.00, 54.60] span. What differs is *which* clients get reweighted (Norm-
1306 Clip and RFA reweight in a way correlated with adversary status, our permutation deliberately does
1307 not) and we state that as the hypothesis caveat (iii) was written to test, not as a measured quantity:
1308 the shipped disturbance logs record only $c_{\min}$, $c_{\max}$ and ρ per round, not per-client coefficients, so
1309 an adversary-versus-benign coefficient contrast requires a new (cheap, no-ASR) measurement.

1310 **Two checks that passed.** The $\kappa=0$ rung reproduces each arm’s published standalone value exactly
1311 on every shared seed (16 seed-level comparisons across four arms, largest $|\Delta\text{ASR}|$ and $|\Delta\text{acc}|$ both
1312 $0.00\text{e}+00$ at tolerance 10^{-6}), so $\kappa=0$ really is the identity and the fresh rung and the published base-
1313 line are the same quantity. **The first-rung fall we rest the (P4)↯(P5) witness on**, in full: *cos_krum*
1314 / pixel at $\rho=2.72$ falls 0.173 with clean accuracy moving only $0.599 \rightarrow 0.567$, paired $t=-1.976$,
1315 one-sided $p=0.044$, falling in 5/8 seeds. And *cos_krum*’s selection never changed in any of the
1316 120 measured cells, with maximum relative score drift 1.1×10^{-4} : Proposition 1’s classification of
1317 which statistics a positive per-client rescaling can disturb is verified directly, independently of the
1318 outcome the ladder then failed to predict.

1319 O Secondary results: baselines, the adaptive stress test, scale, and switching

1320 **The margin behaves like a certificate, not just a switch.** C3 is binary, but the margin m is
1321 measurable and its *strength* tracks security: over the 29 measured rounds with $w_a > 0$, higher
1322 margin goes with lower realized adversarial mass (Spearman $\rho = -0.48$, $p = 0.009$), Corollary 1’s
1323 bound holding in every one, a number rather than a verdict, though the link is directional, not fitted,
1324 and post-hoc.

1325 **Single-defense baselines.** Measured under the identical protocol (5 seeds, 50 rounds), no individ-
1326 ual defense achieves low max-committed ASR: CM 0.52, FG 0.73, Rep 0.84, TM 0.85, RFA 0.89,
1327 NC 0.94. Against the best single defense (CM, 0.52), FG→RFA (0.093) is $5.6\times$ lower and FG→CM
1328 (0.107) is $4.9\times$ lower. FoolsGold alone reaches ASR 0.20 on model-scaling only by collapsing the
1329 model to 10% accuracy, so it is not suppression: consistent with its C1 failure.

1330 **Adaptive attacks against FG→RFA.** We evaluate one whitebox adaptive attack family that
1331 jointly evades both defenses: disjoint coordinate blocks minimize cosine similarity (evade FG) while




Figure 6: **The criterion is a mechanism test, and the mechanism is measurable.** (a) What C2 checks: the upstream transform T and the statistic S that d_2 reads determine an *invariance class* (Proposition 1), and the class determines the verdict with no composition run, which is why norm clipping does not blind FoolsGold but does blind Reputation. (b) The measured chain for FG→RFA, per round, on the transformed points RFA actually receives; each marker is one round with $w_a > 0$ (rounds where FoolsGold zeroes every adversary are Lemma 1 and carry no payload), large markers are mean $\pm$ s.d., and values below 10^{-4} are drawn at 10^{-4} . Corollary 1’s bound holds in every round, and the $\alpha=0.1$ collapse is visible in the theorem’s own quantity: m falls from 4.71 ± 1.51 ($n=17$) to 2.06 ± 2.35 ($n=12$), with several rounds at $m < 1$ where the bound degrades to $n_a/(n_b+n_a)$, and realized Λ_a rises $0.003 \rightarrow 0.085$ (max 0.50) as composition ASR goes $0.093 \rightarrow 0.569$.

Table 7: Framework validation (representative pairs; full tallies in text). Threshold: LOW < 0.5 , HIGH ≥ 0.5 . [§]Pre-registered top-up to $n=30$ on the arm that sets the max: mean 0.463, interval $[0.412, 0.514]$, which still crosses 0.5, so this pair is certified on the frozen mean rule and not resolved below threshold.

Composition	Category	Predicted	Max ASR	Correct?
<i>PASS: $C1 \wedge C2 \wedge C3$ (predicted LOW)</i>				
FoolsGold $\rightarrow$ RFA	PASS	LOW	0.093	✓
FoolsGold $\rightarrow$ CoordMedian	PASS	LOW	0.107	✓
Reputation $\rightarrow$ CoordMedian	PASS	LOW	0.463 [§]	✓
<i>C1-FAIL: no constituent suppresses (predicted HIGH)</i>				
Reputation $\rightarrow$ RFA	C1-FAIL	HIGH	0.825	✓
NormClip $\rightarrow$ FoolsGold [†]	C1-FAIL	HIGH	0.887	✓
RFA $\rightarrow$ FoolsGold [†]	C1-FAIL	HIGH	0.914	✓
<i>C2-FAIL: discriminative property destroyed (predicted HIGH)</i>				
NormClip $\rightarrow$ Reputation	C2-FAIL	HIGH	0.819	✓
<i>C3-FAIL: margin compressed (predicted HIGH)</i>				
FoolsGold $\rightarrow$ TrimmedMean	C3-FAIL	HIGH	0.508	✓
NormClip $\rightarrow$ RFA	C3-FAIL	HIGH	0.853	✓
<i>False negatives (predicted HIGH, measured LOW)</i>				
NormClip $\rightarrow$ CM / RFA $\rightarrow$ CM	C3-FAIL	HIGH	0.433 / 0.485	×

[†]Reclassified from C2-FAIL: by Proposition 1(a) these preserve C2.

1332 ϵ -ball projection stays near the geometric median (evade RFA). The adversary has full state access,
1333 and these constraints are in tension. ASR against FG→RFA (5 seeds each, base quoted on the
1334 *same* five seeds *and the same pixel attack the adaptive family is built from*, so the comparison is
1335 seed-matched and same-attack rather than against a denominator at a different n or the pairs worse
1336 scaling arm; the $n=30$ figures below move *both* legs to 30 and so are seed-matched too): base com-
1337 mitted pixel 0.045 ± 0.013 ; Neurotoxin [Zhang et al., 2022] top-20% at scale 2.0 0.048 ± 0.009 ;
1338 projection only, no decorrelation, $\epsilon=2$: 0.050 ± 0.008 . With joint constraints the sweep over ϵ gives
1339 0.036 ± 0.015 ($\epsilon=0.5$), **0.167 ± 0.069** ($\epsilon=1$), 0.130 ± 0.030 ($\epsilon=2$), 0.124 ± 0.025 ($\epsilon=4$). ASR
1340 peaks at $\epsilon=1$: tighter projection over-constrains the attack, larger ϵ fails to evade RFA, and without

decorrelation FoolsGold suppresses it outright, so the competing constraints leave a narrow effective window. This is an empirical finding (Theorem 1 concerns mechanism preservation and does not bound an adaptive adversary whose optimization is untied to its conditions) and it probes one natural family of joint constraints at a handful of ϵ values, not the adaptive attack space. The ϵ sweep above is held at the uniform $n=5$ it compares at, but the $\epsilon=1$ cell has since been topped up to $n=30$ on seeds 42–71, where it reads 0.157 ± 0.073 against the pixel arm’s own $n=30$ mean of 0.064: **both legs recomputed on the same 30 seeds, the ratio is $2.4\times$, and $1.7\times$** against the pairs scaling arm the adaptive family does not attack. At $n=5$ those two ratios read $3.7\times$ and $3.3\times$, and almost all of the fall is the denominator rather than the attack weakening: the base arm’s $n=30$ distribution is right-skewed (mean 0.064, median 0.052), so seeds 42–46 sit in its left tail. A ratio quoting a numerator at one n against a denominator at another is not reported here.

Scale. At $N=100$, $f=0.05$, ResNet18 both pre-registered PASS predictions hold: $\text{FG} \rightarrow \text{RFA}$ max $\text{ASR} = 0.069$, $\text{FG} \rightarrow \text{CM} = 0.152$. We flag that these runs reach only ≈ 0.29 clean accuracy at 50 rounds (ResNet18 is far from converged at this budget), so they establish that suppression transfers across scale and architecture, not that the composition is deployable at $N=100$ as configured.

0.1 Heterogeneity, and where the criterion breaks

Earlier drafts swept heterogeneity for single defenses only; no composition had been run at any $\alpha \neq 0.5$. We therefore ran all three PASS pairs across $\alpha \in \{0.1, 1.0, 10.0\}$ (3 seeds, both attacks, 50 rounds; max-committed ASR / clean accuracy):

Composition	$\alpha=0.1$	$\alpha=0.5$	$\alpha=1.0$	$\alpha=10.0$
FG→RFA	0.569 / 0.20	0.093 / 0.50	0.051 / 0.57	0.063 / 0.62
FG→CM	0.433 / 0.40	0.107 / 0.68	0.133 / 0.74	0.129 / 0.78
Rep→CM	0.548 / 0.48	0.463 / 0.76	0.377 / 0.79	0.326 / 0.81

In our sweep, which is three seeds per cell and fixes direction rather than magnitude, the certified compositions stay below threshold for $\alpha \in \{0.5, 1.0, 10.0\}$, including the near-IID setting where one might expect FoolsGold to founder on legitimately similar benign updates: $\alpha=10$ is indistinguishable from $\alpha=0.5$. **Under extreme heterogeneity ($\alpha=0.1$) two of the three PASS compositions cross the threshold** (FG→RFA 0.569, Rep→CM 0.548) and only FG→CM survives, at 0.433: a real boundary condition on the framework, with caveats in both directions: FG→RFA at $\alpha=0.1$ is strongly bimodal across seeds (0.065–0.923), so at $n=3$ the direction is clear but the magnitude is not, and clean accuracy degrades badly (0.20–0.48), so the regime stresses utility as well as security. Crucially Theorem 1 supplies a diagnostic quantity that retrospectively explains this boundary: re-measuring the transformed points at $\alpha=0.1$, RFA’s separation margin $m := \Delta'_{\text{sep}} / [2(R'_B + \delta')]$ collapses from an all-round mean of 3.24 (met in 68% of rounds) to 1.01 (met in 20%), over the payload-carrying rounds alone, the ones Figure 6 plots, the same collapse reads 4.71 ± 1.51 ($n=17$) to 2.06 ± 2.35 ($n=12$) with the per-round minimum falling $2.82 \rightarrow 0.00$, the difference being that the all-round means average in the Lemma 1 rounds at $m=0$, exactly the boundary $m \rightarrow 1$ where Corollary 1’s bound degrades to 40% adversarial mass, because extreme heterogeneity inflates the benign radius R'_B until the benign cluster is no longer separated from the adversarial points, which is the precondition the theorem requires. A practitioner running the 5–10 round pilot we recommend would measure $m \approx 1$ and decline to deploy. The criterion is therefore mechanistically diagnostic at this boundary (the property that makes a screening rule usable) though this is a post-hoc explanation of a known failure, not a pre-registered prediction of it.

0.2 Composition vs. randomized switching

We make no formal claim about persistence. An earlier version derived the value of temporal switching from the measured retention $\gamma \approx 0.66$, on the reasoning that triggers decay slowly enough that intermittent suppression cannot clear them, and offered “high- γ attacks blunt switching” as a qualitative prediction. **That derivation is withdrawn:** it rested on a persistence model this paper does not establish, γ is reported only as a measured property of the tested regime, and nothing in the preservation argument rests on it. What remains below is measured rather than derived.

1388 **Composing against mixing, and the confound we label rather than hide.** An earlier version
1389 reported an order-sensitivity control, fixed $FG \rightarrow CM$ and $FG \rightarrow RFA$ against a randomized 50/50
1390 *ordering*, and read a 9–19 $\times$ gap from it. **That comparison is withdrawn and nothing here rests**
1391 **on it:** CoordMedian supplies no per-client transform, so every $CM \rightarrow FG$ round reduces to FoolsGold
1392 alone, and what the gap measures is one defense per round against two rather than any cost of
1393 randomizing. What we retain are two same-menu comparisons of composing against mixing. Pixel
1394 ASR at $n=5$: $FG \rightarrow CM$ every round gives 0.099 ± 0.020 against 0.679 ± 0.159 for drawing one of
1395 the two per round, a 6.9 $\times$ gap. Oracle ASR at $n=30$: on reputation+TrimmedMean, composing
1396 gives 0.316 ± 0.162 against 0.915 ± 0.065 for the rep 30%/tm 70% mix. **Compute is not matched**
1397 **in either, and we say so rather than imply otherwise:** mixing applies one defense per round where
1398 composition applies two, so what these gaps bound is the cost of using an available defense pair one
1399 at a time, and *not* the cost of randomization at fixed defense computation.

1400 P The metric-swap suite’s two-sided prediction

1401 Because C1 binds for all eight pairs, every frozen label in the metric-swap suite was HIGH (a
1402 consequence of the measured inputs rather than a later choice), so the suite bears on specificity only.
1403 One prediction inside it was genuinely two-sided. *cos_krum*’s C1 verdict on model-scaling turns
1404 on whether the accuracy gate is applied (0.078 ASR at 0.146 clean accuracy), so we froze *both*
1405 readings in advance: primary (gated) $\Rightarrow$ HIGH, secondary (numeric only) $\Rightarrow$ LOW, with the rule
1406 fixed beforehand that a LOW outcome at accuracy ≥ 0.35 would mean the gate was wrong. Both
1407 cells came out **HIGH** (0.508, 0.566), so the gate is vindicated against a pre-registered alternative
1408 that would have overturned it. This is the only place in the paper where our own accuracy gate was
1409 put at risk by a frozen rule.

1410 Q The prospective suite, arm by arm

1411 (i) **FLTrust as d_1 (7 pairs, 2/7): out of scope.** When d_1 alone suppresses, the composition in-
1412 herits that suppression whatever d_2 does, so all five predicted-HIGH pairs are misses even though
1413 their C2/C3 failures are real: FLTrust genuinely destroys Reputation’s consensus-distance signal,
1414 but destroying a mechanism that was never load-bearing costs nothing. This is C0, and it was invis-
1415 ible on the original 42 pairs, where the strongest single defense reaches only 0.519 and the lowest
1416 ASR among all 39 non-PASS pairs is 0.433; no pair was ever out of scope. We apply C0 here
1417 *retrospectively*, which is weaker than a pre-registered exclusion.

1418 (ii) **FLTrust as d_2 (6 pairs, 6/6): in scope, but nearly an identity.** Proposition 1(a) predicts
1419 $\cdot \rightarrow FLTrust$ preserves C2, and it does, but FLTrust gates on a scale-invariant cosine and then *renor-*
1420 *malizes every update to* $\|g_{srv}\|$, discarding magnitude outright, so a positive rescaling $u_i \mapsto c_i u_i$
1421 is *exactly* cancelled: FLTrust discards precisely what d_1 modifies. The per-seed data show this di-
1422 rectly. $Krum \rightarrow FLTrust$ and $MKrum \rightarrow FLTrust$ are **bit-identical** to FLTrust alone on all six runs
1423 (pass-through d_1); $NC \rightarrow FLTrust$ is bit-identical *on the pixel attack*, where no client exceeds $\tau=5$
1424 so $c_i \equiv 1$, and differs only within seed noise on model-scaling, where clipping fires; $Rep \rightarrow$ and
1425 $RFA \rightarrow FLTrust$ likewise differ only within noise. The lone real deviation is $FG \rightarrow FLTrust$ (0.064,
1426 up to 0.084 per seed): the one d_1 producing exact *zeros* rather than positive rescalings, changing
1427 FLTrust’s client set instead of its magnitudes, which is the Lemma 1 interaction. So this arm con-
1428 firms that theory and implementation agree, and demonstrates the degeneracy mechanism cleanly,
1429 but it is **not falsifiable**: every d_1 in our menu is a *non-negative* per-client rescaling, FLTrust cancels
1430 the positive part exactly, and so the only channel left open is the *zero pattern*, which is precisely
1431 where the one deviation appears. Proposition 1(a) permits no other outcome on the strictly-positive
1432 part, and Lemma 1 accounts for the rest.

1433 (iii) **Neither defense is FLTrust (7 pairs, 5/7): the only arm that could have surprised us.**
1434 All were predicted HIGH; correct are $rep \rightarrow krum$ 0.880, $nc \rightarrow mkrum$ 0.943, $rfa \rightarrow mkrum$ 0.875,
1435 $krum \rightarrow rfa$ and $mkrum \rightarrow rfa$ both 0.870. Of those five, only $rep \rightarrow krum$ composes two real mecha-
1436 nisms: MultiKrum at $K=5$ is bit-identical to FedAvg, so $nc \rightarrow mkrum$ and $rfa \rightarrow mkrum$ are Norm-
1437 Clip and RFA standalone, and $krum \rightarrow rfa$ and $mkrum \rightarrow rfa$ are one measurement of RFA standalone
1438 made twice (per-seed $|\Delta ASR| = 0.000$). Neither miss is a mechanism error: $fg \rightarrow krum$ reads 0.078
1439 at 9.9% clean accuracy (a collapsed model, the same hollow-suppression artifact as FoolsGold alone)
1440 and $krum \rightarrow cm$ reads 0.450 against a 0.5 threshold as a degenerate pair reducing to CoordMedian

1441 alone, whose own baseline is 0.519: threshold noise on a borderline constituent. This arm contains
 1442 *zero* predicted-LOW pairs, so like Tier 2 it bears on specificity only.

1443 **The suggestive contrasts, withdrawn.** We had read three matched pairs, d_1 held fixed while d_2
 1444 moved from FLTrust to a Euclidean rule, as a 19–20 $\times$ gap turning on exactly the property Propo-
 1445 sition 1 identifies. Two of them ran against MultiKrum, which at $K=5$ selects *all five* participants
 1446 and is therefore bit-identical to FedAvg (1.9×10^{-6} per aggregation; a genuine selector again at
 1447 $K=10$), so those arms carried no downstream defense at all. Substituting the genuine selector Krum
 1448 does not repair them, and that is the finding: all three Krum arms take their max-committed ASR
 1449 from the pixel backdoor, on which Krum’s own standalone ASR is 0.583, so the gap measures C1
 1450 failure rather than disturbed invariance, and $\text{nc} \rightarrow \text{krum}$ is bit-identical to Krum alone on all three
 1451 seeds because under that attack no client exceeds $\tau=5$ and d_1 is the exact identity. Counting the
 1452 pass-through d_1 rows as well, 8 of the 20 rows reduce to a single defense rather than the 5 labelled
 1453 DEGEN at freeze, and arm (iii) composes two real mechanisms in 2 pairs, not 7. The frozen labels
 1454 and the 13/20 score stand; what we retract is our reading of them. The confound is accordingly
 1455 worse than we reported, in the way Lemma 2 predicts: the invariant-signal d_2 is also the only d_2 in
 1456 this menu that suppresses alone, so effectiveness and mechanism co-vary across the contrast. This
 1457 is why breaking the confound requires holding standalone effectiveness fixed rather than choosing a
 1458 different defense.

1459 R Adversary-correlated dispersion Δ_c , measured

1460 R.1 Heterogeneity is not the mechanism: ρ against Δ_c

1461 **The finding, stated first.** What costs a downstream defense its suppression is not how heteroge-
 1462 neous the upstream transform is but **whether its reweighting is correlated with adversary iden-**
 1463 **tity**, and not even the size of that correlation, which has **two opposite regimes**. This is why our own
 1464 headline predictor failed.

1465 **Two different things a coefficient vector does, and the theory sees only one.** (Symbols: Ap-
 1466 pendix G.) Proposition 1 turns on **dispersion**, $\rho := \max_i c_i / \min_i c_i$ (how far apart the coefficients
 1467 are, which decides whether an ordering survives) and says nothing about *which* clients receive which
 1468 coefficient. That second property is **adversary-correlated dispersion**,

$$\Delta_c := \mathbb{E}[c_i \mid i \in A] - \mathbb{E}[c_i \mid i \notin A], \quad A \text{ the adversarial participants,}$$

1469 and the two are **independently controllable**, not merely uncorrelated in our sample: a transform
 1470 can spread the coefficients arbitrarily at $\Delta_c = 0$ exactly, or leave them nearly equal while shifting
 1471 the adversary’s share, and Proposition 3 exhibits a family of each. **Proposition 1 and Theorem 1**
 1472 **constrain ρ ; Lemma 1’s attenuation channel is about Δ_c** , and §5 finds the real transforms differ
 1473 on the second. Δ_c is *not* a screening quantity: computing it requires knowing A (§6).

1474 **Two mechanism-preserving regimes, at opposite ends of ρ .** Lemma 1 is a second, inde-
 1475 pendent route to suppression, and the two together explain all three PASS compositions. Let
 1476 $r := w_a / \max_b w_b$ and S be the attack’s coordinate extremeness. **(A) Ordering preservation**
 1477 keeps the adversary in the tail where a rank aggregator discards it ($S \cdot r > 1$). **(B) Attenuation**
 1478 instead drives the adversarial transformed value *inside* the benign spread, where it cannot be an ex-
 1479 treme order statistic and its payload is scaled by r ($S \cdot r \ll 1$; $r = 0$ is exact annihilation). The
 1480 hazard is $S \cdot r \approx 1$, and our two weighting defenses sit at opposite ends: FoolsGold’s weights are
 1481 bounded ($\rho \leq 2.1$, $S \cdot r \geq 4.8$), regime A; Reputation’s are wildly unbounded (ρ mean 1.4×10^5)
 1482 precisely *because* it crushes the adversary ($S \cdot r = 0.017$), regime B. So Reputation $\rightarrow$ CoordMedian
 1483 violates the ordering condition by five orders of magnitude and still suppresses (Figure 3, App. N,
 1484 exhibits both regimes). Hence ρ does not order Δ_c (Proposition 3), and Δ_c ’s *magnitude* does not
 1485 order suppression either, the same large $|\Delta_c|$ being protective at one end and catastrophic at the
 1486 other. **Neither quantity is a predictor**; they buy an account of *why* a heterogeneity-based reading
 1487 mispredicts.

1488 $\Delta_c := \mathbb{E}[c_i \mid i \in A] - \mathbb{E}[c_i \mid i \notin A]$ for every per-client transform in the paper,
 1489 over the same 5 seeds $\times$ 3 rounds and, within a round, the *same* raw update stack for every
 1490 transform, so the rows are comparable by construction (measure_coefficient_targeting.py,

1491 results/coefficient_targeting.json). **No ASR is computed and no model is trained per**
1492 **transform**, which is what lets this be added without rescoring anything; the figures are **post-hoc and**
1493 **exploratory** and no frozen prediction depends on them. Coefficients are read back from transformed
1494 norms, $c_i = \|T(u_i)\|/\|u_i\|$, so Table 8 measures the shipped code path rather than the builder’s
1495 description of it. As a provenance check the synthetic families’ adversarial coefficient shares repro-
1496 duce results/admission_measurement.json to 0.0 absolute on all 13 cells, so these rows may
1497 be read on the same axis as the Mode-S and Mode-A figures of §4.

Table 8: ρ says how much a transform spreads the coefficients; Δ_c says who receives the spread. They are not ordered together: RFA and NormClip target the adversary $2\text{--}3.6\times$ harder than the ladder’s most dispersed rung at ρ that is $3.7\text{--}13\times$ *smaller*. Means over non-degenerate rounds; the synthetic rows are attack-pooled since their coefficients do not depend on the attack.

transform	attack	ρ	$\mathbb{E}[c \mid A]$	$\mathbb{E}[c \mid \neg A]$	Δ_c
NormClip	model-scaling	4.19	0.268	0.989	−0.721
NormClip	pixel	1.00	1.000	1.000	0.000
Reputation	model-scaling	8.3×10^5	0.003	1.387	−1.383
Reputation	pixel	1.68	0.915	1.019	−0.105
FoolsGold	model-scaling	1.9×10^{12}	1.497	0.834	+0.663
FoolsGold	pixel	1.8×10^{12}	1.365	0.986	+0.378
RFA	model-scaling	14.8	0.125	1.332	−1.208
RFA	pixel	1.69	0.923	1.012	−0.089
<i>Synthetic families of §5–§4, same code path:</i>					
ladder, $\kappa=0.5/1.0/2.0$	—	2.72/7.39/54.6	—	—	−0.140/−0.240/−0.336
Mode S, all rungs	—	1.00–54.6	—	—	0.000
Mode A, $\nu=-2 \dots +2$	—	7.39–7.39	—	—	−1.138 $\dots$ + 2.469

1498 Three readings, in decreasing confidence. **(1) ρ is not a proxy for Δ_c :** the two are anti-ordered
1499 on the anchors, and NormClip on the pixel backdoor has $\rho = 1$ exactly because no client exceeds
1500 $\tau=5$, so $\Delta_c = 0$ there for a reason the invariance classification cannot see. **(2) Both transforms**
1501 **carrying a missed anchor are targeted far beyond any ladder rung**, which puts the residuals
1502 of §5 on a measured axis rather than a range mismatch; Δ_c does not sort the anchors themselves,
1503 since NormClip carries one that misses and one that lands. **(3) Δ_c is nonetheless not a predictor of**
1504 **suppression:** RFA is strongly negative at −1.208 and rfa→krum fails at 0.779, while Reputation is
1505 more negative still at −1.383 and rep→coord_median suppresses. That is the two-regime structure
1506 of §2 again: strong negative Δ_c can mean either annihilation or mere dilution, and Δ_c ’s magnitude
1507 alone does not say which. We report it as the axis on which the ladder’s anchors were unmatched,
1508 not as the quantity that replaces the statistic.

1509 S Confusion matrix in full

1510 Exactly 5 of the 42 pairs have max-committed ASR < 0.5 : the 3 PASS pairs plus the 2 false
1511 negatives, all of which appear among the representative rows of Table 7. Pooling both in-distribution
1512 tiers (each pair scored against the label committed for it at the time) gives TP = 3, FP = 2, TN = 35,
1513 FN = 2 (38/42; specificity 94.6%, which we report but do not lead with, since 37 of 42 pairs are
1514 negatives and specificity is inflated by that imbalance; precision 1.0 at base rate 11.9% is the $8.4\times$
1515 enrichment that matters, and recall is 40% under the tightened C1, since FG→RFA now counts as a
1516 miss). We report this pooled figure with its decomposition rather than on its own, because the tiers
1517 are not equivalent evidence: it is 16/18 development *fit* plus 22/24 held-out *prediction*. The four
1518 sub-tiers the body refuses to pool are development *fit*, held-out *specificity*, out-of-sample *positives*
1519 *with no informative negatives*, and the 13/20 *prospective* suite. Scoring instead against the criterion’s
1520 mechanism labels, which classify both failed LOW calls as HIGH, gives TP = 3, FP = 0, TN = 37,
1521 FN = 2 (40/42); we take the more conservative pooled figure as primary. The two false positives,
1522 fg→NC (0.927) and fedavg→CM (0.519), were hedged as speculative at pre-registration time and
1523 neither satisfies C1; every pair that actually satisfies C1∧C2∧C3 came out LOW. Category tally
1524 after the Proposition 1 correction: degenerate 18, C1-fail 11, C2-fail 3, C3-fail 7, PASS 3, and all 3
1525 surviving C2 failures are $\cdot \rightarrow$ Reputation, exactly the non-invariant class of Proposition 1(c). With
1526 a positive set of only 2 certified in-distribution pairs and recall 40%, the 94.6% specificity figure

1527 is high partly because the positive class is rare, so we report it as a supporting statistic rather than
 1528 a headline. The conditions are sufficient, not necessary, and they do reject working compositions.
 1529 Extension to 3-way compositions: 9/10 triples correct (too small for confidence).

1530 T Diagnosing the two false negatives in full

1531 NormClip $\rightarrow$ CM (0.433) and RFA $\rightarrow$ CM (0.485) were labelled C3-FAIL on the reading that *any* het-
 1532 erogeneous upstream rescaling compresses a rank aggregator’s margin. Theorem 1(1) says some-
 1533 thing weaker and quantitative: ordering survives when the attack’s coordinate extremeness exceeds
 1534 the realized weight ratio, $S > \rho$. Measured, NormClip’s ratio is mild. Over 9 live rounds it clips
 1535 only 0–2 of 5 clients (11 of 45 client-rounds), so $\approx 76\%$ of clients receive $c_i = 1$ *exactly*, and the
 1536 largest raw norm we observe is 24.5 against $\tau=5$, bounding $\rho \leq 4.9$; against $S=10$ the margin
 1537 $S/\rho \geq 2.0$ still clears 1. RFA-reweighting is likewise near-uniform on benign clients. Both pairs
 1538 therefore sit close to the *uniform* rescaling case, where ordering is preserved exactly, which is why
 1539 their ASR lands in an intermediate band just under threshold rather than in the high band the binary
 1540 label predicts. **The refinement does not recover them.** Both also fail C1: CoordMedian alone is
 1541 already at 0.519, and neither NormClip (0.935) nor RFA (0.885) suppresses, so no reading of C3
 1542 makes either pair certifiable. What the diagnosis buys is a sharper condition and an account of the
 1543 residual: a d_2 whose own baseline sits at 0.519 needs only a mild order-preserving upstream trans-
 1544 form to cross 0.5, so these two are threshold noise around a marginal aggregator, not suppression
 1545 the framework failed to recognise. The operational refinement is to evaluate C3 for rank aggrega-
 1546 tors against the *realized* ρ from a 5–10 round pilot rather than the binary question of whether d_1 is
 1547 heterogeneous: the same quantity Theorem 1(1) already requires, which the labelling had not been
 1548 using.

1549 U Condition-by-condition scoring, and the run-count saving

1550 The results below are §3’s, moved here for space. None changes a reported number.

1551 **The out-of-sample bookkeeping, in full.** All 4 out-of-sample PASS predictions the tightened cri-
 1552 terion licenses held, as did both in-distribution C1 $\wedge$ C2 $\wedge$ C3 pairs. **But all 7 are *positive* predictions**
 1553 **whose FAIL controls are uninformative** (DBA fails even undefended), which is why §3 declines
 1554 to read them as an out-of-sample classification result: five positive predictions against an 11.9%
 1555 base rate (5/42, Appendix S), with the positive direction separately unvalidated (§4).

1556 **Does every condition earn its place? C3 does not, reported as a finding.** Scoring all 42 pairs
 1557 on each condition *independently*: C1 alone selects 4 at 50% precision, C2 alone 14 at 36%, C1 $\wedge$ C2
 1558 selects 2 at 100%, and **C3 alone selects 35 at 9%**, below the 12% base rate, never binding on a
 1559 pair C1 $\wedge$ C2 already admits. So we demote C3 to a margin diagnostic and recommend C1 $\wedge$ C2,
 1560 retaining it on theoretical grounds only: Theorem 1 makes margin compression a real failure mode
 1561 and the $\alpha=0.1$ failure of §3 is exactly a margin collapse. It changes no number: a finding about the
 1562 criterion, not an improvement to it. **One scope limit on the C2 column that we describe as read**
 1563 **off the two defenses’ definitions:** for the conditional classes (coord_median, trimmed_mean, rfa)
 1564 it is not, and c2.provenance in the same generator censuses the difference: 3 of the 14 pairs C2
 1565 admits rest on algebraic invariance and 11 on a per-pair call made in sample, those 11 containing all
 1566 5 low-ASR pairs. *Within this menu* the algebraically settled region holds 3 pairs and no positive, so
 1567 here C2’s own coverage is carried by judgement, not by Prop. 1. **That second half we then removed**
 1568 **by enlarging the region rather than relabelling it:** those 3 pairs are one column of a 3×3 census,
 1569 and closing it prospectively puts a low-ASR pair inside the algebraic region and separates the two
 1570 strategies more sharply there than anywhere in this table (§K.2). What is disclosed and not repaired
 1571 is the *first* half, which needs an invariance result for the conditional classes; we do not have one.

1572 **What the saving is.** Counted in runs rather than pairs (C1’s standalone evaluations are runs too)
 1573 the saving is 420 **unscreened** against 80 **screened**, 81% **fewer composition-evaluation runs**, not
 1574 the 95.2% of *pairs* we previously reported; **nothing is eliminated before any run**, and the structural
 1575 claim is the scaling $|A| S D(D-1)$ (*quadratic* in the defense count) to $|A| S ((D-1) + |\text{admitted}|)$.
 1576 GPU-hours are measured (median 11.1 min/run over 228 logged runs): 63 h saved.

V What the boundary permits, claim by claim

The inventory below separates the three kinds of *no* this paper reports: *structural* (no design satisfying the gate can answer the question, Prop. 4), *open* (answerable in principle, but not by an oracle instrument), and *scope* (outside what the committed-attack setting measures). Nothing here is a new claim; each row points at the section that establishes it.

Table 9: **What this design can and cannot conclude.** The middle column is about *identifiability*, not statistical power: a “no” marked *structural* cannot be repaired by more seeds, defenses or attacks (Proposition 4). Every entry is scoped to the measured setting of §5.

question	identifiable?	answer, and its evidence
Does disturbing d_2 ’s statistic change its <i>decision</i> ?	measured	Yes: Krum’s selection flips in up to 0.80 of rounds
Does it change the <i>admitted adversarial mass</i> ?	by construction	No: unchanged in every one of the 60 measured rounds
Does it change <i>suppression</i> ?	instrumented	No: $ \Delta =0.026$, inside margin (Prop. 2(a))
Does <i>adversary-correlated</i> reweighting change suppression?	instrumented	Yes, non-monotonically: single-peaked (Mode A)
Does admitted mass <i>order</i> outcomes across defenses?	pre-registered	No: refuted: +0.098 against a frozen +0.150
Is (P4) <i>sufficient</i> for (P5)?	control arm	No: <code>cos_krum</code> admits 0.000 while ASR falls
Can (P4) be screened <i>without</i> adversary identity?	no: open	Both instruments are oracles: the stated open problem
Does C1 identify C2’s contribution?	no: structural	Empty cell by the gate, for every menu (Prop. 4)
Does any of this certify <i>deployed</i> security?	no: scope	Committed attack only; one adaptive test breaks it

W Anticipated objections

Five objections we expect, answered against text that is already in the paper rather than by new argument. Nothing here is a new claim; each answer is a pointer plus the scope it turns on.

1. This is small-scale federated learning, not decentralized foundation-model training. Agreed, and we claim **methodological** relevance only (§1): no transformer arm, no foundation-model measurement, and no sentence implying either. What transfers is a statement about *causal structure*: any criterion whose eligibility gate is defined on the same outcome it predicts conditions on a descendant of both mechanism and effectiveness (Prop. 4), and that argument does not depend on model scale or participant count. What does *not* transfer is every empirical magnitude, which is why each is reported per arm, tiered at **very unequal weight and never pooled** (Table 2), and scoped in Table 1. The result forbids an inference; it does not measure a system.

2. The interventions are oracles: no deployed system can pin an adversary’s coefficient or read adversary identity. Correct, and that is why Mode S is an **instrument, not a screen**. The Conclusion says so in those terms: Mode S measures level (P4) only by reading adversary identity, which no deployable defense can, so our experiments **identify the causal channel and are not a (P4) screen**. Pinning $c=1.0$ has one purpose, to cut the confounded path of Figure 1(a) by construction, exactly as an oracle-labelled control isolates a channel in an ablation. The deployable question is left open as the paper’s closing sentence. The one control that *removes* an oracle component (score-only, which scores the transformed stack but aggregates the untransformed selected update) does not change the verdict (App. E, D).

3. $n=5$ seeds is small. It is: $n=5$ per arm, $n=8$ for `cos_krum`, with Student- t 95% intervals throughout and no claim that a small effect is zero. Where suppression does not move we report **no evidence of a practically meaningful change under the pre-specified ± 0.15 equivalence margin**, a margin frozen before the runs (Table 1), never a claim that suppression is unchanged. Two design choices do more for precision than seeds would: the identity rungs are imported **bit-identical per seed** (largest deviation 0.00e+00), so the contrast is within-seed rather than across noise, and the central claim is a **sign reversal** (-0.272 against $+0.098$ on the identical cell, Table 3), which is not a null result. Finally, Prop. 4 says the cell that would validate the criterion is *empty by design*: more seeds cannot repair it (Table 9).

4. If ASR does not move, why does a change of decision matter? Because the change of decision is what a preservation criterion *measures*, and the gap between the two is the finding. An

1613 invariance argument certifies the top three levels of Def. 1 while the security content sits at level
 1614 (P4), the adversarial contribution admitted. A criterion reading the Krum arm would call that com-
 1615 position broken (the decision flips in up to 80% of rounds) for a composition whose suppression
 1616 shows no measurable change: a false alarm in the direction that costs coverage (Prop. 5 caps such a
 1617 screen’s recall at 80%). The converse fails too: `cos_krum`’s selection is bit-identical at every rung
 1618 yet it never suppresses (0.584 at the identity) and its ASR *falls* 0.173 at the first accuracy-gate-
 1619 passing rung, the full-ladder endpoint of 0.425 lying below the accuracy floor, so exact invariance
 1620 of the statistic neither implies suppression nor holds it fixed (App. J.5).

1621 **5. The theory covers only positive per-client rescalings.** Yes, and we say so wherever it is used:
 1622 **a laboratory for studying composition, not a general model of it** (§1; Prop. 1, Thm. 1, App. I).
 1623 The class is not arbitrary (it is what norm clipping, contribution weighting and reputation weighting
 1624 do to a client’s update) but it excludes coordinate-wise filters and any transform that changes an
 1625 update’s direction, and the theorem’s own margin closes under heterogeneity ($m=3.24 \rightarrow 1.01$ at
 1626 $\alpha=0.1$), so **it certifies nothing where heterogeneity is high**. The negative does not rest on the
 1627 theory: it is a *measured* counterexample, and a narrower theory makes the counterexample sharper
 1628 rather than weaker, since **the theorem certifies mechanism preservation, not security**, it bounds
 1629 no ASR.

1630 X Experimental Details

1631 Figure 7 places the evaluated settings in the (γ, p_{eff}) plane. The mix plotted at its top is a Norm-
 1632 Clip/FedAvg menu at realized ASR 0.956, and **we do not rest the compose-rather-than-switch**
 1633 **reading on it**: that policy draws plain FedAvg in 26% of rounds, so a quarter of them are effectively
 1634 undefended and the contrast is not between two defended pipelines. **Nor is it the only mix we ran**,
 1635 and §O.2 reports the other two: FG/CM at $n=5$ and reputation+TrimmedMean at $n=30$, both over
 1636 menus of two real defenses, which are what carry that reading. The figure records the measured
 1637 points and draws no boundary between them.

1638 **FL configuration.** $N=10$ clients, $K=5$ per round, 50 rounds, SGD with momentum 0.9, learning
 1639 rate 0.01, local epochs 3, batch size 32, Dirichlet $\alpha=0.5$ (non-IID). Scale experiments: $N=100$,
 1640 $K=10$, $f=0.05$, ResNet18 with GroupNorm.




Figure 7: **The evaluated settings, plotted. No regimes, and one line, which is a measurement.** X-axis: persistence γ . Y-axis: best single-defense ASR (p_{eff}). Each marker is a setting we ran and carries its own measured value; the dotted line is the measured FL persistence $\hat{\gamma} \approx 0.66$. **An earlier version shaded four regime bands and drew threshold guides at $\gamma=0.15$, $p_{\text{eff}}=0.3$ and $p_{\text{eff}}=0.6$; none of those numbers was measured, so all six are gone.** Five points cannot locate a boundary, and nothing here is evidence of a phase transition in (γ, p_{eff}) space.

1641 **Attacks.** Model-scaling: scale factor 10, 50% poison fraction, target class 0. Pixel backdoor:
 1642 5×5 bottom-right trigger, target class 0. DBA: 4 sub-triggers at image corners, $f=0.4/0.6$. Joint-
 1643 constraint: $\epsilon \in \{0.5, 1.0, 2.0, 4.0\}$, disjoint coordinate blocks per adversary, single-step ϵ -ball pro-
 1644 jection per round.

1645 **Defenses.** NormClip $\tau=5.0$, i.e. $u_i \mapsto u_i \min(1, \tau/\|u_i\|_2)$ over the whole flattened update;
 1646 TrimmedMean $\beta=0.2$ (trim 20% each tail); RFA: unweighted Weiszfeld, 50 iterations, tolerance
 1647 10^{-6} , initialised at the mean; FoolsGold: $w_i \propto \log((1 - s_i)/(s_i + \epsilon) + \epsilon)$ with s_i the client’s
 1648 maximum absolute pairwise cosine similarity normalised by the round maximum, $\epsilon=10^{-5}$, then
 1649 shifted by $-\min_i$ and normalised, so the least-trusted client of each round receives weight exactly
 1650 0 (empirically 2 of 5 clients per round).

1651 **Reputation.** Despite the name this defense is *memoryless*, it keeps no cumulative per-client state
 1652 and is recomputed from scratch each round from that round’s K participants. Consensus is the
 1653 coordinate-wise median m of the K updates; scores are ℓ_2 distances $d_i = \|u_i - m\|_2$; the bandwidth
 1654 is adaptive, $\sigma = \text{median}_i(d_i)$ (clamped $\geq 10^{-6}$); weights are $w_i = \exp(-d_i/\sigma)$, normalised to sum
 1655 to 1. No clipping, no initialisation. Its discriminative property is therefore *consensus distance*, not
 1656 similarity spread, which is why Proposition 1(c) predicts that $\cdot \rightarrow$ Reputation compositions fail C2.
 1657 Weights are strictly positive, so Reputation never produces the zero vectors of Lemma 1.

1658 **Adversary count per round.** Adversaries are the fixed client identities $\{0, \dots, fN - 1\}$; each
 1659 round samples K participants uniformly without replacement from all N , with no stratifica-
 1660 tion. For $N=10$, $f=0.2$, $K=5$ the per-round adversary count is hypergeometric: $P(n_a=0) =$
 1661 $P(n_a=2) = 56/252 = 0.222$, $P(n_a=1) = 140/252 = 0.556$; measured over 100 rounds we ob-
 1662 serve 0.19/0.54/0.27. Two consequences, both matching the data. CoordMedian’s count majority
 1663 $n_a < K/2 = 2.5$ holds for every attainable $n_a \in \{0, 1, 2\}$, i.e. in 100% of rounds. TrimmedMean’s
 1664 budget is $\lfloor \beta K \rfloor = \lfloor 1.0 \rfloor = 1$, so it is exceeded whenever $n_a=2$: 22.2% of rounds, ≈ 11 of 50.
 1665 This is why FG $\rightarrow$ TrimmedMean fails (0.508) while FG $\rightarrow$ CoordMedian succeeds (0.107) under an
 1666 otherwise identical upstream transform. We retain the pre-registered C3-FAIL label for FG $\rightarrow$ TM;
 1667 the count-budget violation is an additional mechanism, not a relabelling.

1668 **Theorem quantities.** FoolsGold weight ratio across 5 seeds $\times$ 20 rounds ($N=10$, $f=0.2$, model-
 1669 scaling): $\max \rho = 2.1$, giving extremeness ratio $S/\rho \geq 4.8$ with $S=10$ (regime A). Reputation,
 1670 same protocol (5 seeds $\times$ 20 rounds, 81 rounds containing both adversarial and benign clients): ρ
 1671 mean 1.44×10^5 , max 2.84×10^6 , p95 6.1×10^5 ; mean adversarial weight 4.83×10^{-4} against mean
 1672 benign weight 0.278, and $w_a < 10^{-3}$ in 85% of rounds. Reputation therefore fails the ordering
 1673 condition by ~ 5 orders of magnitude and suppresses via regime B instead. RFA, measured on
 1674 the transformed points $u'_i = w_i u_i$ that RFA actually receives (25 adversary-present rounds over 3
 1675 seeds): FG zeroes every adversary in 32% of rounds, where Lemma 1 applies and the measured
 1676 contraction $1 - \Lambda_a$ is 0.29 on average (range 0.001–0.52). In the other 68% ($w_a > 0$) the separation
 1677 condition of Theorem 1(2) is satisfied in *every* round, with margin $\Delta'_{\text{sep}}/2(R'_B + \delta') \geq 2.8$ (mean
 1678 3.2) and δ'/R'_B mean 0.19, max 0.72. The two cases are exhaustive, so the analysis covers all
 1679 measured rounds. Proposition 1(a) verified in the shipped implementation: FoolsGold’s weight
 1680 vector changes by at most 3×10^{-8} (float32 rounding) under NormClip across 1000 synthetic trials
 1681 and live rounds in which 1–2 of 5 clients were genuinely clipped (raw norms up to 24.5 against
 1682 $\tau=5$). Exactly 2 of the 5 clients per round receive zero FoolsGold weight. **Corollary 1 in full:**
 1683 with $D_b := R'_B + \delta'$ bounding every benign residual and $D_a := \Delta'_{\text{sep}} - R'_B - \delta' = (2m - 1)D_b$
 1684 lower-bounding every adversarial one, the Weiszfeld coefficients satisfy $\sum_j 1/d_j \geq n_b/D_b$ and
 1685 $1/d_a \leq 1/D_a$, so $\Lambda_a \leq (n_a/D_a)/(n_b/D_b + n_a/D_a) = n_a/(n_b(2m - 1) + n_a)$. At $m \rightarrow 1$ this
 1686 is $n_a/(n_b + n_a)$, i.e. 40% for $n_a=2$, $n_b=3$; at the measured $m \geq 2.82$ it is $\leq 12.6\%$, and at the
 1687 all-round mean $m = 3.24$, $\leq 10.8\%$.

1688 **Pre-registration provenance.** We distinguish the tiers explicitly. The 18-pair development set is
 1689 where the criterion was *developed*; its predictions were not committed in advance and we report it
 1690 as fit. The 24-pair held-out set was pre-registered in `pre_registration_wave2.md` and committed
 1691 before those experiments ran; it contains no predicted-PASS pair, so it bears only on specificity. The
 1692 out-of-sample PASS predictions (DBA transfer, $N=100$ /ResNet18 scale transfer, composition domi-
 1693 nance) were pre-registered separately in `pre_registration_oos_pass.md`, also committed in ad-

1694 vance. The 10 three-way predictions were pre-registered in `pre_registration_three_way.md`.
1695 Two held-out LOW calls, hedged as speculative at pre-registration time, failed; we score
1696 them as misses rather than dropping them. The prospective suite of §5 was pre-registered
1697 in `pre_registration_prospective.md` at commit 262cf35, which precedes the first write
1698 to `results/prospective_suite/`; single-defense baselines for the three new defenses were
1699 measured *before* the freeze, because C1 and C0 take them as inputs (FLTrust 0.048, Krum
1700 0.583, Multi-Krum 0.783 max-committed; FedAvg 0.667 scaling / 0.768 pixel, clearing the
1701 sanity gate). The predicted labels in that file are reproduced verbatim in the PAIRS list of
1702 `run_prospective_suite.py` and the confusion matrix is recomputed from per-seed ASR in
1703 `results/prospective_suite/summary.json`, not transcribed.

1704 **Freeze commits and run tallies for the interventional suites.** Every design,
1705 grid and threshold below was git-committed before the corresponding results direc-
1706 tory existed, and each hash is verifiable in the repository history. *Dose-response*
1707 *ladder* (§5): `pre_registration_dose_response.md` at commit e711a95, before
1708 `results/dose_response/`; 80 runs. *Metric-swap mirrors* (§5): 8 frozen labels at commit
1709 356897b. *Prospective suite* (§5): commit 262cf35, as above. *Targeted intervention, Modes S*
1710 *and A* (§4, §5): design and thresholds at commit 5130cec before `results/targeted_dose/`
1711 existed; 97 runs plus 25 identity runs verified bit-identical per seed (largest deviation 0.00e+0);
1712 the adversarial coefficient share was verified with no ASR involved, before any run. *Rank-*
1713 *discordant replication* (§4): `pre_registration_dose_replication.md` at commit 06d9bce
1714 before `results/dose_replication/` existed; 15 new runs, identity rung imported bit-identical
1715 from the ladder and verified by `--harness-check`. *Second dataset and architecture* (§4):
1716 `pre_registration_dose_femnist.md` at commit 478b555 before `results/dose_femnist/`
1717 existed; 12 new runs, with $\kappa=0$ run in-suite because no EMNIST-byclass ladder exists to
1718 import from. *Score-only control* (§4, Appendix E): `pre_registration_score_only.md` at
1719 commit 35788d9 before `results/score_only/` existed; 15 new runs, identity rung imported
1720 bit-identical and verified by `--harness-check`, and the `cos_krum` corollary checked at one seed
1721 as pre-registered.