
Statistic Preservation Does Not Identify Attack-Suppression Preservation: Limits of Composed Defenses in Collaborative Learning

Anonymous Author(s)

Affiliation

Address

email

Abstract

Collaborative learning composes defenses that were never designed together, so upstream processing can destroy what the downstream defense reads. The natural repair is to require the upstream stage to **preserve** it; for the broad class of *positive per-client rescaling* compositions, where the invariance question is exactly algebraic, that is not enough. The obstruction is structural: a criterion gated on the outcome it predicts conditions on a descendant of both the mechanism and the defense’s standalone strength, so **no screen of this form identifies the mechanism’s contribution, at any sample size**. Ignoring this does not cost power, it costs the sign: on the identical cell the outcome-gated ladder moves -0.272 where the instrument moves $+0.098$. Preservation is a hierarchy, and we break its last two links by intervening rather than comparing: Oracle intervention pinning every adversary’s coefficient closes the attenuation channel, and disturbing the statistic Krum reads then flips its decision in up to 80% of rounds while the *support* of the admitted adversarial mass never changes in any measured round of any arm, and suppression stays inside our pre-specified ± 0.15 equivalence margin. We prove sufficient conditions for the weighting→aggregator class and replicate the negative in a second invariance class and on a second dataset and architecture. The implication for decentralized training is methodological: evaluate composed mechanisms by intervening on adversarial contribution, not by inferring it from downstream statistics.

1 Introduction

Federated learning [McMahan et al., 2017] is vulnerable to data poisoning [Chen et al., 2017, Bagdasaryan et al., 2020], and an arms race has produced defenses [Fung et al., 2020, Pillutla et al., 2022, Yin et al., 2018, Blanchard et al., 2017] with accuracy/robustness trade-offs and none universally dominant (Appendix F). *Can we compose them for stronger robustness?* Composition does not even commute: FoolsGold→CoordMedian suppresses a pixel backdoor to ASR 0.076 while the reverse order reaches 0.732: a rank aggregator emits one vector rather than per-client updates, so it cannot be an *upstream* stage and the pair collapses to FoolsGold alone, bit-identically. Naive layering fails this way whenever upstream processing destroys what the downstream defense reads, and it bites hardest in decentralized pipelines, foundation models included.

Why the question is sharper without a trusted coordinator. Such stacks compose independently developed mechanisms among heterogeneous, partially trusted, often pseudonymous participants, so preservation is the obvious discipline to demand of them. Concretely: a contributor clips its own update, a scheduler reweights contributors by reputation, and whoever merges runs a robust




Figure 1: **The paper in one figure: the statistic moves, admission does not, suppression does not follow.** (a) Two paths run from an upstream rescaling T to suppression: the *chain* a preservation criterion reasons along, and a *confounded* path on which T merely attenuates whichever client carries the poison (Lem. 1) and reaches ASR **without passing through d_2 's statistic at all**. **Mode S cuts the second by construction**, pinning every adversary at $c=1$ (the one channel the dose closes, and not the only thing it moves). The admission count pools the four aggregators of §5; Krum admits no adversary at any rung, so its zero is a floor (App. C). (b) Mode S, Krum under model-scaling (§4): as ρ grows to 54.6 the *decision* flips in up to 80% of rounds while the *support* of the admitted adversarial mass does not, and max-committed ASR stays inside the pre-specified ± 0.15 equivalence margin ($|\Delta| = 0.026$). Error bars in this panel are Student- t 95% intervals on the levels ($n=5$). (c) **Six cells, one contrast each, estimated two ways** (circle: outcome-gated ladder; diamond: Mode S, adversary pinned at $c=1$). **The designs disagree on 4 of 6 cells, and on coord_median/pixel they have opposite signs:** Δ ASR = -0.272 (paired 95% t interval $[-0.422, -0.122]$) against $+0.098$ ($[+0.018, +0.178]$), intervals excluding zero and not overlapping, so **closing the attenuation channel does not shrink the effect, it reverses its sign**. Above the dashed rule a pre-registered mechanism was read off (*training data, not evidence for it*); below are its two out-of-sample cells, one of which **refutes it, so we withdraw it**.

aggregator: three parties, three mechanisms, nobody owning the stack. The only check at each seam is that the next stage still reads what it expects, so a preserved statistic is read as evidence the stack is robust. **That inference is what this paper refutes. The identification problem we isolate is defined independently of model scale**, so we study it in a controlled federated setting rather than claiming transfer to foundation-model training; statistic-preservation checks are therefore not security evidence in larger decentralized stacks, and assuming otherwise in a trustless pipeline is worse than uninformative.

Contributions. We propose no new defense, and **our central contribution is a negative identification result**: an outcome-gated comparison of defense pairs cannot determine whether preserving a downstream defense's mechanism is what keeps the composition secure. Everything below is that result's demonstration or its consequence. *Attack-suppression preservation* names the object throughout: suppression of the committed attack surviving composition, instantiated as ASR, never the statistic or the decision that follows, and we claim nothing about confidentiality, availability, or robustness beyond that suppression; Fig. 1 is the paper in one picture. **The boundary.** Outcome-gated cross-defense screening is structurally unidentified: its eligibility gate is defined on the same outcome it predicts, so no cross-defense contrast *satisfying this outcome-gated design* identifies the mechanism's contribution and the admissible designs are within-defense (§3). **The demonstra-**

tion. Statistic preservation is accordingly not a valid causal proxy for attack-suppression preservation, established by intervention rather than correlation: on the identical cell the outcome-gated design estimates -0.272 where the intervention estimates $+0.098$, reversing the direction and not merely the precision (§4). **The consequence.** The screening rule follows from the boundary and is nothing more: evaluation prioritization, measurably broken under heterogeneity (App. J.2). **Also.** The preservation hierarchy (Def. 1); mechanism preservation as invariance under positive per-client rescalings (Prop. 1), with margins on the weighting→aggregator class (Thm. 1), supporting machinery, safe at *both* extremes of the weight ratio, so dispersion alone mispredicts; counterexamples across four aggregators of three structural kinds; and the negative replicated in a second invariance class and on a second dataset and architecture, while refuting the cross-arm admission ordering we had predicted. Scope is inventoried once, in Tables 1 and 2: what each claim does *not* license, and the evidence tiers (App. J.1).

2 What Preservation Could Mean: A Hierarchy of Five Levels

What does it mean for a composition to preserve a defense? Every per-client defense we consider acts as a positive scalar rescaling $T(u_i) = c_i u_i$, which gives the invariance question a sharp *algebraic* answer (Prop. 1), but writing C2 as an invariance of the statistic S leaves a gap our own experiments exposed. Three different things can be preserved: what the downstream defense *reads* (P1–P2), what it *decides* and what *survives* it (P3–P4), and whether the attack is still *suppressed* (P5). They come apart, so we name the levels and ask of each whether it determines the one below:

Level	What is preserved	Determines the level below?
(P1)	the statistic’s value	yes, by definition
(P2)	the ordering it induces	yes, by definition
(P3)	the decision d_2 makes	no , by counterexample (§4)
(P4)	the adversarial contribution admitted	no , by counterexample (§5)
(P5)	suppression of the committed attack	—

Both breaks are measured, at fixed d_2 and attack: Krum’s decision flips in up to 80% of rounds with the adversarial mass it admits unchanged (§4), and `cos_krum`’s decision and admitted mass are bit-identical while ASR falls 0.173 (§5); Fig. 2 draws the chain and both breaks.

The algebra fixes which level a given composition sits at: cosine similarity is exactly invariant, so C2 holds for every $\cdot \rightarrow$ FoolsGold composition; coordinate ordering survives only under a bounded-ratio condition; consensus distance and residual magnitude are not invariant at all, so $\cdot \rightarrow$ Reputation compositions are C2 failures; norm clipping does not blind FoolsGold, though it does blind Reputation.

Definition 1 (Preservation levels). *For an upstream transform T and a downstream defense d_2 reading statistic S :*

- (P1) **Value invariance:** $S(T(U)) = S(U)$.
- (P2) **Ordering invariance:** $S(T(U)) = \varphi(S(U))$ for some φ preserving the ordering S induces on clients. (this is C2).
- (P3) **Decision invariance:** the client or coordinate d_2 actually selects is unchanged.
- (P4) **Adversarial contribution (admission) invariance:** the adversarial input d_2 admits into its aggregate is unchanged. Admission is the pre-registered term throughout for how much adversarial input survives d_2 : a level, not one measurable quantity, instantiated by whatever each rule exposes: Krum’s and `cos_krum`’s selected-client indicator, `coord_median`’s selected-coordinate share, reputation’s adversarial weight share, which is why App. C’s influence column is aggregator-specific.
- (P5) **Attack-suppression preservation:** ASR stays below threshold.

We target (P3)↯(P4) in §4 and (P4)↯(P5) in §5. The conditions (C0)–(C3) the levels induce (**the rule being $C1 \wedge C2$**) are in App. A and B.

What the theory delivers: supporting machinery, not the claim. For a per-client *weighting* into a *rank-based* or *geometric* aggregator, Thm. 1 gives sufficient conditions for d_2 ’s ordering to survive: it **certifies mechanism preservation, not security** and bounds no ASR. Its one prediction, admitted adversarial mass $\Lambda_a \leq 12.6\%$ against a realized $\leq 1.8\%$, goes vacuous as the separation

margin approaches 1, which heterogeneity reaches, and does not follow from a (d_1, d_2) definition pair alone, as a screen would need (App. I).

3 The Testability Boundary

Why can we not simply compare many defense pairs? The natural test keeps pairs whose constituents already suppress the attack and asks whether the composition preserves it. That is exactly where identification is lost. **The boundary in words** (Prop. 4, formal in App. J.4): let a screen certify $(d_1 \rightarrow d_2)$ only if an *eligibility gate* holds: some constituent already suppresses the attack, measured by the same outcome metric and threshold the screen predicts for the composition, and a mechanism condition P holds, fixed by the two defenses' definitions alone. Four cells follow, and that design never fills two of them:

	Mechanism condition	d_2 suppresses a alone	it does not
	P holds (preserved)	observed: every certified pair	empty by the gate
	$\neg P$ (destroyed)	observed: predicted-FAIL pairs	empty by the gate

Reading down a column is the contrast that would identify P 's contribution, and **the only column in which d_2 is not already doing the work is empty**, for every menu a screen of this form could range over: excluded by design, not undersampled: more data cannot fill a cell the design never permits. The gate cannot be dropped either, because without it C1 is vacuous. So wherever our conditions certify a pair, a low composition ASR is attributable to d_2 's standalone strength rather than to preservation. Equivalently, this is selection on a collider [Elwert and Winship, 2014, Pearl, 2009]: the gate is a descendant of both P and d_2 's standalone effectiveness, so conditioning on it makes them co-vary. The hazard is standard; the gate reading as hygiene rather than as the step that removes identification is not, and we made that mistake ourselves (§5).

The admissible design, derived: the practical contribution. Read forwards the boundary is *constructive*, each clause forced by a term in Prop. 4: (1) hold d_2 and the attack fixed, so the gate's input is the same quantity on both sides of every contrast; (2) intervene on the upstream transform rather than selecting a different d_2 , whose variation is what empties the cell; (3) pin the adversarial coefficient channel, since an upstream that attenuates the adversary moves the outcome by a route unrelated to d_2 's statistic, the confound that inverted our own sign; (4) measure every channel: aggregate displacement, decision, admission, influence. §4 is clause (3) as an instrument; the six-step form is in App. J.3. A second boundary bounds what such a screen can *find*: **no emergent composition can satisfy an inherited-suppression gate**, our own lowest-ASR composition being the witness: FG→RFA reaches 0.093 where FoolsGold alone reaches 0.732, so at $|\mathcal{E}|/|\mathcal{S}|=1/5$ the bound caps recall at 80%, above the 40% we measure (App. J.2).

4 Statistic Preservation Does Not Identify Attack-Suppression Preservation

Does disturbing what a defense reads actually cost suppression? **The claim: it need not, and the design that says otherwise gets the sign backwards.** Our pre-registered pooled hypothesis (H1), that ASR rises with disturbance of the statistic d_2 reads, is rejected, the effect *negative*: that ladder moved adversarial influence and the statistic together (§5). We adjudicate C2's implication that losing (P3) means losing (P5) by oracle intervention on two arms in two invariance classes, not by a general law.

The instrument, as a counterfactual. The clean question asks what happens when the downstream statistic changes but the adversary cannot become weaker. **Mode S** enforces exactly that: it pins every adversary's coefficient at $c=1.0$ and spreads only the benign clients over $\rho = e^{2\kappa}$, holding the adversarial coefficient share constant at 0.266667 to within float32 read-back noise, closing Figure 1(a)'s attenuation path, so a rise in ASR here cannot be attenuation. Mode S and Mode A (§5) are causal instruments built from adversary identity: they measure mechanism-level effects rather than emulate deployable defenses (Table 1, App. J.5).

Two readings of C2, frozen against each other. *The statistic reading*: what matters is disturbance of d_2 's statistic, read off its *decisions*. *The admission reading*: what matters is the *adversarial input*

146 d_2 admits. Mode S disturbs Krum’s selection heavily while leaving the mass it admits untouched,
 147 so that arm adjudicates (Table 4, the pre-registration’s H-statistic and H-admission): as ρ climbs
 148 to 54.6 we find the decision changing in up to 0.80 of rounds while the adversarial mass admitted
 149 changes in none of them, and ASR showing no evidence of the predicted increase, remaining in-
 150 side the pre-specified equivalence margin ($p_{\uparrow}=0.70$, $|\Delta|=0.026 < 0.15$). **The statistic reading is**
 151 **refuted**: statistic and decision disturbance need not change suppression once the adversarial coeffi-
 152 cient share is held fixed. *On this arm the zero is a floor*: Krum admits no adversary at any rung, so
 153 the dissociation is read across all four aggregators, each under its own committed attack:

Aggregator	Kind	Δ dec.	Δ adm.	$\Delta\Lambda_a$	Δ ASR
Krum	selector	0.733	0.000	0.000	−0.026
Reputation	weighted averager	1.000	0.000	0.020	+0.178
Cosine-Krum	cosine-invariant selector	0.000	0.000	0.000	−0.425
Coord. median	coordinate order statistic	0.482	0.000	0.033	+0.098

155 Δ adm. is the *support* of the admitted adversarial mass: it never moves in any of the 60 measured
 156 rounds of any arm, including the 48 per arm where reputation and coord.median carry adversarial
 157 mass at baseline. $\Delta\Lambda_a$ is the separate question of how much weight the adversary’s direction carries.
 158 cos_krum’s −0.425 is the full-ladder endpoint, below our accuracy floor; the witness we rest on is
 159 its first, gate-passing rung (0.173, §5); App. C adds EMNIST-byclass and the score-only control.

160 **What the instrument does not hold fixed, and the pre-registered control that closes it.** Mode S
 161 pins the *adversarial* coefficients and nothing else: the dose also moves Krum’s pairwise distances,
 162 which benign client it selects, and the **magnitude and geometry of the update actually aggre-**
 163 **gated**. That third channel is measured: App. C’s Δ agg. column, the aggregate the defense *emits*,
 164 reads **0.892** here (App. D). **Score-only Krum** closes it: Krum scores the transformed stack and so
 165 makes the same selection, but the *untransformed* selected update is aggregated: the decision channel
 166 preserved, the magnitude and geometry change removed. Rules were frozen at 35788d9 before any
 167 score-only ASR existed, $\Delta > +0.15$ meaning the published result was partly a magnitude artifact.
 168 We find $\Delta = -0.023$ ($0.062 \rightarrow 0.039$) against −0.026 uncontrolled on the identical cell, at the
 169 same Δ dec. of 0.733, the refuting threshold 0.212 nowhere approached and every rung clearing the
 170 accuracy floor, so the central claim is instrumented against two independent confounds, attenuation
 171 and magnitude (App. E).

172 **A pre-registered replication that refutes the positive half, and the sign reversal.**
 173 coord.median under the pixel backdoor is *rank-discordant*: of the interpretable arms it has the
 174 smallest decision change (0.482) and the largest admission-level change ($\Delta\Lambda_a=0.033$), so the two
 175 readings predict opposite rank positions for its rise, and thresholds were frozen against the two pub-
 176 lished rises (confirm if $\Delta > +0.178$, refute if $\Delta < +0.150$). We find $\Delta = +0.098$, which **refutes**
 177 **the admission ordering**: the arm with the largest admission-level change rose *less* than one with a
 178 smaller change. The decision reading does not predict it either ($R^2=0.239$ pooled). What replicates
 179 is the negative on a coordinate-wise order statistic: a decision change of 0.482 costs no suppression
 180 by the frozen margin, at accuracy 0.698. Identity ASR here is 0.443, so equivalence is to the *identity*
 181 *run*; the rise is positive in 5/5 seeds (post-hoc): small and consistent, not zero.

182 **The confounded design gets the causal sign wrong.** On this identical cell (matched displace-
 183 ment, decision change and clean accuracy, Table 3), closing one channel reverses the sign:

Design	The adversarial coefficient	Δ ASR
Outcome-gated ladder (§5)	free to attenuate	−0.272 ($p_{\downarrow}=0.0015$)
Mode S intervention (§4)	pinned at $c=1$	+0.098

185 **The confounded design does not lose power, it gets the sign wrong.** Closing the attenuation chan-
 186 nel inverts the effect rather than sharpening it: a criterion calibrated against the confounded contrast
 187 would have been calibrated against the confound, so §3’s boundary is not a technicality about power.
 188 **Downstream-statistic preservation is not, in general, sufficient to establish attack-suppression**
 189 **preservation** (Prop. 2): no predicate of the statistic alone determines the admitted adversarial mass,
 190 and none of that mass alone determines suppression: the title claim, by counterexample *within this*
 191 *class* (Table 1).

5 What the Boundary Licenses, and What We Found

Does the pattern hold beyond the adjudicating arm, and what does a screen then buy? We validate on CIFAR-10 with a 4-layer CNN, $N=10$ clients, $f=0.2$ adversarial, two committed attacks (model-scaling, pixel backdoor), at $n=5$ seeds per cell ($n=8$ for `cos_krum`, $n=3$ on EMNIST-byclass); ASR means *max-committed ASR*; the rest of the protocol is fixed in App. X. The core evidence is a seed-matched intervention on one cell, not a mean comparison across cells: what §3 excludes is a cell, not a sample size.

The confounded ladder that motivated the instrument, and the sign it got wrong. Our earlier suites tried to validate C2 by *choosing a different downstream defense*, the move Prop. 4 rules out. Varying only the *dose* at fixed d_2 and attack ($\kappa=0$ bit-identical to d_2 standalone) refutes that suite’s pooled prediction that disturbance predicts each arm’s rise: $\rho_s=+0.351$, one-sided $p=0.091$ against a frozen $p < 0.05$, and in three of four arms ASR *falls* as the dose rises (80 runs, one arm confirming decisively; Table 5). **The fault is ours and is now measured:** fixing $\text{mean}(c)=1$ leaves relative weights free, so a large spread mostly *dilutes* whichever client carries the poison: attenuation, not statistic disturbance (App. M, R).

What the instrument does to that ladder, across four aggregators and a second dataset. Most of the one confirmation was attenuation: with the adversary’s share pinned, reputation’s rise at $\rho=54.6$ shrinks from $+0.762$ to $+0.178$, and what remains tracks a small *admission* change, not the large decision change (App. J.5). The class-(a) control refutes us the other way: `cos_krum`’s selection is bit-identical at every rung yet ASR is 0.584 at the identity (it never suppresses) and *falls* 0.173 at the first accuracy-gate-passing rung (one-sided $p=0.044$, 5/8 seeds); the larger fall across the full ladder sits below our own accuracy floor, so **the gate-passing rung is the witness we rest on.** **Mode A** (payload share at fixed dispersion) separates Thm. 1’s two regimes: confirmed for `coord_median`, refuted for reputation. A *pre-registered* replication on EMNIST-byclass/simple_cnn moves the decision on 0.643–0.714 of rounds with the adversary never admitted, suppression inside the pre-specified ± 0.15 equivalence margin: $\Delta = -0.017$ beside CIFAR-10’s -0.026 : **replicated** on a second dataset and architecture, never “generalizes” (App. C, K).

And what a screen then buys. Only the 13/20 prospective sub-tier risks falsification (Table 7), and what limits it is a property of our menu rather than of the screen: FLTrust suppresses both attacks alone (0.048), so its pairs are out of scope by C0 or near-identities. Two pre-registered attempts to validate the *screen* were defeated by that same confound, which makes them evidence for Prop. 4. **Scale:** the pre-registered PASS predictions hold at $N=100/\text{ResNet18}$ but far from converged accuracy, so they bound nothing about converged training. Under *one* adaptive stress test $\text{FG} \rightarrow \text{RFA}$ rises $3.3\times$: **inherited robustness is not adaptive robustness** (App. K, O, S).

6 Conclusion

The finding. Downstream-statistic preservation is not sufficient to identify preserved attack suppression. An intervention changes a defense’s statistic and decision substantially without changing the adversarial mass it admits or the suppression that follows; conversely a bit-identical statistic and decision coexist with a change in suppression, each direction with a measured witness.

Why. An outcome-gated cross-defense comparison conditions on the same outcome it seeks to explain, so it cannot separate mechanism preservation from standalone effectiveness: a structural identification boundary (Prop. 4) no amount of additional data within that design removes. Within-defense intervention fixes d_2 and the attack instead: on the identical cell **the confounded design does not lose power, it gets the sign wrong.**

The implication. Evaluate composed defenses by measuring the adversarial contribution that survives the downstream mechanism, not by inferring security from the statistic it consumes. That object is (P4), which Mode S reads only by knowing adversary identity, so we identify the channel rather than screen for it (App. J.3). The open question: **can (P4) be estimated without knowing who is malicious?**

References

- E. Bagdasaryan, A. Veit, Y. Hua, D. Estrin, and V. Shmatikov. How to backdoor federated learning. In *International Conference on Artificial Intelligence and Statistics*. PMLR, 2020.
- A. N. Bhagoji, S. Chakraborty, P. Mittal, and S. Calo. Analyzing federated learning through an adversarial lens. In *International Conference on Machine Learning*. PMLR, 2019.
- P. Blanchard, E. M. El Mhamdi, R. Guerraoui, and J. Stainer. Machine learning with adversaries: Byzantine tolerant gradient descent. In *Advances in Neural Information Processing Systems*, 2017.
- X. Cao, M. Fang, J. Liu, and N. Z. Gong. FLTrust: Byzantine-robust federated learning via trust bootstrapping. In *Network and Distributed System Security Symposium (NDSS)*, 2021.
- X. Chen, C. Liu, B. Li, K. Lu, and D. Song. Targeted backdoor attacks on deep learning systems using data poisoning. *arXiv preprint arXiv:1712.05526*, 2017.
- F. Croce, M. Andriushchenko, V. Sehwag, E. Debenedetti, N. Flammarion, M. Chiang, P. Mittal, and M. Hein. RobustBench: A standardized adversarial robustness benchmark. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2021.
- F. Elwert and C. Winship. Endogenous selection bias: The problem of conditioning on a collider variable. *Annual Review of Sociology*, 40:31–53, 2014.
- C. Fung, C. J. M. Yoon, and I. Beschastnikh. The limitations of federated learning in sybil settings. In *23rd International Symposium on Research in Attacks, Intrusions and Defenses (RAID)*, pages 301–316, 2020.
- P. Kairouz, H. B. McMahan, B. Avent, A. Bellet, M. Bennis, A. N. Bhagoji, K. Bonawitz, Z. Charles, G. Cormode, R. Cummings, et al. Advances and open problems in federated learning. *Foundations and Trends in Machine Learning*, 14(1–2):1–210, 2021.
- H. B. McMahan, E. Moore, D. Ramage, S. Hampson, and B. A. y. Arcas. Communication-efficient learning of deep networks from decentralized data. In *Proceedings of the 20th International Conference on Artificial Intelligence and Statistics (AISTATS)*, volume 54, pages 1273–1282. PMLR, 2017.
- T. D. Nguyen, P. Rieger, H. Chen, H. Yalame, H. Möllering, H. Fereidooni, S. Marchal, M. Miettinen, A. Mirhoseini, S. Zeitouni, F. Koushanfar, A.-R. Sadeghi, and T. Schneider. FLAME: Taming backdoors in federated learning. In *USENIX Security Symposium*, 2022.
- J. Pearl. *Causality: Models, Reasoning, and Inference*. Cambridge University Press, 2nd edition, 2009.
- K. Pillutla, S. M. Kakade, and Z. Harchaoui. Robust aggregation for federated learning. *IEEE Transactions on Signal Processing*, 70:1142–1154, 2022.
- V. Shejwalkar and A. Houmansadr. Manipulating the byzantine: Optimizing model poisoning attacks and defenses for federated learning. In *Network and Distributed System Security Symposium (NDSS)*, 2021.
- F. Tramèr, A. Kurakin, N. Papernot, I. Goodfellow, D. Boneh, and P. McDaniel. Ensemble adversarial training: Attacks and defenses. In *International Conference on Learning Representations (ICLR)*, 2018.
- H. Wang, K. Sreenivasan, S. Rajput, H. Vishwakarma, S. Agarwal, J.-y. Sohn, K. Lee, and D. Papailiopoulos. Attack of the tails: Yes, you really can backdoor federated learning. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 33, 2020.
- C. Xie, K. Huang, P.-Y. Chen, and B. Li. DBA: Distributed backdoor attacks against federated learning. In *International Conference on Learning Representations*, 2020.

- 286 D. Yin, Y. Chen, K. Ramchandran, and P. L. Bartlett. Byzantine-robust distributed learning: To-
287 wards optimal statistical rates. In *Proceedings of the 35th International Conference on Machine*
288 *Learning*, 2018.
- 289 Z. Zhang, A. Panda, L. Song, Y. Yang, M. W. Mahoney, P. Mittal, K. Ramchandran, and J. E.
290 Gonzalez. Neurotoxin: Durable backdoors in federated learning. In *International Conference on*
291 *Machine Learning (ICML)*, 2022.

292 Table 1 sits ahead of the appendices because it governs all of them: it is the one place each scope
 293 condition is stated, and no later section restates them at each use.

Table 1: **Scope of claims.** The one place each scope condition is stated. **Claims do not transfer upward:** the formal guarantee does not certify the general framing, and the empirical figures certify neither level above them.

kind	content	status, and what it does <i>not</i> license
General	An evaluation gated on the outcome it predicts cannot identify the mechanism’s contribution; admissible designs are within-defense (Prop. 4, Prop. 5).	Proved, and not specific to FL, to backdoors, or to our conditions. Licenses no claim that a composition <i>is</i> secure, nor that mechanism preservation implies suppression.
Formal	Mechanism preservation as <i>invariance</i> under positive per-client rescalings $T(u_i) = c_i u_i$, $c_i > 0$, into specific rank and geometric aggregators (Thm. 1, Lem. 1, Cor. 1).	Proved on that class only: a <i>laboratory</i> for composition, not a model of FL defenses in general. A d_1 that alters participation, carries cross-round state, mixes clients, or aggregates first has no such T and is covered by nothing here (C0, applied retrospectively to arms frozen before it). Bounds no ASR , and Regime A’s condition is not operationally checkable (S is not a stable per-attack constant).
Empirical	§4’s interventions and §5’s 42-pair menu: CIFAR-10, CifarCNN/ResNet18, $N \in \{10, 100\}$, $f \in [0.05, 0.6]$, Dirichlet $\alpha=0.5$, star topology , three attack families; one arm on EMNIST-byclass/simple-cnn.	Measured at $n=5$ seeds per cell ($n=8$ for <code>cos_krum</code> , $n=3$ on EMNIST-byclass), fixed before the runs, so intervals are wide and no null is read as flatness without a pre-registered equivalence margin. Dataset transfer is tested on the flagship arm alone: “replicated on a second dataset and architecture,” never “generalizes”. “Refuted” always means a pre-registered <i>directional</i> hypothesis was rejected against a rule frozen for it, never that a mathematical claim was disproved.
Both instruments are oracles	Mode S and Mode A build c_i from adversary identity, as does Δ_c .	They identify a causal channel and are not screens ; computing (P4) without adversary identity is our open problem, and admission is <i>a direction to measure</i> , not a validated predictor. C1 and all built on it assume a committed adversary: what is bounded is <i>inherited</i> robustness under a fixed attack, never adaptive robustness .

294 A The composition model and the invariance classification

295 **The two witnesses under Definition 1, and which rung each rests on.** (P1) $\Rightarrow$ (P2) $\Rightarrow$ (P3) hold
 296 **by definition: logical implications, not empirical claims;** their content is that they stop at (P3), and
 297 §2’s table states both breaks below it. *Holding* decision *and* admission invariance does not deliver
 298 suppression: `cos_krum` falls 0.173 by the first rung alone through the magnitude channel, at $\rho=2.72$,
 299 where clean accuracy moves only $0.599 \rightarrow 0.567$ and our accuracy gate passes (Appendix N).
 300 Across the whole ladder it falls 0.425, but that endpoint sits below our accuracy floor, so **the gate-**
 301 **passing rung is the witness we rest on** (§5). *Losing* decision invariance does not cost suppression:
 302 Krum’s selection flips in up to 80% of rounds under an intervention pinning the adversarial mass, and
 303 its ASR stays inside the pre-specified equivalence margin ($|\Delta|=0.026 < 0.15$) (§4). So a criterion
 304 stated at (P1)–(P2) constrains security in *neither* direction, **(P4), where the security content sits, is**
 305 **not sufficient for (P5) either, and the hierarchy answers “which level does a security claim have**
 306 **to be stated at?”, not “which level predicts ASR?”**

307 **Definition 2** (Defense composition). *Given defenses d_1 and d_2 , composition ($d_1 \rightarrow d_2$) operates*
 308 *as: d_1 produces per-client transformed updates $U' = \{T_{d_1}(u_i)\}_{i=1}^K$; d_2 aggregates U' using its*
 309 *standard algorithm. For FoolsGold $\rightarrow$ CoordMedian, FG outputs $\{w_i u_i\}$ and CoordMedian takes*
 310 *the unweighted coordinate-wise median of the K transformed values, not a “weighted median”.*
 311 *Clients with $w_i = 0$ are **not removed** but reach d_2 as zero vectors, so at least one point of U' is*
 312 *always the origin, which is what Lemma 1 is about.*

313 Each defense suppresses malicious updates through a *discriminative property* (a structural feature
 314 separating adversarial from benign contributions) and the property is defense-class-specific: *simi-*
 315 *larity spread* (FoolsGold), *coordinate ordering* (TrimmedMean, CoordMedian), *residual magnitude*
 316 (RFA), *consensus distance* (Reputation, in ℓ_2 from the coordinate-wise median) and *norm magni-*
 317 *tude* (NormClip). Mechanism preservation is then governed by whether d_1 ’s transform leaves d_2 ’s
 318 property intact, and because every per-client defense we consider acts as a positive scalar rescaling,
 319 that question has a sharp answer:

320 **Proposition 1** (Invariance classes under positive per-client rescaling). *Let $T(u_i) = c_i u_i$ with $c_i > 0$,*
 321 ***strictly:** the form taken by NormClip ($c_i = \min(1, \tau/\|u_i\|)$), Reputation ($c_i \propto e^{-d_i/s}$) and RFA-*
 322 *reweighting ($c_i \propto 1/\max(\|u_i - \mu'\|, \varepsilon)$), each positive by construction. **FoolsGold is deliberately***

323 *not in that list: its shift-and-clamp gives the least-trusted client weight exactly 0 in every round (2 of*
 324 *5 clients here), so FoolsGold-as- d_1 falls outside the hypothesis and is governed by Lemma 1 rather*
 325 *than by (a)–(c). Then:*

- 326 (a) **Cosine similarity is exactly invariant:** $\cos(c_i u_i, c_j u_j) = \cos(u_i, u_j)$. Hence FoolsGold’s
 327 weight vector is unchanged, $w(T(U)) = w(U)$, and **C2 holds for every $\cdot \rightarrow$ FoolsGold**
 328 **composition.**
- 329 (b) **Coordinate ordering is not invariant** in general; it survives under the bounded-ratio con-
 330 dition of Theorem 1(1).
- 331 (c) **Consensus distance and residual magnitude are not invariant:** heterogeneous c_i can in-
 332 vert their ordering, so $\cdot \rightarrow$ Reputation compositions are C2 failures and $\cdot \rightarrow$ RFA requires
 333 the separation condition of Theorem 1(2).

334 So mechanism preservation is a question about the *invariance class* of the downstream signal, and
 335 the classification corrects a natural but wrong intuition: norm clipping does not blind FoolsGold,
 336 though it does blind Reputation, whose signal is consensus distance rather than similarity.

337 **The framework in general form.** The structure does not depend on rescaling. Write d_2 ’s dis-
 338 criminative property as a statistic $S(U)$ and let T be d_1 ’s per-client transform; then $(d_1 \rightarrow d_2)$ is
 339 **mechanism-preserving** when $S(T(U)) = \varphi(S(U))$ for some φ preserving the ordering d_2 acts on.
 340 C2 is exactly that requirement, C3 asks in addition that φ not compress the gap below d_2 ’s operating
 341 threshold, and Proposition 1 is the case $T(u_i) = c_i u_i$ (Appendix I), so the conditions are alge-
 342 braic rather than experimental. It also states the real limit: **a d_1 whose transform is not per-client**
 343 **(mixing updates, or carrying cross-round state) has no T of this form, and nothing here applies.**

344 **The full causal chain.** Definition 1 is a prefix of a longer one: statistic $\rightarrow$ decision $\rightarrow$ admission $\rightarrow$
 345 aggregate adversarial influence $\rightarrow$ per-round update $\rightarrow$ persistence $\rightarrow$ ASR. **We measure the four**
 346 **emphasized links, not update geometry or persistence dynamics:** an acknowledged gap, not an
 347 assumed-benign one. A selection flipping from one benign client to another disturbs the statistic
 348 maximally and the mechanism not at all, which is why (P3) and (P4) come apart. **“Mechanism”**
 349 **always means the downstream defense’s discriminative operation, never its security behaviour,**
 350 that implication being what this paper refutes.

351 B The signal-preservation conditions in full

352 **Definition 3** (Signal-Preservation Conditions). *Composition $(d_1 \rightarrow d_2)$ is predicted to suppress*
 353 *attack a when: (C0) d_2 ’s mechanism is load-bearing, $\text{ASR}(d_1, a) \geq 0.5$ for some committed a : a*
 354 *condition on the criterion’s domain that **never certifies a pair**, and one we formulated only after*
 355 *§5 made it bind, so it is applied **retrospectively** and is **not part of what the 42 pairs validate**; (C1)*
 356 *some constituent suppresses every committed attack alone, $\min_{d \in \{d_1, d_2\}} \text{ASR}(d, a) < 0.5$ for all*
 357 *a , defined through the same threshold that constitutes the outcome, hence an **empirical prerequisite***
 358 ***and not a mechanistic property**; (C2) d_1 ’s transform does not destroy d_2 ’s discriminative property,*
 359 *decided from definitions alone by Proposition 1; (C3) it does not compress the quantitative margin*
 360 *below d_2 ’s operating threshold. **The prioritization rule is $C1 \wedge C2$, C3 demoted to a margin diag-***
 361 *nostic (Appendix U); the pre-registered labels of §5 were frozen with C3 in the gate and are not*
 362 *relabelled. **Theorem 1 bounds no ASR: that mechanism preservation plus C1 predicts low ASR is a***
 363 *separate empirical hypothesis, tested in §5, and **neither stage alone establishes low ASR, nor does***
 364 ***their conjunction** (stage structure, Appendix B).*

365 Definition 3 in the body states the four conditions and every disclosure attached to them; what
 366 follows is the same criterion with its *stage structure*, which is worth separating because conflating
 367 the stages is the main source of confusion about what the framework claims.

368 Composition $(d_1 \rightarrow d_2)$ suppresses attack a when the conditions below hold. The criterion is a **two-**
 369 **stage screen**, and we label the stages because conflating them is the main source of confusion about
 370 what the framework claims: stage 1 is empirical screening of the *constituents*, stage 2 is mechanistic
 371 screening of the *composition*.

372 **Stage 0: applicability (C0).** Whether the framework has anything to say about this pair at all. This
 373 is a condition on the criterion’s *domain*, and Lemma 2 is stated over it together with C1.

374 **(C0) The downstream mechanism is load-bearing.** $\text{ASR}(d_1, a) \geq 0.5$ for at least one committed
 375 attack a . C2 and C3 concern preservation of d_2 's mechanism, so they carry information only when
 376 that mechanism does work the composition needs: if d_1 alone suppresses everything, the composi-
 377 tion inherits that suppression and a correct prediction is correct for a reason the framework did not
 378 supply. The asymmetry is deliberate: a strong d_2 is the intended case. **Failing C0 removes a pair**
 379 **from the criterion's domain; it never certifies one.** *Provenance:* C0 was vacuous on the menu
 380 the framework was developed on (strongest single defense 0.519) and was formulated only after the
 381 prospective suite of §5 made it bind, so it is applied **retrospectively** and is **not part of what the 42**
 382 **pairs validate** (Appendix Q).

383 **Stage 1: empirical constituent screening** (C1). Read off standalone evaluations. This screens; it
 384 does not explain, and it is not a mechanistic claim.

385 **(C1) Individual effectiveness.** For *every* committed attack a in the evaluated set,
 386 $\min_{d \in \{d_1, d_2\}} \text{ASR}(d, a) < 0.5$. The quantifier matters: our headline metric is the *maximum*
 387 over attacks, so requiring a suppressing constituent for only one attack would license compositions
 388 that fail on the other. C1 is defined through the same ASR threshold that constitutes the outcome,
 389 so it is an **empirical prerequisite, not a mechanistic property**: the one condition that cannot be
 390 checked without running something.

391 **Stage 2: mechanistic composition screening** (C2, with C3 as a diagnostic). Determined by d_1 's
 392 and d_2 's definitions alone, no experiment required, and what Theorem 1 proves. **The evaluation-**
 393 **prioritization criterion is $C1 \wedge C2$,** with C3 retained as a **margin diagnostic** rather than a gate for
 394 the reasons scored in Appendix U. *The pre-registered labels of §5 were frozen with C3 in the gate*
 395 *and are not relabelled,* and the demotion changes no reported number.

396 **(C2) Discriminative-property preservation.** d_1 's transform does not destroy d_2 's discriminative
 397 property, as determined by Proposition 1: automatic for cosine-similarity downstream defenses, con-
 398 ditional for coordinate ordering, and violated for consensus distance under heterogeneous rescaling.

399 **(C3) Margin preservation.** d_1 does not compress the quantitative margin below d_2 's operating
 400 threshold. For rank aggregators: extremeness gap survives reweighting. For geometric aggrega-
 401 tors: separation condition holds on transformed points.

402 **What is proved and what is hypothesised.** Stage 1 is empirical and outcome-defined, stage 2
 403 mechanistic, and **neither alone establishes low ASR, nor does their conjunction**: Theorem 1 es-
 404 tablishes that $C2 \wedge C3$ preserve d_2 's discriminative mechanism on the reweighted inputs and **bounds**
 405 **no ASR**. That mechanism preservation plus C1 predicts *low ASR* is a separate empirical hypothesis,
 406 tested in §5.

407 C The four channels, all four aggregators

408 Endpoint change ($\kappa=0$ against $\kappa=2$) in each measured channel under Mode S, for every
 409 downstream defense in the paper. **The first column, Δ agg., is the relative displace-**
 410 **ment of the aggregate the defense emits,** $\|\text{agg}(T(U)) - \text{agg}(U)\|/\|\text{agg}(U)\|$: what enters
 411 training, and **not** the value of a scoring statistic. The measured channels are therefore ag-
 412 gregate displacement, decision, admission and influence. **Recomputed, not transcribed:**
 413 every cell is emitted by `experiments/build_channel_table.py` from `admission_`
 414 `measurement.json`, `femnist_admission.json`, `targeted_dose/summary.json`, `dose_`
 415 `replication/summary.json` and `dose_femnist/summary.json` under `results/`, which
 416 the script only reads; it also asserts that the identity rung displaces nothing (0.00e+00 for all
 417 six rows). **The sixth row is the score-only control** (Appendix E): the first row's arm with the
 418 magnitude channel closed, so its aggregate column is the re-selection component alone (from
 419 `displacement_decomposition.json`) and its decision, admission and influence columns are
 420 the first row's own numbers, score-only leaving the selection untouched; its ASR comes from
 421 `score_only/summary.json`, which the emitter refuses to merge into the ASR sources it shares
 422 cell keys with, so it cannot shadow the first row's published number. The influence column is
 423 **aggregator-specific by necessity** (a selector has no weights) and its instantiation is: adversarial
 424 share of the Weiszfeld/reputation weight mass for the averager, adversarial share of selected
 425 coordinates for the order statistic, and the indicator that an adversarial update is selected at all for

the two selectors. Earlier drafts labelled this arm “FEMNIST” in the tables while the prose said EMNIST-byclass: the loader reads torchvision EMNIST(split="byclass") under a feminist key and partitions it by the label-Dirichlet split used throughout, *not* LEAF’s by-writer split, so the tables now read EMNIST-byclass while the frozen artifact and pre-registration filenames keep the feminist key. The EMNIST-byclass row is the second-dataset replication of §5; its channels were measured prospectively, before its ASR existed.

Aggregator	Kind	Δ agg.	Δ dec.	Δ adm.	$\Delta\Lambda_a$	Δ ASR
Krum	selector (pairwise distance)	0.892	0.733	0.000	0.000	−0.026
Reputation	weighted averager (cross-round state)	0.659	1.000	0.000	0.020	+0.178
Cosine-Krum	selector (cosine distance)	1.297	0.000	0.000	0.000	−0.425
Coord. median	coordinate-wise order statistic	0.670	0.482	0.000	0.033	+0.098
Krum (EMNIST-byclass)	selector (pairwise distance)	0.842	0.643	0.000	0.000	−0.017
Krum, score-only	selector, magnitude channel closed	0.986	0.733	0.000	0.000	−0.023

What this table does and does not show. It shows that the dissociation between the decision channel and the admission channel is not a property of one aggregator: it holds for a selector, for a cosine-invariant selector, for a coordinate-wise order statistic and for a weighted averager, on two datasets. It does **not** order the arms ($\Delta\Lambda_a$ and Δ ASR disagree in rank, which is the refutation of §4 stated in one row-pair) and it is *not* an out-of-sample test: every row is a Mode-S ladder reported elsewhere in this paper.

D Decomposing the emitted displacement: re-selection against rescaling

The Δ agg. column above mixes two things the causal claim needs apart. For a selector, $\text{agg}(U) = u_{s_0}$ and $\text{agg}(T(U)) = c_{s_1} u_{s_1}$, with s_0 the client selected on the raw stack and s_1 the client selected on the transformed one, so exactly

$$c_{s_1} u_{s_1} - u_{s_0} = \underbrace{(u_{s_1} - u_{s_0})}_{\text{re-selection}} + \underbrace{(c_{s_1} - 1) u_{s_1}}_{\text{rescaling}},$$

both components reported relative to $\|u_{s_0}\|$. They are **not orthogonal** (the measured angle runs from $\cos = -0.99$ to -0.43 , mean ≈ -0.65), so their norms need not sum to the total, and the table prints the angle rather than leaving that discrepancy unexplained. What is asserted instead is the exact vector identity: worst residual $2.66\text{e-}07$ over all 240 (arm, family, rung, seed, round) cells, float32 noise. **This decomposes the published number, not a second similar one:** every recomputed total agrees with the corresponding Δ agg. value in `admission_measurement.json` to $0.00\text{e}+00$, on the same seeds, rounds and configuration. Selectors only: a weighted averager and a coordinate-wise order statistic emit no single selected client, so the split is undefined for them. `experiments/measure_displacement_decomposition.py` trains no ASR, writes a new file, and leaves every frozen artifact byte-identical.

Design	Selector	κ	total	re-selection	rescaling	cos	sel. chg.
Confounded ladder	krum	0	0.000	0.000	0.000	—	0.000
		0.5	0.637	0.685	0.333	−0.65	0.600
		1	0.731	0.851	0.685	−0.66	0.733
		2	0.880	1.675	1.674	−0.72	0.867
Instrument (Mode S)	krum	0	0.000	0.000	0.000	—	0.000
		0.5	0.663	0.646	0.308	−0.64	0.533
		1	0.847	1.075	0.666	−0.69	0.800
		2	0.892	0.986	0.973	−0.66	0.733
Confounded ladder	cos_krum	0	0.000	0.000	0.000	—	0.000
		0.5	0.270	0.000	0.270	—	0.000
		1	0.515	0.000	0.515	—	0.000
		2	0.834	0.000	0.834	—	0.000
Instrument (Mode S)	cos_krum	0	0.000	0.000	0.000	—	0.000
		0.5	0.395	0.000	0.395	—	0.000
		1	0.754	0.000	0.754	—	0.000
		2	1.297	0.000	1.297	—	0.000

454 **What it says, and what it does not.** Mode S into Krum at $\kappa=2$: the published 0.892 is 0.986
455 of re-selection and 0.973 of rescaling, so the dose moves *which* benign client is aggregated and
456 *how large* it arrives, in comparable measure. Of the 11/15 rounds whose selection changes there,
457 **none** lands on an adversary: the quantitative form of the claim that Krum’s decision changes are
458 benign-to-benign, and the reason the support of the admitted adversarial mass changes in **none of**
459 **the 60 measured rounds** while the decision channel moves by 0.733, on this arm no adversary is
460 admitted at *any* rung, so that count is a floor, and the arms carrying adversarial mass at baseline are
461 reputation and `coord_median` (48 of 60 rounds each, App. C). `cos_krum` is the informative contrast:
462 its statistic is exactly invariant to positive rescaling, so 0/15 selections change at every rung of *both*
463 designs and its entire displacement (1.297 under the instrument, 0.834 under the ladder) is rescaling.
464 Its fall therefore cannot travel through the decision channel at all, which corroborates §5’s attribution
465 of that fall to the magnitude channel from a quantity measured independently of any ASR. **This is**
466 **a decomposition of a displacement, not of an ASR effect:** it says what the dose moves in the
467 update entering training and makes no claim about suppression, and it is an analysis of runs already
468 reported rather than a new experiment.

469 E The score-only control in full

470 §4 reports the primary outcome: $\Delta = -0.023$ (mean ASR $0.062 \rightarrow 0.039$, seeds 42–46) against
471 -0.026 for the uncontrolled arm on the identical cell, inside the frozen ± 0.15 equivalence margin,
472 with the accuracy floor cleared at every rung (0.565/0.629/0.624/0.630). The verdict wording is
473 the pre-registration’s own for that branch: *the negative survives with the magnitude channel closed*.

474 **What the arm is.** Krum scores on the transformed stack (statistic and decision disturbed ex-
475 actly as in the uncontrolled arm, since the scoring input is the same) and aggregates the *untrans-*
476 *formed* selected update, so what enters training is bit-for-bit what a client produced. It is the
477 reviewer-suggested control in its own terms: the selected update’s magnitude and direction are
478 held fixed while the selection statistic is perturbed independently. Implementation is a defaulted
479 `score_only=False` kwarg, so no existing call site changes and every frozen computation stays
480 bit-identical; the suite’s own harness check verifies that on the published arms. The $\kappa=0$ rung is
481 *imported* rather than re-run, because at the identity the transform returns the stack unwrapped and
482 score-only Krum *is* Krum alone: an exact claim, verified bit-identically at one seed before the arm
483 started.

484 **Three disclosures, from the freeze rather than from hindsight.** (i) The secondary Jonckheere–
485 Terpstra test was specified to detect a monotone *rise* (the outcome that would have refuted flat-
486 ness with attenuation already closed by construction) and the observed trend is a *fall* ($J=58.5$,
487 $z=-1.110$, $p_{\uparrow}=0.867$, permutation $p_{\uparrow}=0.874$). We report it and do **not** score it as support for
488 flatness; the pre-registration says in advance that at $n=5$ a null here is not evidence of flatness.
489 (ii) The branches had unequal headroom. With the identity rung at 0.062, refuting us required mean
490 ASR($\kappa=2$) > 0.212 against a ceiling of 1 and was fully reachable, whereas the indeterminate branch
491 ($\Delta < -0.15$) was *arithmetically unreachable*: ASR cannot fall below 0, so the largest possible fall
492 on this cell is 0.062. (iii) **Selectors only.** A weighted averager and a coordinate-wise order statistic
493 emit no single selected client, so “the selected update” is undefined for them and the control does not
494 cover them: the `coord_median` arm carrying the $+0.098$ rise is therefore **not** magnitude-controlled,
495 and neither is reputation’s $+0.178$. Score-only Krum is also an instrument and not a defense: it
496 aggregates an update its own scoring stage never saw. And closing the magnitude channel still does
497 not hold the model trajectory fixed, because which client is selected continues to change, by design,
498 that being the channel left open.

499 **The `cos_krum` corollary, and exactly what it settles.** Section 5 of the same pre-registration as-
500 serted, before the run, that score-only `cos_krum` must be *bit-identical* across rungs: its statistic is
501 exactly invariant under positive rescaling, so the selection cannot move, and with the emitted update
502 untransformed the whole trajectory must repeat. It does (accuracy 0.528400 and ASR 0.003111 at
503 all four rungs, $|\Delta| = 0.00e+00$ throughout) and the $\kappa=0$ value reproduces the published cell at that
504 seed exactly, so the two pipelines agree at 0.000000 rather than approximately. **What the check**
505 **does not do is confirm the arm’s published fall.** It runs one seed, as frozen, and that seed is un-
506 representative of the arm: under the uncontrolled dose its own ladder *rises* $+0.224$ where the eight-
507 seed mean *falls* -0.425 . Closing the magnitude channel removes all 0.224 of that seed’s movement,
508 which corroborates the *channel*, for `cos_krum` the emitted displacement is entirely rescaling (Ap-

pendix D), so no other route is open, but the attribution of the arm’s fall rests on that decomposition rather than on this check, and settling it empirically would need the arm’s other seven seeds. Expanding a pre-registered seed count after seeing the data is what pre-registration exists to prevent, so we report the invariance as confirmed and the arm-level attribution as **corroborated mechanically**, not established at the arm mean.

F Related work

The contrast in one sentence. Prior work asks whether a defense withstands an attack. We ask whether an *evaluation criterion* can causally attribute the security of a composition to preservation of the downstream mechanism.

Byzantine-robust aggregation, studied in isolation. Krum [Blanchard et al., 2017], Trimmed Mean [Yin et al., 2018], RFA [Pillutla et al., 2022], FoolsGold [Fung et al., 2020], FLTrust [Cao et al., 2021] and FLAME [Nguyen et al., 2022] are each evaluated as a single stage; we propose no new defense and reason about *ordered pairs* of existing ones. Ensemble adversarial training [Tramèr et al., 2018] *assumes* upstream processing preserves downstream signals; we make that assumption explicit, checkable, and (as it turns out) **the wrong abstraction for security**.

Adaptive attacks are why an inherited-robustness guarantee is not a security guarantee. Optimized model poisoning defeats Krum, trimmed mean and FoolsGold once the attack is chosen *against* the defense [Shejwalkar and Houmansadr, 2021, Bhagoji et al., 2019], as do distributed and durable backdoors [Xie et al., 2020, Wang et al., 2020, Zhang et al., 2022]. C1 quantifies over a *committed* attack, so this literature bounds what any inherited-suppression criterion (ours included) can promise; our single whitebox stress test (§6) measures that gap rather than defending the criterion.

Evaluation methodology, which is where this paper sits. The adjacent precedent is adversarial robustness, where reported numbers proved unreliable until evaluations were required to be adaptive [Croce et al., 2021]. Our point is methodological in the same way but about **identification** rather than attack strength: composition studies gate candidates on standalone effectiveness [Kairouz et al., 2021] and then read the composition’s outcome as evidence about the mechanism, which Proposition 4 shows cannot work at *any* sample size. What is new is the level hierarchy at which the preservation intuition fails and the design class that can test it (Appendix J.4).

The attack side of the same literature engineers *persistence* rather than magnitude: model-scaling [Bagdasaryan et al., 2020] and Neurotoxin [Zhang et al., 2022] both aim at triggers that survive rounds of clean aggregation, which is the quantity we measure as γ (§1) and the reason temporal defense switching does not help in our regime (Appendix O). Nothing in this paper proposes a new aggregator, a new attack, or a new defense; the object of study is the *composition* of defenses that already exist, and the contribution is an identification argument about how such compositions can and cannot be evaluated.

G Notation

symbol	meaning
c_i	upstream coefficient applied to client i ’s update, $T(u_i) = c_i u_i$, $c_i > 0$
A	the set of adversarial participants in a round
ρ	statistic-level dispersion, $\max_i c_i / \min_i c_i$: <i>how far apart</i> the coefficients are
Δ_c	adversary-correlated dispersion, $\mathbb{E}[c_i \mid i \in A] - \mathbb{E}[c_i \mid i \notin A]$: <i>who</i> receives them
r	adversary-to-benign weight ratio, $w_a / \max_b w_b$; the Mode-A dial, $r = e^\nu$
S	the attack’s coordinate extremeness (Theorem 1(1))
m	RFA separation margin, $\Delta'_{\text{sep}}/2(R'_B + \delta')$
Λ_a	aggregate adversarial influence: adversarial share of the Weiszfeld mass
κ	the Mode-S / ladder dispersion dial, $\rho = e^{2\kappa}$
γ	attack <i>persistence</i> , $\text{ASR}(t+1)/\text{ASR}(t)$ after a clean round (≈ 0.66); never a weight

H Deployment guidance in full

Deployment guidance. (0) Check C0: if one candidate defense already suppresses every attack you have, deploy it alone: composition has nothing to add. (1) Check C1: does at least one defense suppress known attacks? (2) Check C2 by invariance class (Proposition 1): a cosine-similarity downstream defense has its signal preserved behind any positive rescaling, a consensus-distance one does not. (3) For weighting→aggregator compositions Theorem 1 gives verifiable sufficient conditions for C2∧C3. Gate on C2; read C3 as a *quantitative diagnostic*: a 5–10 round pilot returns ρ and the separation margin m , and it is $S > \rho$ (rank aggregators) or $m \gg 1$ (RFA) that matters, not whether d_1 happens to be heterogeneous, the binary reading that produced both of our false negatives. (4) Best candidate on our evidence, for further evaluation rather than deployment: **FG→CM** (accuracy 0.68, max ASR 0.107). FG→RFA is marginally lower on the point estimate (0.093 at $n=30$) but statistically indistinguishable, and carries the *highest* adaptive ASR (0.167) and an 18-point accuracy cost, so it is not the more secure choice on our evidence. (5) In our persistent-backdoor setting fixed-order composition substantially outperforms randomized ordering; we have not tested regimes where randomization might pay.

I Proofs

What the two arms jointly establish. Both halves are witnessed by runs already reported, so the following formalizes measurements rather than adding a claim.

Proposition 2 (Statistic-level non-identifiability). *Within the positive per-client rescaling class there exist transform families T_1, T_2 and downstream defenses d_2 such that (a) T_1 and T_2 induce arbitrarily different statistics and decisions at d_2 while inducing the same admitted adversarial mass and the same suppression; and (b) T_1 and T_2 induce the same admitted adversarial mass (indeed the same statistic and the same decision) while inducing different suppression. Consequently no predicate of the downstream statistic alone determines the admitted adversarial mass, and no predicate of the admitted mass alone determines suppression.*

Both parts are proved by exhibited witnesses (Appendix I): Mode S into Krum under model-scaling for (a), `cos_krum` for (b). **The corollary is the paper’s title claim:** no metric defined *solely* on preservation of the downstream statistic can identify attack-suppression equivalence, (a) exhibiting statistic disturbance without suppression change and (b) the converse. **Scope:** a non-existence claim *within the positive-rescaling class on these defenses and attacks*, by exhibited counterexample: **not** that no metric of any kind can order security.

Proposition 3 (ρ and Δ_c are independently controllable). *Within the same class, the statistic-level dispersion $\rho = \max_i c_i / \min_i c_i$ and the adversary-correlated dispersion $\Delta_c = \mathbb{E}[c_i \mid i \in A] - \mathbb{E}[c_i \mid i \notin A]$ are not functionally related: there is a family with $\Delta_c \equiv 0$ exactly while ρ ranges over $[1, 54.6]$, and a family with the statistic’s dispersion held fixed while Δ_c ranges over $[-1.138, +2.469]$. Hence neither quantity determines the other.*

This makes §2’s warning precise: **a transform reported by ρ alone leaves Δ_c entirely free.** It does not follow that Δ_c is a predictor; its magnitude fails to order suppression either (§5). The instrument is elaborate because §3 leaves nothing simpler. §5 answers what remains: *what happened when we tried it the obvious way first?*

I.1 Formal Guarantee: Weighting → Aggregator Class

For the practically important class containing our certified compositions (a per-client *weighting* defense d_1 followed by a *rank-based* or *geometric* aggregator d_2) C2 and C3 can be formally established. **Read what follows as a *boundary* theorem and deliberately not a security theorem:** it says exactly how much an invariance argument about composition can prove, which is preservation of d_2 ’s discriminative mechanism at levels (P1)–(P3). It bounds no ASR, and the counterexamples in the box above are counterexamples to reading it as though it did.

Theorem 1 (Bounded reweighting preserves the downstream discriminative mechanism). *Let d_1 apply per-client scalar weights $w_i \geq 0$ with ratio $\rho := \max_i w_i / \min_{i:w_i>0} w_i$ (over strictly positive weights), operating on n_b benign and n_a adversarial updates in a round of size K . Assume $w_a > 0$ for the adversarial clients; the $w_a = 0$ case is Lemma 1. Then:*

- 597 1. **Rank aggregators.** The composition passes $\{w_i u_i\}$ to an unweighted order statistic over
598 K transformed values per coordinate. The mechanism is preserved when:
- 599 (i) Count majority: $n_a < K/2$.
600 (ii) Signed ordering (C2): on suppression-relevant coordinate k , the adversarial trans-
601 formed value lies strictly beyond every benign one: $w_a u_{a,k} > \max_b w_b u_{b,k}$ (upper
602 tail) or $w_a u_{a,k} < \min_b w_b u_{b,k}$ (lower tail).
- 603 Sufficient magnitude condition for (ii): if $\text{sign}(u_{a,k})$ is common to all adversaries and
604 $|u_{a,k}| > \rho \cdot \max_b |u_{b,k}|$, then (ii) holds. This branch is one of two mechanism-preserving
605 regimes, the other (strong attenuation, $w_a \rightarrow 0$) being Lemma 1.
- 606 2. **Geometric median (RFA):** the mechanism is preserved when separation dominates dis-
607 placement: the minimum adversarial–benign separation among the transformed updates
608 strictly exceeds twice the benign spread plus the displacement the adversaries induce in
609 the geometric median (Appendix I defines the four quantities and proves it), under which
610 every adversarial residual exceeds every benign one, so RFA assigns adversarial updates
611 strictly lower inverse-residual weight.

612 *Proofs are in Appendix I. Both parts are sufficient conditions for d_2 to preserve the downstream*
613 *suppression ordering, not to exclude adversarial contributions outright; **neither is sufficient for low***
614 *ASR, and part (2) constrains the aggregate adversarial coefficient only through Corollary 1.*

615 Theorem 1(2) orders residuals *per client*, which alone does not bound the adversary’s *total* share
616 of the aggregate. A realized *separation margin* $m \geq 1$, the factor by which that condition is met,
617 defined with the bound in Corollary 1 (Appendix I), does bound it, but only quantitatively: as $m \rightarrow 1$
618 the bound degrades to the adversarial count fraction, so **the separation condition on its own is not**
619 **enough**. This is the one place the theory predicts a number and the measurement lands inside it: at
620 the measured $m \geq 2.82$ in every round with $w_a > 0$ (mean 4.71 over those rounds) the bound is
621 $\Lambda_a \leq 12.6\%$, against a realized $\Lambda_a \leq 1.8\%$. The $m \rightarrow 1$ boundary is where §3’s heterogeneity
622 failure sits, and it is reached ($m=1.01$ at $\alpha=0.1$).

623 Λ_a is level (P4) in measurable form. (P4) is the first level whose content is the adversary’s
624 influence on the aggregate, and for a geometric aggregator Λ_a is that level made measurable:
625 Corollary 1 bounds it, Lemma 1 drives it to a benign-only convex combination, and §4’s instru-
626 ments intervene on it directly. **So the level the evidence points to is one the theory can bound,**
627 **just not from a (d_1, d_2) definition pair alone**, which is what a screen would need. Unqualified,
628 “admission” means *admitted adversarial mass*, whether an adversarial update enters d_2 ’s aggregate
629 at all; the three aggregator-specific quantities Definition 1 sets out are never used interchangeably
630 (Appendix G).

631 The remaining case dominates in practice: FoolsGold drives the adversarial weight to zero. This
632 does *not* remove the adversary from RFA’s input (Definition 2): RFA receives a zero vector, a point
633 of the geometric-median problem like any other. The correct statement is stronger than removal:

634 **Lemma 1 (Zero-weight annihilation).** *Our implementation runs the ε -smoothed Weiszfeld iteration,*
635 *whose fixed point $\mu' = \sum_i \lambda_i u'_i$, $\lambda_i \propto 1/\max(\|u'_i - \mu'\|, \varepsilon)$, is a convex combination for every*
636 *configuration, so no non-degeneracy assumption is needed (Appendix I). If $w_a = 0$ for every ad-*
637 *versarial client then $u'_a = 0$ for each, and writing $\Lambda_a := \sum_{a \text{ adv}} \lambda_a$ the adversarial terms vanish*
638 *identically:*

$$\mu' = \sum_{b \text{ benign}} \lambda_b u'_b = (1 - \Lambda_a) \bar{u}, \quad \bar{u} := \sum_b \frac{\lambda_b}{1 - \Lambda_a} u'_b,$$

639 *a convex combination of benign transformed updates. The aggregate is a **contraction of a benign-***
640 ***only weighted average:** u_a does not appear in μ' at all, and the origin point’s only effect is a smaller*
641 *effective step: **a utility cost, not a security failure.***

642 The same argument covers CoordMedian. The two cases are exhaustive over every measured round
643 (FG zeroes all adversaries in 32% of adversary-present rounds, the separation condition holding in
644 the other 68%), and the mechanism survives for *opposite* reasons: under Lemma 1 the adversaries
645 hold most of the Weiszfeld mass ($\Lambda_a = 0.71$) but carry no payload; under Corollary 1 they carry
646 payload but at most 12.6% of it. The contraction $1 - \Lambda_a = 0.29$ is composition’s accuracy cost:
647 FG→RFA reaches accuracy 0.50 and FG→CM 0.66 against 0.78 for the aggregators alone, paid
648 even when the zeroed clients are benign.

649 The proofs below are short; they are collected here rather than inline because what each result *rules*
650 *in or out* is the content, not the derivation. Both witnesses of Proposition 2 are runs reported in §4
651 and §5: Mode S into Krum for part (a), `cos_krum` for part (b).

652 *Proof of Proposition 4.* The load-bearing condition excludes d_1 from attaining the minimum in (i),
653 so certification forces $M(d_2, a) < \theta$; both cells with $M(d_2, a) \geq \theta$ are unoccupied whatever the
654 value of P . A contrast varying d_2 therefore compares pairs differing in P only inside the selected
655 region $M(d_2, a) < \theta$, where P and $M(d_2, a)$ are both functions of the same choice of d_2 ; the
656 missing cell is the one in which they vary independently, and no reweighting of the occupied cells
657 recovers it. The emptiness follows from the gate rather than from $\mathcal{D}$. $\square$

658 *Proof of Lemma 2.* C1 at a^* gives $\min_{d \in \{d_1, d_2\}} \text{ASR}(d, a^*) < 0.5$. By C0 the minimum is not
659 attained at d_1 , so it is attained at d_2 . $\square$

660 *Proof of Proposition 5.* Emergence is the negation of the gate, so each member of $\mathcal{E}$ is rejected while
661 belonging to $\mathcal{S}$: a false negative. The bound is the resulting recall. $\square$

662 *Proof of Proposition 2.* Both halves are exhibited. (a) Mode S into Krum under model-scaling, $\kappa=0$
663 against $\kappa=2$: ρ moves from 1 to 54.6 and Krum’s decision changes in up to 0.80 of rounds, while
664 the support of the admitted adversarial mass changes in none of the 60 measured rounds (a floor
665 on this arm: no adversary is admitted at any rung), the adversarial coefficient share is constant at
666 0.266667, and suppression shows **no evidence of a practically meaningful change under the pre-**
667 **specified ± 0.15 equivalence margin** ($|\Delta|=0.026$). (b) `cos_krum` under the pixel backdoor: by
668 Proposition 1(a) its statistic is invariant, and its decision and admission are measured bit-identical
669 (0.000) at every rung, yet ASR falls 0.173 by the first rung (§5). $\square$

670 *Proof of Proposition 3.* By construction: Mode S pins every adversarial coefficient at $c=1$ and
671 spreads only benign clients, giving $\Delta_c \equiv 0$ at every ρ ; Mode A holds the benign weights uni-
672 form and sweeps only the adversary-to-benign ratio, moving Δ_c at fixed dispersion. Realized values
673 are tabulated in Appendix R, computed with no ASR involved. $\square$

674 *Proof of Proposition 1.* (a) $c_i u_i / \|c_i u_i\| = u_i / \|u_i\|$ for $c_i > 0$, so the matrix of pairwise cosines is
675 unchanged; FoolsGold’s weights are a function of that matrix alone (via each client’s normalized
676 maximum similarity), hence unchanged. (c) With $u_1=(1, 0)$, $u_2=(0, 1)$, $u_3=(3, 3)$ the coordinate-
677 wise median is $(1, 1)$ and u_3 is farthest; under $c = (1, 1, 0.1)$ the median becomes $(0.3, 0.3)$ and u_3
678 becomes the *closest* point, reversing the ordering Reputation depends on. $\square$

679 *Proof of Theorem 1.* (1) CoordMedian operates on $\{w_i u_{i,k}\}_{i=1}^K$ per coordinate. If the adversar-
680 ial values lie strictly beyond all benign values on one side (ii) and $n_a < K/2$, then at most
681 $n_a < \lceil K/2 \rceil$ of the extreme order statistics are adversarial, so the median is a benign value. For
682 the magnitude condition: $w_a > 0$ gives $\rho \geq \max_b w_b / w_a$, so $|u_{a,k}| > \rho \max_b |u_{b,k}|$ implies
683 $|w_a u_{a,k}| > w_a \rho \max_b |u_{b,k}| \geq (\max_b w_b) \max_b |u_{b,k}| \geq \max_b |w_b u_{b,k}|$; with a common sign this
684 places $w_a u_{a,k}$ strictly beyond every benign value in that direction. Model-scaling sets $u_a = S \tilde{u}_a$
685 with $S=10$, which amplifies the adversary’s own direction and so satisfies the sign assumption by
686 construction; with measured $\rho \leq 2.1$ the margin is $S/\rho \geq 4.8$. TrimmedMean additionally needs
687 the trimming budget to cover the adversarial count.

688 (2) Every benign residual satisfies $\|u'_b - \mu'\| \leq R'_B + \delta'$. For adversarial u'_a : $\|u'_a - \mu'_B\| \geq \Delta'_{\text{sep}} - R'_B$,
689 so $\|u'_a - \mu'\| \geq \Delta'_{\text{sep}} - R'_B - \delta'$. Requiring this to exceed $R'_B + \delta'$ gives (i). $\square$

690 Theorem 1(2) orders residuals *per client*, which alone does not bound the adversary’s *total* share of
691 the aggregate: several adversaries could each be individually down-weighted yet collectively hold
692 substantial mass. The separation margin controls this, quantitatively:

693 **Corollary 1** (Bounded aggregate adversarial influence). *Let $m := \Delta'_{\text{sep}} / [2(R'_B + \delta')]$ be the realized*
694 *separation margin, so Theorem 1(2) reads $m > 1$. Benign residuals are at most $D_b := R'_B + \delta'$ and*
695 *adversarial residuals at least $(2m - 1)D_b$, so bounding $\sum_j 1/d_j$ below by the benign terms alone*
696 *gives*

$$\Lambda_a = \sum_a \lambda_a \leq \frac{n_a}{n_b(2m-1) + n_a}.$$

697 *The aggregate is thus a convex combination whose adversarial share is at most Λ_a : writing $\mu' =$*
698 *$(1 - \Lambda_a)\bar{u}_B + \Lambda_a\bar{u}_A$ for the benign and adversarial parts, the displacement from the benign-only*
699 *combination $\bar{u}_B$ is at most $\Lambda_a\|\bar{u}_A - \bar{u}_B\|$.*

700 The bound is only as good as the margin: as $m \rightarrow 1$ it degrades to $n_a/(n_b + n_a)$, approaching
701 40% of the mass at $n_a=2, n_b=3$. We measure $m \geq 2.82$ in every round with $w_a > 0$ (mean 4.71
702 over those rounds), capping the adversarial share at 12.6% for $n_a=2$; the realized Λ_a never exceeds
703 1.8%, and the bound holds in all 17 measured rounds.

704 J Body material, at full length

705 The body compresses the following; nothing here is a new claim, a new number or a new run.
706 Figure 2 summarizes the preservation chain and the two directions in which it breaks.

707 J.1 Contributions, organization and scope in full

708 §1 states these compressed. Nothing here is a new claim.

709 J.1.1 The three evidence tiers

Table 2: The three tiers of evidence, of very unequal weight and never pooled. §1 states them in prose; every empirical claim in the paper is tagged with one of the three words where it is made.

tier	design	what it licenses
Causal	within-defense intervention at fixed d_2 , attack and C1 input: Mode S, Mode A (§4)	statements about the <i>mechanism</i> , scoped to these cells
Replication	EMNIST-byclass, second dataset <i>and</i> architecture ($\Delta=-0.017$), and <code>coord_median</code> , a second invariance class (+0.098)	“replicated on a second dataset and architecture”, never “generalizes”
Screening	the outcome-gated 42-pair suite, 38/42	evaluation prioritization only ; by Proposition 4 it cannot identify the mechanism’s contribution at any sample size

710 J.1.2 The five contributions

- 711 1. **The preservation hierarchy, and the two-directional break above.** Because the per-client
712 defenses in use are all positive scalar rescalings, mechanism preservation is a decidable *invari-*
713 *ance* question (Proposition 1), and on the weighting→aggregator class we prove it with explicit
714 margins (Theorem 1), a **boundary** theorem, not a security theorem. Definition 1 gives the five
715 levels the box states, and §4 measures both breaks. One corollary is worth stating alone: weight
716 heterogeneity is safe at *both* extremes and dangerous in the middle, so **the weight ratio ρ alone**
717 **mispredicts** (Lemma 1, §R.1).
- 718 2. **An identification boundary for outcome-gated screening, and the design class it leaves**
719 **admissible.** Proposition 4 proves the general claim above; its FL instance is that $C0 \wedge C1$
720 force the downstream defense of every certified pair to suppress the attack alone, **mechanism-**
721 **effectiveness confounding** (§3), and its evidential cost is real: a matched ablation built to break
722 the confound had its premise refuted by its own pre-freeze baselines, so **the confound re-**
723 **materialized inside the experiment designed to remove it.** A second boundary limits coverage
724 rather than validation: **emergent** suppression (present in the pair, absent in both constituents)
725 cannot satisfy an inherited-suppression gate, capping such a screen’s recall (Proposition 5), and
726 our own best composition is its witness.
- 727 3. **Three central empirical findings, all causal interventions, not a validation of any screen.** (*i*)
728 *The negative, instrumented and replicated (Causal, then Replication).* Mode S pins every ad-
729 versary’s coefficient, closing Lemma 1’s attenuation channel *by construction*; a **pre-registered**
730 **replication** in a *different* invariance class and on a *second dataset and architecture* reproduces it

while **refuting the cross-arm admission ordering we had predicted** ($R^2=0.239$ over 26 cells), so **admission is a direction to measure, not a quantitative predictor**. (ii) *A sign reversal under a change of intervention (Causal)*. The same cell yields -0.272 ($p_{\downarrow}=0.0015$) with adversarial influence free to co-vary and $+0.098$ with it pinned: the confounded design does not lose power, **it gets the sign wrong**; hence we withdraw the pooled form of C2’s causal claim (80 runs; $\rho_s = +0.351$, $p = 0.091$; §5). (iii) *Heterogeneity closes the theorem’s own margin*. At Dirichlet $\alpha=0.1$ RFA’s all-round mean separation margin collapses from $m=3.24$ to $m=1.01$, the limit at which Corollary 1’s bound goes vacuous, and two of three certified compositions cross the threshold (§3). A fourth predictor falls with them: **adversary-correlated dispersion** Δ_c is independently controllable from ρ (Proposition 3) and locates the two real anchors our synthetic curve misses, but **orders suppression no better**. Statistic disturbance, admission magnitude and coefficient targeting are each refuted *as a predictor*: what survives is a *preservation level*, not a statistic of the coefficient vector.

4. **The screening rule is a consequence of (2), not a claim of its own** (*Screening* throughout). What survives the negative is a conservative *evaluation-prioritization* rule: read C2 off the downstream signal’s invariance class, then evaluate the survivors first. We report it as an application (§3), never as a validated security predictor; it certifies only *inherited* robustness under a *committed* attack, at recall 40%, and outputs an *evaluation ordering*, never a deployment decision.
5. **Secondary results**. Fixed-order composition beats randomized ordering $9\times$, because preservation conditions are order-sensitive; and **one** whitebox composition-aware *adaptive* stress test raises $FG \rightarrow RFA$ to 0.167 ± 0.069 , $3.3\times$ its seed-matched committed 0.051, **an inherited-robustness guarantee is not adaptive robustness** (§6, Appendix O).

J.1.3 How the paper is organized, as an argument

Composition creates an **identification problem**. §2 defines the object to be identified: mechanism preservation as invariance, the five levels, the conditions; §3 characterizes the **boundary** on identifying it and *derives* the class of designs that can; §4 runs the experiment that class dictates; §5 reports what the boundary **falsifies**, our own pooled hypothesis included; §3 states what survives.

J.1.4 Why this class, and the scope discipline

Everything we *prove* is instantiated for positive per-client rescalings, deliberately: the class is a **laboratory for studying composition, not a general model of it**, it covers the per-client defenses in current use (norm clipping, FoolsGold, RFA reweighting, FLTrust) and makes invariance decidable from definitions. **Every claim here is general, formal or empirical, and they do not transfer upward**. Table 1, first in the appendix, is the one place each scope condition is stated (what falls outside the class, what the theorems do and do not bound, the sample sizes, what “refuted” means, why our instruments are not deployable) and we do not restate them at each use.

J.1.5 Threat model and scope

We consider persistent backdoor attacks where an adversary controlling fraction f of clients embeds triggers that survive across rounds, and define *persistence* γ as the fraction of backdoor ASR retained per round of clean aggregation ($\gamma \approx 0.66$ for CIFAR-10/CNN). Attacks include model-scaling [Bagdasaryan et al., 2020], pixel-pattern backdoors [Chen et al., 2017], distributed backdoor attacks [Xie et al., 2020], and adaptive composition-aware adversaries [Shejwalkar and Houmansadr, 2021]. The framework constrains the aggregation pipeline, not the communication graph, and is evaluated only in **star-topology FL**. **A scope condition on the criterion itself, not a limitation of our runs**: C1 and every prediction built on it are defined against a *committed* adversary, so what is bounded is *inherited* robustness under a fixed attack, **never adaptive robustness** (Table 1).

J.2 The prioritization criterion, and where it breaks

This section is a consequence of the boundary, and what survives is an evaluation-prioritization criterion, not a security screen. $C1 \wedge C2$ answers one question: given a menu, which ordered compositions are worth running first. **We disclaim predictive validity explicitly**; it does not predict whether a composition is secure, its positive direction is unvalidated, and the cross-arm admission ordering a predictive reading would need was *refuted* by our own pre-registered replication (§4). It

also costs us our own best composition: under the per-attack quantifier **FG→RFA no longer satisfies C1**, so the lowest-ASR composition we found is a **false negative** and Proposition 5’s witness (uncertifiable by *any* inherited-suppression screen rather than missed by ours) while the certified in-distribution set (FG→CM 0.107, Rep→CM 0.423) is correct and the miss costs nothing detectable (0.093 against 0.107, Welch $p=0.52$). Out of sample it licenses *positive* predictions only, against un-informative FAIL controls, so **this is not an out-of-sample classification result**: the *Screening*-tier 38/42 at recall 2/5 rests on the in-distribution tiers. The defensible claim is *conservative prioritization*; scoring condition by condition, the confusion matrix, the false negatives and what the rule can and cannot do are in Appendices S–T, U and Table 9 (App. V).

The procedure, and where it breaks. (0) If one defense already suppresses every attack you have, stop (C0); (1) require a constituent that suppresses the committed attack (C1); (2) read C2 off the downstream signal’s *invariance class*, which §3 says you cannot test directly; (3) read C3 as a margin from a short pilot, never as a gate. Then **evaluate the survivors first, and evaluate them**: none of this certifies security. **Where the rule breaks is measured rather than argued: under extreme heterogeneity (Dirichlet $\alpha=0.1$) two of three certified compositions cross the threshold and RFA’s all-round mean separation margin collapses from $m=3.24$ to $m=1.01$:** the $m \rightarrow 1$ limit at which Corollary 1’s bound says nothing. Theorem 1’s guarantee is therefore not merely conservative under heterogeneity but *empty* there, and the theorem earns its keep by predicting its own failure region: a short pilot measuring $m \approx 1$ declines to deploy. **Post-hoc, not pre-registered** (Appendix O.1).

802 J.3 What is proved, what is supported, and the six-step procedure

What is proved, what is supported, and what is not. **Proved:** the invariance classes (Prop. 1); sufficient conditions for mechanism preservation on the weighting→aggregator class (Thm. 1, Lem. 1); the identification boundary and the design class it leaves admissible (Prop. 4, Lem. 2); the recall bound on inherited-suppression screens (Prop. 5). **Empirically supported** (*Causal*, then *Replication*): that substantial disturbance of the downstream statistic need not produce a corresponding change in suppression when the adversarial coefficient share is held fixed, on the Krum arm under Mode S and, by pre-registered replication, in a *different* invariance class and on a *second dataset and architecture*; and that the criterion behaves as a conservative prioritizer on 42 pairs. **Not established:** its positive predictive direction ($R^2=0.239$, cross-arm ordering refuted), and generalization beyond the scope Table 1 fixes, $\gamma \approx 0.66$ being architecture-conditional. The criterion **trades discovery for reliability**: it rejects working compositions (recall 40%, one of them ours), cannot in principle do otherwise (Prop. 5), and draws negative-direction evidence only from the in-distribution tiers. FG→CM is an evaluation *ordering*, **not a deployment recommendation** (Appendix H).

What to do instead: six steps for any composition claim. (1) Name the downstream mechanism and its invariance class. (2) State *which* preservation level you claim, (P1)–(P5). (3) Check the upstream transform lies in the class your argument covers: C0 here; one altering participation, carrying cross-round state or mixing clients is covered by no invariance argument. (4) Hold d_2 , the attack, and therefore the gate’s input fixed; **do not select candidates on the outcome**. (5) Intervene on the upstream transform and **measure the channel**: aggregate displacement, decision, admission, influence, not the outcome-conditional contrast. (6) Report the level at which the chain broke. Steps (4)–(5) are what Proposition 4 licenses, and following them turned our -0.272 into $+0.098$; **this, not the criterion, is what we ask readers to carry away**. One boundary it does not cross: under *one* whitebox composition-aware *adaptive* stress test FG→RFA rises to $ASR = 0.167 \pm 0.069$, $3.3\times$ its seed-matched committed 0.051, so **an inherited-robustness guarantee is not adaptive robustness**: the condition the framework most lacks (Appendix O).

828 J.4 The identification boundary, formally

This section states §3’s two propositions and their FL instance at full length; the body gives the boundary in words and the design it dictates.

Proposition 4 (Why constituent-effect contrasts cannot identify composition effects under inherited-effect gating). *Let a screen certify $(d_1 \rightarrow d_2)$ against attack a only if both (i) an **eligibility gate** $E(d_1, d_2, a) := \min_{d \in \{d_1, d_2\}} M(d, a) < \theta$ holds, where M and θ are the same outcome metric*

834 and threshold the screen predicts for the composition, and (ii) a **mechanism condition** $P(d_1, d_2)$
835 holds, determined by the two defenses' definitions alone. Suppose the screen is applied only where
836 the downstream mechanism is load-bearing, $M(d_1, a) \geq \theta$. Consider estimating P 's contribution
837 to $M(d_1 \rightarrow d_2, a)$ from contrasts across a defense menu $\mathcal{D}$. Then every certified pair satisfies
838 $M(d_2, a) < \theta$, so of the four cells of $\{P, \neg P\} \times \{M(d_2, a) < \theta, M(d_2, a) \geq \theta\}$ the certified set
839 occupies at most two, and **the cell needed to separate P from d_2 's standalone effectiveness is empty**
840 **by the gate, for every $\mathcal{D}$** . No contrast varying d_2 identifies P 's contribution. Identification requires
841 holding d_2 , a and therefore E fixed while intervening on the upstream transform: **the admissible**
842 **designs are within-defense**.

843 The structure is a collider: the gate E is a function of constituent effectiveness, effectiveness affects
844 the composition's outcome, and the mechanism condition P is fixed by the same pair of definitions,
845 so conditioning on E makes P and effectiveness co-vary across every certified candidate. The design
846 table is the whole argument, and **the missing counterfactual is a cell, not a sample size**:

	$M(d_2, a) < \theta$ (d_2 suppresses a alone)	$M(d_2, a) \geq \theta$ (it does not)
847 P holds (mechanism preserved)	observed: every certified pair	empty by the gate
$\neg P$ (mechanism destroyed)	observed: predicted-FAIL pairs	empty by the gate

848 Reading down a column contrasts P against $\neg P$ at fixed effectiveness (the comparison that would
849 identify P 's contribution) and the right-hand column, the only one in which d_2 is not already doing
850 the work, is empty for every menu $\mathcal{D}$. **The counterfactual needed is excluded by design, not**
851 **undersampled**. The emptiness follows from the gate, not the menu, and nothing in the argument is
852 specific to FL, to backdoors, or to our conditions (Appendix I). Its instance for *our* criterion is one
853 line:

854 **Lemma 2** (C0 and C1 place the effective constituent downstream). *Let $(d_1 \rightarrow d_2)$ satisfy C0 and C1,*
855 *and let a^* be a committed attack witnessing C0, so $\text{ASR}(d_1, a^*) \geq 0.5$. Then $\text{ASR}(d_2, a^*) < 0.5$:*
856 *the downstream defense suppresses a^* on its own.*

857 Instantiated: in *every* in-scope pair our conditions certify, d_2 already suppresses the attack d_1 fails,
858 so a certified composition's low ASR is always attributable to d_2 's standalone strength (C1) rather
859 than to d_1 having preserved d_2 's mechanism (C2, C3): confounded **by the conditions themselves**,
860 not by an unlucky menu. We name this **mechanism–effectiveness confounding**.

861 **This is selection on the outcome; where it hides is the point.** The hazard is standard (condition-
862 ing on a descendant of both the putative cause and the outcome [Elwert and Winship, 2014, Pearl,
863 2009]) and **we claim no novelty for it**. What is not standard is that in composition studies the gate
864 is presented not as conditioning but as a *reasonableness* filter (“we study defenses that work”), so it
865 reads as hygiene rather than as the step that removes identification. Two consequences are routinely
866 assumed away: **a better estimator cannot fix it**, no reweighting of occupied cells recovering an
867 unoccupied one, and **the gate cannot be dropped**, because without it C1 is vacuous.

868 **Both worked instances are ours, and one of them inverted a sign.** We name no other paper;
869 our own two designs are the instances. The C1-gated study of §5 is the template at the level of the
870 *criterion*; the dose ladder of §5 is a second instance one level down, its gated quantity being the
871 dose's incidental attenuation of whichever client carries the poison. **We ran the outcome-gated**
872 **design first, and on the identical cell it returned the opposite sign** (Table 3): the hazard cost us
873 the sign of our own effect, in the direction that flattered us.

874 **The admissible design, derived: this, not the criterion, is the practical contribution.** Read
875 forwards the boundary is *constructive*: each clause below is forced by a term in Proposition 4, not
876 chosen for convenience. (1) **Hold d_2 and the attack fixed**, so the gate's input $M(d_2, a)$ is numeri-
877 cally the same quantity on both sides of every contrast and cannot co-vary with the mechanism. (2)
878 **Intervene on the upstream transform** instead of selecting a different d_2 off a menu, because it is
879 varying d_2 that empties the identifying cell. (3) **Pin the adversarial coefficient channel**, because
880 an upstream that attenuates the adversary moves the outcome by a route with nothing to do with d_2 's
881 statistic: the confound that inverted our own sign above. (4) **Measure every channel: aggregate**
882 **displacement, decision, admission, influence, rather than the outcome-conditional contrast, the**

outcome being exactly the variable the gate conditioned on. §4 is clause (3) built as an instrument; Appendix J.3 states all four as a six-step procedure for any composition claim.

A second boundary: what such a screen cannot find. The gate that limits what a screen can be *validated* on also limits what it can *cover*, which converts the most obvious objection to our criterion, that its recall is poor, into a scope result.

Proposition 5 (Inherited-suppression screens cannot certify emergent synergy). *Call $(d_1 \rightarrow d_2)$'s suppression of a **emergent** if $M(d_1 \rightarrow d_2, a) < \theta$ while $\min_{d \in \{d_1, d_2\}} M(d, a) \geq \theta$. No emergent composition satisfies the eligibility gate of Proposition 4. Hence, writing $\mathcal{S}$ for the compositions with $M < \theta$ and $\mathcal{E} \subseteq \mathcal{S}$ for the emergent ones, any such screen's recall obeys $\text{recall} \leq 1 - |\mathcal{E}|/|\mathcal{S}|$ on every menu and every attack.*

This is a scope theorem, not a defect found by tuning, and it is specific neither to C1 nor to FL: any screen requiring a constituent to suppress the attack independently inherits it. Our own lowest-ASR composition is the witness (FG→RFA reaches 0.093 where FoolsGold alone reaches 0.732 and RFA alone 0.845), so at $|\mathcal{E}|/|\mathcal{S}| = 1/5$ the bound caps any such screen's recall at 80%, above the 40% we measure (§3). Screening for inherited robustness and searching for synergistic compositions are different problems; only the first is ours.

What the boundary rules out, our own previously-decisive test included. These proofs are short and their value is what they rule in and out. They rule *out* the experiment we had previously named as decisive: a d_2 whose signal is invariant under positive rescaling but *weak in isolation* cannot satisfy C1 in scope, so no such certified pair exists to test. §4–§5 run four members of the class they rule in. **The strongest disturbs the statistic while pinning the adversarial mass, isolating statistic disturbance from adversarial coefficient attenuation, so we built it (§4).**

J.5 The targeted suite in full

The body compresses these paragraphs; nothing here is a new claim or a new run.

Reputation's confirmation, in full. With the adversary's share pinned, reputation's rise at $\rho=54.6$ shrinks from $+0.762$ to $+0.178$ (one-sided Welch $t=-5.22$, $p=0.0017$); what remains tracks a small *admission* change ($0.000 \rightarrow 0.020$) rather than the large decision change ($0.000 \rightarrow 1.000$).

The arc that produced this section. Our pre-registered pooled hypothesis (H1) (ASR should rise with measured disturbance of the statistic the downstream defense reads) was **rejected**, the effect coming out *negative* (-0.272 , $p_1=0.0015$ on the flagship cell), because that ladder moved adversarial influence and the statistic *together* and so never tested H1 at all (§5). This section builds the instrument that fixes it, pinning the adversarial coefficient share while the statistic moves, so statistic disturbance is isolated from adversarial coefficient attenuation. C2 is written at level (P2); if losing (P3) meant losing (P5) the criterion would transfer straight to security, and the experiment below adjudicates that implication by **oracle intervention on two arms in two invariance classes** (*Causal* tier), not by a law about defense compositions in general.

An instrument that closes the attenuation channel by construction. A positive rescaling reaches d_2 through two channels at once: it disturbs the *statistic*, and it *attenuates* whichever client carries the poison (Lemma 1). Separating them needs an upstream that moves the first while pinning the second. **Mode S** does this: it fixes every adversary's coefficient at $c=1.0$ exactly and spreads only the benign clients over $\rho = e^{2\kappa}$, so the **adversarial coefficient share is constant at 0.266667 across all four rungs** (spread 3.5×10^{-7}), and a **rise in ASR here cannot be attenuation**. Design and thresholds were frozen before any run, the share verified with no ASR involved (Appendix X). **Mode A** sweeps the payload instead (§5). **Both are oracle instruments, not defenses and not models:** they read adversary identity, which no deployable defense can, and they do **not** quantitatively reproduce any real defense's behaviour, at matched disturbance two real upstream anchors sit $+0.767$ and $+0.789$ above the synthetic curve (§5), the measured size of that gap. Their purpose is causal identification, not defense evaluation.

Two readings of C2, frozen against each other. *H-statistic:* what matters is disturbance of d_2 's statistic, read off its *decisions*. *H-admission:* what matters is the *adversarial input* d_2 *admits*. On the Krum arm these predict opposite outcomes, since Mode S disturbs Krum's selection heavily

while, the adversary’s weight being pinned, leaving the adversarial mass it admits untouched. That arm adjudicates (Table 4, App. N): as ρ climbs to 54.6, Krum’s **decision changes in up to 0.80 of rounds** while its **admission changes in none of the 60 measured rounds**, and ASR shows **no evidence of the predicted increase, remaining inside the pre-specified equivalence margin** ($p_{\uparrow}=0.70$; $|\Delta|=0.026 < 0.15$). **H-statistic is refuted**: on this arm **substantial statistic and decision disturbance does not produce a corresponding change in suppression once the adversarial coefficient share is held fixed, and the evidence points to admission as the relevant channel in these experiments**.

What the instrument does *not* hold fixed, named in full, and what it pins exactly. Mode S pins the *adversarial* coefficients; it does not pin the benign ones and was never claimed to. So the dose also moves (a) Krum’s pairwise distances, hence the statistic itself; (b) *which* benign client Krum selects, in 0.533/0.800/0.733 of rounds at $\rho=2.72/7.39/54.6$; (c) the **magnitude and geometry of the update actually aggregated**, since the selected client enters scaled by its own coefficient; and (d) the model trajectory, which diverges across rungs in consequence. Channel (c) is **already measured and printed**: the Δ agg. column of Appendix C is $\|\text{agg}(T(U)) - \text{agg}(U)\|/\|\text{agg}(U)\|$, the relative displacement of the aggregate the defense *emits*, and reads **0.892** here. Appendix D splits that 0.892 into 0.986 of *re-selection* and 0.973 of *rescaling*, with none of the re-selections landing on an adversary. So the claim is sharper: **that displacement, with decision changes in up to 80% of rounds, leaves suppression inside the pre-specified equivalence margin, provided the adversarial coefficient share does not move. What is pinned is pinned exactly**: $c_{\text{adv}}=1.0$ at every rung with maximum deviation 0.00e+00 and the share constant to 3.5e−07 on CIFAR-10, 3.9e−07 on EMNIST-byclass, float32 read-back noise, not slack in the instrument (`measure_admission.py`).

Closing the magnitude channel as well: the pre-registered score-only control. Channel (c) is the one confound Mode S leaves open, so we closed it too. **Score-only Krum** scores on the transformed stack exactly as before and aggregates the *untransformed* selected update, so the dose changes *which* client contributes and nothing about *what* it contributes: Δ dec. is the same 0.733 while Δ agg. is no longer 0.892 but the re-selection component alone, 0.986 relative to $\|u_{s_0}\|$. Rules were frozen at 35788d9 before any score-only ASR existed: primary Δ against the same ± 0.15 margin, where $\Delta > +0.15$ would have said **the published result was partly an artifact of the magnitude channel** and narrowed the central claim. **The outcome is $\Delta = -0.023$** (mean ASR 0.062 $\rightarrow$ 0.039) against -0.026 uncontrolled on the identical cell, so **closing the magnitude channel moves the effect by 0.003**; the refuting threshold 0.212 is nowhere approached and every rung clears the accuracy floor. **The central claim is thus instrumented against two independent confounds**, attenuation and magnitude (row six of Appendix C). The secondary trend test, the headroom asymmetry, the selectors-only scope and the `cos_krum` corollary are in Appendix E.

A pre-registered replication, on a different defense class, that refutes the positive half. The two existing Mode-S arms move together on both predictors, so neither adjudicates. `coord_median` under the pixel backdoor is *rank-discordant*: of the interpretable arms it has the **smallest decision change** (0.482) and the **largest admission-level change** ($\Delta\Lambda_a=0.033$), so the two readings predict opposite rank positions for its rise. Thresholds were frozen against the two published rises: confirm if $\Delta > +0.178$, refute if $\Delta < +0.150$ (Appendix X). **The outcome is $\Delta = +0.098$, so the admission ordering is refuted**: the arm with the largest admission-level change rose *less* than one with a smaller change. Nor does the decision reading predict it, having called the *smallest* rise when $+0.098$ is the middle one; the added cells leave the pooled model unconfirmed ($R^2=0.239$, the share coefficient’s sign failing the pre-specified conjunction). **What replicates is the negative**, on a coordinate-wise order statistic rather than Krum’s selector: a decision change of 0.482 costs no suppression by the frozen margin, at accuracy 0.698. Two qualifications: this cell’s identity ASR is 0.443, so equivalence is to the *identity rung*, not a low absolute ASR; and the rise is positive in 5/5 seeds (post-hoc), so the cost is **small and consistent, not zero. The instrument earns its keep**: the confounded ladder on this identical cell *falls* 0.272 ($p_{\downarrow}=0.0015$), so closing the attenuation channel **inverts the sign of the effect**.

What the two arms jointly establish. Both halves are witnessed by runs already reported, so what follows formalizes measurements rather than adding a claim. **Statistic-level non-identifiability** (Proposition 2, proved from the exhibited witnesses in Appendix I): within the positive per-client rescaling class there are transform families inducing arbitrarily different statistics and *decisions* at

the *same* admitted adversarial mass and the same suppression (Mode S into Krum under model-scaling) and families inducing the same statistic, decision *and* admitted mass at *different* suppression: `cos_krum`. Consequently **no predicate of the downstream statistic alone determines the admitted adversarial mass, and no predicate of the admitted mass alone determines suppression. That corollary is the paper’s title claim:** no metric defined *solely* on preservation of the downstream statistic can identify attack-suppression equivalence. **Scope:** a non-existence claim *within this class on these defenses and attacks*, by exhibited counterexample: **not** that no metric of any kind can order security. A second proposition sharpens §2’s warning: statistic-level dispersion ρ and adversary-correlated dispersion Δ_c are **independently controllable** ($\Delta_c \equiv 0$ while ρ ranges over $[1, 54.6]$; Δ_c over $[-1.138, +2.469]$ at fixed dispersion), so **a transform reported by ρ alone leaves Δ_c entirely free** (Proposition 3), which does *not* make Δ_c a predictor, its magnitude failing to order suppression either (§5). The instrument is elaborate because §3 leaves nothing simpler. §5 answers what remains: *what happened when we tried it the obvious way first?*

K What the boundary licenses: the suites in full

This section holds §5’s three suites at full length: the two attempts the identification boundary defeated, the confounded dose ladder that motivated Mode S, and the remaining targeted arms. Nothing here is rescored; the body states the outcomes.

K.1 Two attempts that the confound defeated

Both were built to validate the *screen*; neither could, so they are evidence for Proposition 4 rather than for the framework.

Two pre-registered suites, and the two ways the boundary asserted itself. (i) *A prospective suite on three unused defenses: 13/20, and no positive test.* Everything scored in §3 was developed on the menu it is scored on, so we ran a prospective test on FLTrust [Cao et al., 2021], Krum and Multi-Krum [Blanchard et al., 2017], for which the theory predicts *opposite* outcomes. All 20 labels were git-committed before the output directory existed and both attacks clear the FedAvg sanity gate, so the negatives are informative (11/13 in scope, 5/7 on the falsifiable arm; Appendix Q). **Two disclosures dominate it:** FLTrust alone suppresses both attacks (0.048), so its 7 pairs are **out of scope** by C0 (applied *retrospectively*) and its 6 in-scope predicted-LOW pairs are **near-identities**, four bit-identical to FLTrust alone, so **the suite supplies no falsifiable prospective positive test.** (ii) *A matched metric-swap ablation, whose premise its own baselines refuted.* To break the confound by construction we held the suppression *principle* fixed and varied only the invariance class, mirroring Reputation and Krum scoring from *directions* alone. Measuring the C1 input before freezing **refuted the design’s premise:** `cos_reputation` collapses (0.017 $\rightarrow$ 0.982) and `cos_krum` reaches 0.078 only at 0.146 accuracy, failing the C1 gate applied everywhere else (**magnitude information is what lets this family handle a magnitude attack**), so C1 bound for all eight pairs, every label was HIGH, and the suite bears on specificity only (8/8; App. L, P). **The confound re-materialized inside the experiment built to remove it.**

The one contrast that escapes, and the design it dictates. One comparison evades the boundary: hold d_2 *and* the attack fixed and vary only the upstream, so C1 is identical by construction. On *Krum*/model-scaling, NormClip disturbs its selection 0/9 and reaches 0.062 while RFA disturbs it 6/9 and reaches 0.779 (Welch $p=0.012$, $n=5$), with no overlap across the six cells where standalone d_2 genuinely suppresses. **This is post-hoc** (Appendix L) but has what nothing above it does: C1 held exactly fixed, disturbance measured on the mechanism rather than read off the outcome.

K.2 H1, and the confounded design that motivated the instrument

Both suites above tried to validate C2 by *choosing a different downstream defense*: the move Proposition 4 rules out. The admissible alternative varies only the *dose* at fixed d_2 and attack, so the C1 input is the same quantity at every rung and the confound is **impossible by construction** rather than matched away.

1037 **The instrument, the abscissa, and the predictions frozen against them.** The synthetic d_1 has
1038 one parameter, the *dispersion* of a positive per-client rescaling: $c_j \propto \exp(\kappa(2j/(K-1) - 1))$ at
1039 position j of a (seed, round)-keyed permutation, normalized so $\text{mean}_j(c_j) = 1$. It is Proposition 1’s
1040 form, the **dose** is $\rho = e^{2\kappa}$, $\kappa=0$ is bit-identical to d_2 standalone, and normalizing the mean while
1041 keying the permutation on (seed, round) holds aggregate scale fixed and stops the dose parking ad-
1042 versaries at a fixed end: neither of which, as it turns out, pins Δ_c . Frozen grid $\kappa \in \{0, 0.5, 1.0, 2.0\}$,
1043 four cells across three invariance classes. What matters to C2 is not ρ but what the rescaling does to
1044 the statistic d_2 reads, so we measured that first, with no ASR involved, establishing that *there is a*
1045 *dose* (0 at identity, rising to 0.87/0.93 for class (c)) and that *class-(a) invariance holds in code*. *Pri-*
1046 *mary, pooled:* across all 16 (arm, κ) cells the ASR rise over that arm’s own $\kappa=0$ rung is positively
1047 rank-correlated with measured disturbance (Spearman, one-sided, $\alpha=0.05$). *Secondary, per arm:*
1048 monotone rise for class (c), **FLAT** (the frozen label, scored as equivalence) for class (a), shallower
1049 for class (b); one-sided Jonckheere–Terpstra at $n=5$, “flat” scored as *equivalence* ($|\Delta| < 0.15$, every
1050 rung below 0.5), a margin fixed in advance. **The ladder was falsifiable both ways:** `cos_krum`’s
1051 aggregate is still $c_j u_j$ for the selected j , so suppression can be lost through the magnitude channel,
1052 and a class-(c) arm inside the margin would refute the central fragility claim. **No frozen prediction**
1053 **uses S :** measured, S is **not a stable per-attack constant** (medians 0.57/0.92, range to 12.59), a
1054 limit on the theorem’s checkability. Design, grid, arms, shapes and refutation criteria were frozen in
1055 advance (Appendix M).

1056 **Outcome: one arm confirmed, two refuted, one indeterminate; the pooled test fails.** All 80
1057 runs completed and were scored per seed against the frozen rules (Table 5, App. N). The **primary**
1058 pooled prediction is **refuted:** Spearman $\rho_s = +0.351$ over 16 cells, one-sided $p = 0.091$, against
1059 a frozen $\rho_s > 0$ at $p < 0.05$; **in three of four arms ASR falls as the dose rises.** Arm 2 confirms
1060 decisively: reputation $0.017 \rightarrow 0.779$ ($p = 1.7 \times 10^{-6}$) at nearly constant accuracy, the paper’s
1061 only prospective positive for C2, and it establishes the *destruction* half only. Arm 1 refutes more
1062 than its own shape: Krum, whose statistic the fragility claim is stated over, shows no rise (inside
1063 the equivalence margin) despite measured disturbance 0.867, and does not reproduce our within-
1064 defense contrast ($0.062 \rightarrow 0.779$ under RFA at a *lower* disturbance), so selection-flip rate cannot be
1065 its mechanism. Arm 3 is indeterminate and uninterpretable anyway (`cos_krum/pixel falls` $0.490 \rightarrow$
1066 0.086 , the last rung below the accuracy gate) and arm 4 is refuted outright (`coord_median/pixel`
1067 falling $0.443 \rightarrow 0.170$ at accuracy ≥ 0.69). Both fell where a rise was frozen, which is why neither
1068 is scored in our favour.

1069 **Two faults, both ours, and now measured rather than conjectured.** *The instrument.* Normal-
1070 izing $\text{mean}(c)=1$ fixes the aggregate’s scale but not any client’s *relative* weight, and because the
1071 dose ignores adversary status a $55 \times$ spread mostly *dilutes* whichever client carries the poison: on
1072 the same rounds the ladder’s Δ_c drifts monotonically negative with the dose ($0.000 \rightarrow -0.140 \rightarrow$
1073 $-0.240 \rightarrow -0.336$), so three arms measure attenuation rather than C2 disturbance. **The adversary-**
1074 **targeted dose is the fix:** Mode S is the Δ_c -**pinned** instrument (Δ_c exactly 0 at every rung, ρ up to
1075 54.6) and Mode A the Δ_c -**swept** one. *The abscissa.* At matched disturbance two of six *real* upstream
1076 anchors miss the synthetic curve badly (RFA $\rightarrow$ Krum $+0.767$, NormClip $\rightarrow$ Reputation $+0.789$), and
1077 not through a range mismatch; Δ_c locates the discrepancy, as Proposition 3 says it may, those two
1078 being the most extremely targeted of the four real transforms ($-1.21, -0.72$) at ρ of only 14.8 and
1079 4.19 against the ladder’s 54.6: **matched on the statistic, badly unmatched on the targeting. But**
1080 **this rescues no predictor:** Δ_c ’s magnitude does not order suppression either, and FoolsGold’s is
1081 *positive: post-hoc and exploratory*, computed with no ASR and rescoring nothing (Appendices N,
1082 R). Two checks did pass: the $\kappa=0$ rung reproduces every standalone baseline exactly per seed, and
1083 `cos_krum`’s selection never moved (0/120).

1084 K.3 Mode A, the four aggregators, and the pre-registered second dataset

1085 **Mode A: the two regimes, and where they hold.** **Mode A (payload-only)** closes the other chan-
1086 nel: it keeps the benign weights uniform and sets the adversary-to-benign weight ratio to $r = e^\nu$
1087 at fixed mean, so the adversarial share sweeps $0.049 \rightarrow 0.710$ with the statistic’s dispersion held
1088 fixed: the mirror of Mode S. That gives Theorem 1 a single falsifiable shape: ASR single-peaked
1089 at $r=1$, since an over-attenuated adversary carries no payload (Regime B) and an over-amplified
1090 one is discarded by a rank aggregator (Regime A). It is **confirmed for** `coord_median`, the rank
1091 aggregator the regimes are stated for: ASR rises $0.124 \rightarrow 0.315 \rightarrow 0.443$ into $r=1$ and falls

1092 0.443 $\rightarrow$ 0.296 $\rightarrow$ 0.098 past it, both sides significant ($p_{\downarrow}=3.0\times 10^{-4}$, 6.4×10^{-4}) at accuracy
1093 ≥ 0.74 (Figure 3). It is **refuted for reputation**, which is monotone, maximal at the most attenuated
1094 rung (0.894 at $r=0.135$), because an attenuated adversary reads as *consistent* to a reputation score.
1095 Reputation is not a rank aggregator, so the shape was never licensed for it.

1096 **The same channels, across four aggregators of three structural kinds.** Across selectors on pair-
1097 wise and on cosine distance (`krum`, `cos_krum`), a coordinate-wise order statistic (`coord_median`)
1098 and a weighted averager with cross-round state (`reputation`), the negative is **not Krum-specific**: the
1099 decision channel moves by 0.482–1.000 on three of the four while **the support of the admitted ad-**
1100 **versarial mass changes in none of the 60 measured rounds of any of the four**, and on reputation
1101 and `coord_median` that mass is *present* at baseline in 48 of those 60, so their counts are not floors
1102 the way Krum’s is, and the only two arms whose ASR *rises* are the only two with nonzero $\Delta\Lambda_a$. **An**
1103 **analysis of runs already reported, not a new experiment, and it orders nothing**: the larger $\Delta\Lambda_a$
1104 (0.033) belongs to the arm with the *smaller* rise (Appendix C).

1105 **A second dataset and architecture, pre-registered.** The same cell on EMNIST-
1106 byclass/simple_cnn in place of CIFAR-10/cifar_cnn, frozen in git before any EMNIST-byclass
1107 ASR existed. Premise measured first, no ASR anywhere: the decision moves on 0.643–0.714 of
1108 rounds, the support of the admitted mass unchanged in all 56 measured rounds (again a floor: no
1109 adversary admitted at any rung). **Suppression shows no evidence of a practically meaningful**
1110 **change**: $\Delta = -0.017$ against the frozen ± 0.15 margin, beside CIFAR-10’s -0.026 , at accuracy
1111 $0.767 \rightarrow 0.835$, so the gate never binds. **Replicated on a second dataset and architecture**: the
1112 frozen wording, never “generalizes”. Two limits from the freeze: the Jonckheere–Terpstra was
1113 specified to catch a monotone *rise* and the trend is a *fall* ($z=-1.563$), so it is not read as support for
1114 flatness; and with the identity rung at 0.027 the refuting branch was reachable (> 0.177) while the
1115 indeterminate one was not, so equivalence is to the identity rung, not a low absolute ASR.

1116 It addresses neither recall nor discovery, and nothing in §5 is rescored.

1117 L Verification of the invariance mirrors

1118 All eight metric-swap labels were frozen in advance, and these mirrors are mechanism-
1119 isolating ablations on the same footing as the instruments of §4. We check the statis-
1120 tic directly rather than inferring invariance from an outcome (`verify_cos_invariance.py`,
1121 `results/cos_invariance_check.json`), comparing each d_2 ’s internal state on $\{u_i\}$ against
1122 $\{c_i u_i\}$ for the c_i NormClip and RFA actually produce. Over 18 live rounds (3 seeds $\times$ 2 up-
1123 streams $\times$ 3 rounds) `cos_reputation`’s weight vector matches to $\leq 2.5 \times 10^{-7}$ with its ordering
1124 never changing (0/18) and `cos_krum` selects the same client in 18/18; the Euclidean originals do
1125 not (Reputation’s weights move by up to 0.222, ordering changes 12/18; Krum’s selection changes
1126 6/18). Disturbance tracks the realized weight ratio ρ as Proposition 1 predicts: under NormClip
1127 ($\rho \in [1.00, 4.90]$) Reputation changes in 3/9 and Krum in 0/9, under RFA ($\rho \in [2.11, 17.38]$)
1128 in 9/9 and 6/9. On 1000 synthetic draws the separation is cleaner still (0/1000 vs. 992/1000
1129 for the Reputations, 0/1000 vs. 618/1000 for the Krums), and the mirrors match the shipped im-
1130 plementations they stand in for to $\approx 10^{-8}$ relative aggregate error. **The post-hoc within-defense**
1131 **contrast of §5 re-uses this frozen suite**, so it is not pre-registered, and two seeds were added after
1132 seeing the $n=3$ result: a top-up that cut both ways, degrading the invariant control’s Δ from ≈ 0
1133 to $+0.19$ while leaving the disturbed/undisturbed split without overlap. One asymmetry matters:
1134 `cos_reputation`’s *weights* are exactly invariant but its aggregate is not, since the weighted aver-
1135 age is over the rescaled $c_i u_i$: a magnitude contraction of the Lemma 1 kind, so the prediction is
1136 two-sided, with ASR staying low while clean accuracy pays a contraction cost.

1137 M Dose–response caveats recorded before the freeze

1138 **The pre-freeze disturbance measurement.** The abscissa was measured over 5 seeds $\times$ 3 rounds
1139 per cell with no ASR involved, which is why it precedes the freeze: the class-(a) mirror’s cosine
1140 scores move by at most 1.1×10^{-4} in 0/120 cells, i.e. the invariance holds in the shipped code
1141 and not merely on paper, while the class-(c) arms reach 0.87 and 0.93. **Attack assignment.** The
1142 class-(c) arms are on model-scaling and the class-(a)/(b) arms on the pixel backdoor: not a free

choice, since on model-scaling *no* invariant d_2 suppresses at usable accuracy (§5). “Same attack” is required *within* each ladder, which is where the causal claim lives; across arms we compare curve shapes and slopes, never levels. `coord_median/pixel` sits only 0.057 below threshold, so its binary above/below-0.5 reading is weak either way. The dose is uncorrelated with adversary status by construction but not adversary-*adversarial*; a dose that deliberately up-weights adversaries would be strictly harder and is not run. A synthetic transform is not a real defense, so we overlay the two *real* upstreams already measured on these cells (NormClip, RFA) as anchors at their own measured disturbance; landing far off the synthetic curve would be a finding about the instrument and reported as one. Design, grid, arms, shapes, margin and refutation criteria are frozen in `pre_registration_dose_response.md` at commit `e711a95`, timestamped before `results/dose_response/` existed and reproduced verbatim in the runner’s arm table. **What it does not address:** recall, and discovery: `fg→rfa` remains a false negative of our own rule.

N Dose–response ladder: per-seed detail, intervals, and the anchor residuals

Table 3: **The confounded ladder against the Mode-S instrument, on the one arm that has both.** Same downstream defense, same attack, same seeds, same nominal dose, and the two designs share their $\kappa=0$ cell *numerically* (mean and worst per-seed gap 0.00e+00, asserted at 10^{-12}). At the top rung they are matched on aggregate displacement, decision change and mean clean accuracy, so **the ladder’s fall is not an accuracy collapse**; they differ chiefly in the **adversarial coefficient channel**, and the ASR effect **changes sign**. **The one non-match, disclosed:** aggregate adversarial influence Λ_a moves in *both* designs, 0.081 against 0.033 ($2.4\times$), so neither column is a clean statistic-only intervention, which is why the contrast is “differ chiefly in the coefficient channel”, never “only in it”. *Vocabulary:* this row is the $\Delta\Lambda_a$ column of the channel table (the adversarial share of selected coordinates, the same field in the same artifact) and **not** its Δ adm. column, which counts rounds where the *support* of the adversarial mass changes and reads 0.000 on every row: what moves here is how much weight the adversary’s direction carries, not whether it enters at all. A two-column contrast on one arm: it locates where the two designs differ and attaches no p -value to the difference between the columns. Every cell recomputed from per-seed and per-round rows by `build_comparability_table.py`.

	Confounded ladder dose_kappa	Instrument (Mode S) doseS_kappa
downstream d_2	<code>coord_median</code>	<code>coord_median</code>
attack	committed pixel	committed pixel
seeds	42–46	42–46
nominal dose	κ 0.0 $\rightarrow$ 2.0, ρ 1 $\rightarrow$ 54.6	κ 0.0 $\rightarrow$ 2.0, ρ 1 $\rightarrow$ 54.6
benign coefficients	mean 1 over <i>all</i> clients	mean 1 over <i>benign</i> clients
adversarial coefficients	drawn from the same spread	pinned $c = 1.0$ exactly
adv. coefficient share	0.267 $\rightarrow$ 0.187	0.266667, constant
Δ_c	+0.000 $\rightarrow$ −0.336	0 by construction (read back 9e−08)
aggregate displacement	0.000 $\rightarrow$ 0.722	0.000 $\rightarrow$ 0.670
decision change	0.000 $\rightarrow$ 0.487	0.000 $\rightarrow$ 0.482
adversarial influence $\Delta\Lambda_a$	0.000 $\rightarrow$ 0.081	0.000 $\rightarrow$ 0.033
mean clean accuracy	0.767 $\rightarrow$ 0.690	0.767 $\rightarrow$ 0.698
ASR	0.443 $\rightarrow$ 0.170	0.443 $\rightarrow$ 0.541
ASR effect	−0.272 ($p_\downarrow = 0.0015$)	+0.098 ($p_\uparrow = 0.069$)

Table 6 is every cell of the ladder; Figure 4 plots it on the measured axis and Figure 5 on the nominal one. Both are produced by `analyze_dose_response.py` and `dose_response.figure.py` from `results/dose_response/`; nothing below is transcribed from a log.

The replication arm’s post-hoc seed-matched difference. The frozen primary for the `coord_median` replication (§4) is the *unpaired* $\Delta = +0.098$, and it is not superseded by what follows. Between-seed variance on this cell is large (identity ASR spread 0.287–0.677), which the unpaired test cannot see through, so we also report the seed-matched difference: positive in 5/5 seeds, paired $t = 3.39$, one-sided $p = 0.014$. **This is post-hoc and was not pre-registered;** we

Table 4: Mode S (statistic-only), $n=5$ per arm ($n=8$ for `cos_krum`): mean ASR with the two measured channels. **The decision channel moves far more than ASR throughout.** [†] below the 0.35 accuracy gate. [‡] the replication arm, rank-discordant by design, frozen separately: $\Delta=+0.098$ is below reputation’s $+0.178$, so the cross-arm admission *ordering* is refuted while the negative replicates. Verdicts otherwise scored against readings frozen in advance; per-seed values in Appendix N.

arm	channel	$\rho=1$	2.72	7.39	54.6	test	verdict
krum / scaling	ASR	0.062	0.016	0.023	0.036	$p_{\uparrow}=0.70$	H-admission confirmed, H-stat. refuted
	decision	0.000	0.533	0.800	0.733		
	admission	0.000	0.000	0.000	0.000		
reputation / scaling	ASR	0.017	0.019	0.025	0.195	rise $+0.178$ vs. ladder $+0.762$ ($p=0.0017$)	H-admission confirmed (mostly atten.)
	decision	0.000	0.933	1.000	1.000		
	admission	0.000	0.002	0.008	0.020		
cos_krum / pixel	ASR	0.584	0.411	0.295	0.159 [†]	$p_{\downarrow}=8\times 10^{-4}$ $ \Delta =0.425$	refuted (Prop. 1 practical)
	decision	0.000	0.000	0.000	0.000		
	admission	0.000	0.000	0.000	0.000		
coord_median / pixel [‡]	ASR	0.443	0.442	0.490	0.541	rise $+0.098$ $< +0.150$ ($p_{\uparrow}=0.069$)	ordering refuted; negative repl.
	decision	0.000	0.232	0.358	0.482		
	admission	0.000	0.005	0.013	0.033		

Table 5: The ladder as measured, $n=5$, per-attack mean ASR (mean clean accuracy). *Predicted* shapes were frozen before any run (Appendix X); $p_{\uparrow}$ is the one-sided Jonckheere–Terpstra probability of an increasing trend, the only direction any frozen rule tests. An arm meeting neither its confirmation nor its refutation criterion is reported *indeterminate* rather than scored in our favour. [†] below the 0.35 accuracy gate: a collapsed model, not suppression. Intervals, per-seed values, anchor residuals and both plotted axes: Appendix N.

arm (d_2 / attack)	predicted shape	mean ASR (clean acc.)				JT $p_{\uparrow}$	verdict
		$\kappa=0$	0.5	1.0	2.0		
1 krum / scaling	monotone rise	0.062 (0.57)	0.011 (0.65)	0.014 (0.68)	0.054 (0.62)	0.63	refuted
2 reputation / scaling	monotone rise	0.017 (0.78)	0.022 (0.78)	0.487 (0.77)	0.779 (0.70)	1.7×10^{-6}	confirmed
3 cos_krum / pixel	FLAT (control)	0.490 (0.61)	0.450 (0.57)	0.153 (0.47)	0.086 (0.22) [†]	0.999	indeterminate
4 coord_median / pixel	rise, shallower	0.443 (0.77)	0.428 (0.76)	0.327 (0.75)	0.170 (0.69)	0.998	refuted

1164 report it because it is the honest description of a small consistent effect, and we label it post-hoc
1165 wherever it appears. It does not change the frozen verdict: the admission *ordering* is refuted either
1166 way.

1167 **The primary test, and two post-hoc sensitivities that cannot change its verdict.** The frozen
1168 primary test is a one-sided Spearman correlation of *rise* against *measured disturbance* pooled over
1169 all 16 cells, confirmed only at $\rho_s > 0$ with $p < 0.05$: we obtain $\rho_s = +0.351$, $p = 0.091$, so it
1170 is **refuted**. Two sensitivities are recorded because they are informative about *why*, and are labelled
1171 post-hoc because neither was frozen. (i) The four $\kappa=0$ cells are constructed (0, 0) points (zero
1172 disturbance and zero rise by definition) and dropping them gives $\rho_s = +0.780$, $p = 0.001$ ($n=12$);
1173 so the pooled correlation is carried by the arm-2 rise and destroyed by the two arms whose ASR
1174 falls, not by the identity anchors. (ii) Against the *nominal* dose κ (i.e. $\log \rho$) rather than measured
1175 disturbance, $\rho_s = -0.354$, $p = 0.911$: the wrong sign, which is the same fact as three arms falling.
1176 Neither is scored.

1177 **Per-arm trend tests.** One-sided Jonckheere–Terpstra across the four ordered rungs, 5 seeds each,
1178 $\alpha = 0.05$; the normal-approximation p agrees with a 20,000-draw permutation p to three decimals
1179 throughout. Arm 1: $J = 70$, $z = -0.34$, $p_{\uparrow} = 0.632$. Arm 2: $J = 144$, $z = +4.64$, $p_{\uparrow} =$
1180 1.7×10^{-6} . Arm 3: $J = 29$, $z = -3.10$, $p_{\uparrow} = 0.999$ ($p_{\downarrow} = 9.8\times 10^{-4}$); dropping the gated $\kappa=2.0$
1181 rung leaves $z = -2.06$, $p_{\downarrow} = 0.020$, so the decrease is not an artifact of the collapsed cell alone.
1182 Arm 4: $J = 31$, $z = -2.96$, $p_{\uparrow} = 0.998$ ($p_{\downarrow} = 1.5\times 10^{-3}$). No frozen rule tests the decreasing
1183 direction; $p_{\downarrow}$ is descriptive.

1184 **Clean accuracy against dose separates two ways of losing ASR.** Descriptive, not pre-registered.
1185 Accuracy trends by the same test: arm 1 $z = +1.35$, arm 2 $z = -3.30$, arm 3 $z = -3.50$, arm
1186 4 $z = -4.11$. Arm 3 pays $0.607 \rightarrow 0.220$: a 0.39 collapse through the Lemma 1 contraction




Figure 2: **The preservation chain, and where it breaks.** (P1) $\Rightarrow$ (P2) $\Rightarrow$ (P3) holds by definition; the (P3)–(P4) link breaks in both directions, and both witnesses come from the oracle intervention that pins every adversary’s coefficient and so closes Lemma 1’s attenuation channel by construction (§4). The pooled statistic-level prediction we had frozen is separately **rejected** by a pre-registered ladder (§5), and a pre-registered replication in a different invariance class reproduces the negative but **refutes the cross-arm admission ordering**, so **admission is a direction, not a validated predictor**. Every number in the figure is recomputed from the scoring scripts’ own JSON by `story_chain_figure.py`.

Table 6: All 16 rungs, $n=5$. *disturb* is the independently measured disturbance of that arm’s own statistic at that rung, computed before the freeze with no ASR involved; *rise* is ASR relative to the arm’s own $\kappa=0$ rung and is the primary test’s ordinate. CIs are Student- t at $n=5$ and are wide by design: the ladder was powered to detect a monotone trend, not to resolve single cells. † fails the 0.35 clean-accuracy gate.

arm	κ	ρ	dist.	ASR	95% CI	rise	acc	per-seed ASR (42–46)					
krum scaling class (c)	0	1.00	0.000	0.062	[0.019, 0.105]	+0.000	0.565	0.115	0.021	0.047	0.057	0.069	
	0.5	2.72	0.600	0.011	[−0.001, 0.023]	−0.051	0.652	0.000	0.010	0.011	0.027	0.007	
	1.0	7.39	0.733	0.014	[0.003, 0.024]	−0.048	0.682	0.003	0.015	0.021	0.021	0.008	
	2.0	54.60	0.867	0.054	[0.002, 0.106]	−0.008	0.619	0.005	0.039	0.077	0.113	0.037	
reputation scaling class (c)	0	1.00	0.000	0.017	[0.012, 0.022]	+0.000	0.777	0.014	0.012	0.017	0.023	0.019	
	0.5	2.72	0.867	0.022	[0.014, 0.031]	+0.005	0.779	0.016	0.018	0.024	0.033	0.020	
	1.0	7.39	0.867	0.487	[0.349, 0.624]	+0.470	0.765	0.439	0.547	0.578	0.557	0.313	
	2.0	54.60	0.933	0.779	[0.673, 0.885]	+0.762	0.701	0.710	0.843	0.730	0.896	0.714	
cos_krum pixel class (a)	0	1.00	0.000	0.490	[0.058, 0.922]	+0.000	0.607	0.003	0.269	0.630	0.864	0.686	
	0.5	2.72	0.000	0.450	[0.130, 0.769]	−0.041	0.569	0.764	0.236	0.165	0.648	0.435	
	1.0	7.39	0.000	0.153	[−0.127, 0.433]	−0.337	0.466	0.132	0.090	0.000	0.543	0.001	
	2.0	54.60	0.000	0.086	[−0.100, 0.271]	−0.405	0.220 †	0.083	0.000	0.000	0.345	0.000	
coord_median pixel class (b)	0	1.00	0.000	0.443	[0.263, 0.622]	+0.000	0.767	0.395	0.287	0.449	0.677	0.405	
	0.5	2.72	0.251	0.428	[0.241, 0.615]	−0.015	0.764	0.454	0.215	0.430	0.638	0.402	
	1.0	7.39	0.378	0.327	[0.181, 0.472]	−0.116	0.749	0.351	0.139	0.352	0.463	0.328	
	2.0	54.60	0.487	0.170	[0.113, 0.228]	−0.272	0.690	0.179	0.090	0.204	0.192	0.187	

channel, and the reason its $\kappa=2.0$ cell is uninterpretable rather than a success. Arm 4 pays only 0.767 $\rightarrow$ 0.690 while ASR more than halves, so there the dose genuinely *helps* coordinate median: a refutation of our predicted shape, but a benign mechanism. Arm 2 is the cleanest cell in the ladder, losing 0.076 accuracy while ASR rises $46\times$.

Anchor residuals, and why the instrument is the finding. Residual is real-transform ASR ($n=5$, from `metric_swap_topup`) minus the synthetic rung at the nearest measured disturbance. NormClip $\rightarrow$ Krum: disturbance 0/9, ASR 0.062 vs. rung 0.062, residual +0.001. RFA $\rightarrow$ Krum: 6/9 = 0.667, ASR 0.779 vs. rung 0.011 at disturbance 0.600, residual +0.767. NormClip $\rightarrow$ Reputation: 3/9 = 0.333, ASR 0.806 vs. rung 0.017, residual +0.789.




Figure 3: **The two regimes, exhibited: the payload dose separates them.** *The picture of regimes A and B of §2, produced by the instrument of §5.* With the statistic’s dispersion fixed and only the adversary/benign weight ratio r swept, `coord_median` (a rank aggregator, the class the regimes are stated for) is single-peaked at $r=1$ as Theorem 1 predicts, while reputation is monotone: off-spec, and the theorem is not claimed for it. Bars are Student- t 95% intervals, recomputed per seed by `analyze_targeted_dose.py`; the Mode-S companion is Table 4.




Figure 4: **The measured axis: ASR against disturbance of each arm’s own statistic**, an abscissa fixed before the freeze ($n=5$, Student- t 95% intervals; the figure script imports the arm and anchor construction from the scoring script, so the two cannot diverge). Stars and crosses are the two *real* upstream transforms measured on these same cells: at matched disturbance they sit +0.77 and +0.79 ASR *above* the synthetic curve, though their realized weight ratios (mean 3.61 and 11.93, maximum 4.90 and 17.38) lie *inside* the ladder’s span; §5 attributes the gap to Δ_c , post-hoc. The class-(a) arm’s four rungs coincide at disturbance exactly 0 and its ASR falls anyway: Lemma 1’s magnitude channel, not C2 disturbance. Hollow markers fail the 0.35 accuracy gate.

1196 RFA→Reputation: 9/9, 0.895 vs. 0.779, residual +0.116. Both `cos_krum` anchors sit at distur-
 1197 bance 0 and land (+0.000, −0.044). `coord_median` was not in the metric-swap menu, so arm 4 has
 1198 *no* real-transform anchor and its ecological validity is untested. The two large residuals are not a
 1199 range mismatch: realized ρ is mean 3.61 / max 4.90 for NormClip and mean 11.93 / max 17.38 for
 1200 RFA, both inside the ladder’s [1.00, 54.60] span. What differs is *which* clients get reweighted (Norm-
 1201 Clip and RFA reweight in a way correlated with adversary status, our permutation deliberately does
 1202 not) and we state that as the hypothesis caveat (iii) was written to test, not as a measured quantity:




Figure 5: **The nominal axis:** the same 16 cells against the dose $\rho = e^{2\kappa}$, coloured by Proposition 1 class. Of the two class-(c) arms the fragility claim is stated over, only reputation rises; the other two arms *fall*. This is the axis the frozen per-arm shapes were stated on, and the pooled prediction was stated on the measured axis of Figure 4. Hollow markers fail the 0.35 accuracy gate.

the shipped disturbance logs record only $c_{\min}$, $c_{\max}$ and ρ per round, not per-client coefficients, so an adversary-versus-benign coefficient contrast requires a new (cheap, no-ASR) measurement.

Two checks that passed. The $\kappa=0$ rung reproduces each arm’s published standalone value exactly on every shared seed (16 seed-level comparisons across four arms, largest $|\Delta\text{ASR}|$ and $|\Delta\text{acc}|$ both $0.00\text{e}+00$ at tolerance 10^{-6}), so $\kappa=0$ really is the identity and the fresh rung and the published baseline are the same quantity. **The first-rung fall we rest the (P4) $\nrightarrow$ (P5) witness on,** in full: `cos_krum / pixel` at $\rho=2.72$ falls 0.173 with clean accuracy moving only $0.599 \rightarrow 0.567$, paired $t=-1.976$, one-sided $p=0.044$, falling in 5/8 seeds. And `cos_krum`’s selection never changed in any of the 120 measured cells, with maximum relative score drift 1.1×10^{-4} : Proposition 1’s classification of which statistics a positive per-client rescaling can disturb is verified directly, independently of the outcome the ladder then failed to predict.

O Secondary results: baselines, the adaptive stress test, scale, and switching

The margin behaves like a certificate, not just a switch. C3 is binary, but the margin m is measurable and its *strength* tracks security: over the 29 measured rounds with $w_a > 0$, higher margin goes with lower realized adversarial mass (Spearman $\rho = -0.48$, $p = 0.009$), Corollary 1’s bound holding in every one, a number rather than a verdict, though the link is directional, not fitted, and post-hoc.

Single-defense baselines. Measured under the identical protocol (5 seeds, 50 rounds), no individual defense achieves low max-committed ASR: CM 0.52, FG 0.73, Rep 0.84, TM 0.85, RFA 0.89, NC 0.94. Against the best single defense (CM, 0.52), FG→RFA (0.093) is $5.6\times$ lower and FG→CM (0.107) is $4.9\times$ lower. FoolsGold alone reaches ASR 0.20 on model-scaling only by collapsing the model to 10% accuracy, so it is not suppression: consistent with its C1 failure.

Adaptive attacks against FG→RFA. We evaluate one whitebox adaptive attack family that jointly evades both defenses: disjoint coordinate blocks minimize cosine similarity (evade FG) while ϵ -ball projection stays near the geometric median (evade RFA). The adversary has full state access, and these constraints are in tension. ASR against FG→RFA (5 seeds each, base quoted on the *same* five seeds so the comparison is seed-matched rather than against the $n=30$ mean): base committed 0.051 ± 0.011 ; Neurotoxin [Zhang et al., 2022] top-20% at scale 2.0 0.048 ± 0.009 ; projection only, no decorrelation, $\epsilon=2$: 0.050 ± 0.008 . With joint constraints the sweep over ϵ gives 0.036 ± 0.015 ($\epsilon=0.5$), **0.167 $\pm$ 0.069** ($\epsilon=1$), 0.130 ± 0.030 ($\epsilon=2$), 0.124 ± 0.025 ($\epsilon=4$). ASR peaks at $\epsilon=1$: tighter projection over-constrains the attack, larger ϵ fails to evade RFA, and without decorrelation




Figure 6: **The criterion is a mechanism test, and the mechanism is measurable.** (a) What C2 checks: the upstream transform T and the statistic S that d_2 reads determine an *invariance class* (Proposition 1), and the class determines the verdict with no composition run, which is why norm clipping does not blind FoolsGold but does blind Reputation. (b) The measured chain for FG→RFA, per round, on the transformed points RFA actually receives; each marker is one round with $w_a > 0$ (rounds where FoolsGold zeroes every adversary are Lemma 1 and carry no payload), large markers are mean $\pm$ s.d., and values below 10^{-4} are drawn at 10^{-4} . Corollary 1’s bound holds in every round, and the $\alpha=0.1$ collapse is visible in the theorem’s own quantity: m falls from 4.71 ± 1.51 ($n=17$) to 2.06 ± 2.35 ($n=12$), with several rounds at $m < 1$ where the bound degrades to $n_a/(n_b+n_a)$, and realized Λ_a rises $0.003 \rightarrow 0.085$ (max 0.50) as composition ASR goes $0.093 \rightarrow 0.569$.

Table 7: Framework validation (representative pairs; full tallies in text). Threshold: LOW < 0.5 , HIGH ≥ 0.5 .

Composition	Category	Predicted	Max ASR	Correct?
<i>PASS: C1∧C2∧C3 (predicted LOW)</i>				
FoolsGold → RFA	PASS	LOW	0.093	✓
FoolsGold → CoordMedian	PASS	LOW	0.107	✓
Reputation → CoordMedian	PASS	LOW	0.423	✓
<i>C1-FAIL: no constituent suppresses (predicted HIGH)</i>				
Reputation → RFA	C1-FAIL	HIGH	0.825	✓
NormClip → FoolsGold [†]	C1-FAIL	HIGH	0.887	✓
RFA → FoolsGold [†]	C1-FAIL	HIGH	0.914	✓
<i>C2-FAIL: discriminative property destroyed (predicted HIGH)</i>				
NormClip → Reputation	C2-FAIL	HIGH	0.819	✓
<i>C3-FAIL: margin compressed (predicted HIGH)</i>				
FoolsGold → TrimmedMean	C3-FAIL	HIGH	0.508	✓
NormClip → RFA	C3-FAIL	HIGH	0.853	✓
<i>False negatives (predicted HIGH, measured LOW)</i>				
NormClip → CM / RFA → CM	C3-FAIL	HIGH	0.433 / 0.485	×

[†]Reclassified from C2-FAIL: by Proposition 1(a) these preserve C2.

1234 FoolsGold suppresses it outright, so the competing constraints leave a narrow effective window.
1235 This is an empirical finding (Theorem 1 concerns mechanism preservation and does not bound an
1236 adaptive adversary whose optimization is untied to its conditions) and it probes one natural family
1237 of joint constraints at a handful of ϵ values, not the adaptive attack space. Within it, 0.167 is $3.3\times$
1238 the seed-matched committed-attack ASR ($1.8\times$ the $n=30$ mean).

1239 **Scale.** At $N=100$, $f=0.05$, ResNet18 both pre-registered PASS predictions hold: FG→RFA max
1240 ASR = 0.069, FG→CM = 0.152. We flag that these runs reach only ≈ 0.29 clean accuracy at 50
1241 rounds (ResNet18 is far from converged at this budget), so they establish that suppression transfers
1242 across scale and architecture, not that the composition is deployable at $N=100$ as configured.

1243 O.1 Heterogeneity, and where the criterion breaks

1244 Earlier drafts swept heterogeneity for single defenses only; no composition had been run at any
 1245 $\alpha \neq 0.5$. We therefore ran all three PASS pairs across $\alpha \in \{0.1, 1.0, 10.0\}$ (3 seeds, both attacks,
 1246 50 rounds; max-committed ASR / clean accuracy):

	Composition	$\alpha=0.1$	$\alpha=0.5$	$\alpha=1.0$	$\alpha=10.0$
1247	FG→RFA	0.569 / 0.20	0.093 / 0.50	0.051 / 0.57	0.063 / 0.62
	FG→CM	0.433 / 0.40	0.107 / 0.68	0.133 / 0.74	0.129 / 0.78
	Rep→CM	0.548 / 0.48	0.423 / 0.76	0.377 / 0.79	0.326 / 0.81

1248 In our sweep, which is three seeds per cell and fixes direction rather than magnitude, the certified
 1249 compositions stay below threshold for $\alpha \in \{0.5, 1.0, 10.0\}$, including the near-IID setting where one
 1250 might expect FoolsGold to founder on legitimately similar benign updates: $\alpha=10$ is indistinguish-
 1251 able from $\alpha=0.5$. **Under extreme heterogeneity ($\alpha=0.1$) two of the three PASS compositions**
 1252 **cross the threshold** (FG→RFA 0.569, Rep→CM 0.548) and only FG→CM survives, at 0.433: a
 1253 real boundary condition on the framework, with caveats in both directions: FG→RFA at $\alpha=0.1$ is
 1254 strongly bimodal across seeds (0.065–0.923), so at $n=3$ the direction is clear but the magnitude is
 1255 not, and clean accuracy degrades badly (0.20–0.48), so the regime stresses utility as well as secu-
 1256 rity. Crucially Theorem 1 supplies a diagnostic quantity that retrospectively explains this boundary:
 1257 re-measuring the transformed points at $\alpha=0.1$, RFA’s separation margin $m := \Delta'_{\text{sep}}/[2(R'_B + \delta')]$
 1258 collapses from an all-round mean of 3.24 (met in 68% of rounds) to 1.01 (met in 20%), over the
 1259 payload-carrying rounds alone, the ones Figure 6 plots, the same collapse reads 4.71 ± 1.51 ($n=17$)
 1260 to 2.06 ± 2.35 ($n=12$) with the per-round minimum falling $2.82 \rightarrow 0.00$, the difference being that
 1261 the all-round means average in the Lemma 1 rounds at $m=0$, exactly the boundary $m \rightarrow 1$ where
 1262 Corollary 1’s bound degrades to 40% adversarial mass, because extreme heterogeneity inflates the
 1263 benign radius R'_B until the benign cluster is no longer separated from the adversarial points, which
 1264 is the precondition the theorem requires. A practitioner running the 5–10 round pilot we recommend
 1265 would measure $m \approx 1$ and decline to deploy. The criterion is therefore mechanistically diagnostic at
 1266 this boundary (the property that makes a screening rule usable) though this is a post-hoc explanation
 1267 of a known failure, not a pre-registered prediction of it.

1268 O.2 Composition vs. randomized switching

1269 **High persistence sharply reduces the value of temporal switching.** In our tested regime,
 1270 $\gamma \approx 0.66$ ($\gamma := \text{ASR}(t+1)/\text{ASR}(t)$ after one clean round): triggers decay slowly enough that
 1271 intermittent suppression cannot clear them, so temporal mixing approaches the weaker constituent
 1272 defense’s performance. Whether switching retains any value depends jointly on γ , injection strength,
 1273 switching frequency, and recovery dynamics: persistence alone does not imply switching is worth-
 1274 less. The quantitative boundary is architecture-specific; the qualitative prediction (high- γ attacks
 1275 blunt switching) is what this secondary analysis contributes, and nothing in the preservation argu-
 1276 ment rests on it.

1277 **Order-sensitivity control.** Both variants apply both defenses every round and differ *only* in order,
 1278 so compute is matched and any gap is structural. We report it as a **motivating demonstration**
 1279 **that composition order is structurally consequential, not as evidence for the criterion:** the
 1280 mechanism behind the largest gap is a degeneracy rather than a signal-preservation failure, as the
 1281 diagnosis below makes clear. Pixel ASR, 5 seeds: fixed FG→CM 0.076 ± 0.035 versus randomized
 1282 50/50 ordering 0.681 ± 0.133 ($9\times$) and temporal mixing of one defense per round 0.679 ± 0.159 ;
 1283 fixed FG→RFA 0.037 ± 0.011 versus randomized 0.719 ± 0.113 ($19\times$). The mechanism is the
 1284 degeneracy of §1: CoordMedian supplies no per-client transform, so every CM→FG round reduces
 1285 to FoolsGold alone, leaving half the rounds defended by a single constituent that cannot suppress
 1286 on its own. On a reputation+TrimmedMean menu ($n=30$), composition reaches ASR 0.316 against
 1287 0.915 for temporal mixing.

1288 P The metric-swap suite’s two-sided prediction

1289 Because C1 binds for all eight pairs, every frozen label in the metric-swap suite was HIGH (a
 1290 consequence of the measured inputs rather than a later choice), so the suite bears on specificity only.
 1291 One prediction inside it was genuinely two-sided. `cos_krum`’s C1 verdict on model-scaling turns
 1292 on whether the accuracy gate is applied (0.078 ASR at 0.146 clean accuracy), so we froze *both*
 1293 readings in advance: primary (gated) $\Rightarrow$ HIGH, secondary (numeric only) $\Rightarrow$ LOW, with the rule
 1294 fixed beforehand that a LOW outcome at accuracy ≥ 0.35 would mean the gate was wrong. Both
 1295 cells came out **HIGH** (0.508, 0.566), so the gate is vindicated against a pre-registered alternative
 1296 that would have overturned it. This is the only place in the paper where our own accuracy gate was
 1297 put at risk by a frozen rule.

1298 Q The prospective suite, arm by arm

1299 (i) **FLTrust as d_1 (7 pairs, 2/7): out of scope.** When d_1 alone suppresses, the composition in-
 1300 herits that suppression whatever d_2 does, so all five predicted-HIGH pairs are misses even though
 1301 their C2/C3 failures are real: FLTrust genuinely destroys Reputation’s consensus-distance signal,
 1302 but destroying a mechanism that was never load-bearing costs nothing. This is C0, and it was invis-
 1303 ible on the original 42 pairs, where the strongest single defense reaches only 0.519 and the lowest
 1304 ASR among all 39 non-PASS pairs is 0.433; no pair was ever out of scope. We apply C0 here
 1305 *retrospectively*, which is weaker than a pre-registered exclusion.

1306 (ii) **FLTrust as d_2 (6 pairs, 6/6): in scope, but nearly an identity.** Proposition 1(a) predicts
 1307 $\cdot \rightarrow \text{FLTrust}$ preserves C2, and it does, but FLTrust gates on a scale-invariant cosine and then *renor-*
 1308 *malizes every update to* $\|g_{\text{STV}}\|$, discarding magnitude outright, so a positive rescaling $u_i \mapsto c_i u_i$
 1309 is *exactly* cancelled: FLTrust discards precisely what d_1 modifies. The per-seed data show this di-
 1310 rectly. $\text{Krum} \rightarrow \text{FLTrust}$ and $\text{MKrum} \rightarrow \text{FLTrust}$ are **bit-identical** to FLTrust alone on all six runs
 1311 (pass-through d_1); $\text{NC} \rightarrow \text{FLTrust}$ is bit-identical *on the pixel attack*, where no client exceeds $\tau=5$
 1312 so $c_i \equiv 1$, and differs only within seed noise on model-scaling, where clipping fires; $\text{Rep} \rightarrow$ and
 1313 $\text{RFA} \rightarrow \text{FLTrust}$ likewise differ only within noise. The lone real deviation is $\text{FG} \rightarrow \text{FLTrust}$ (0.064,
 1314 up to 0.084 per seed): the one d_1 producing exact *zeros* rather than positive rescalings, changing
 1315 FLTrust’s client set instead of its magnitudes, which is the Lemma 1 interaction. So this arm con-
 1316 firms that theory and implementation agree, and demonstrates the degeneracy mechanism cleanly,
 1317 but it is **not falsifiable**: every d_1 in our menu is a *non-negative* per-client rescaling, FLTrust cancels
 1318 the positive part exactly, and so the only channel left open is the *zero pattern*, which is precisely
 1319 where the one deviation appears. Proposition 1(a) permits no other outcome on the strictly-positive
 1320 part, and Lemma 1 accounts for the rest.

1321 (iii) **Neither defense is FLTrust (7 pairs, 5/7): the only arm that could have surprised us.**
 1322 All were predicted HIGH; correct are $\text{rep} \rightarrow \text{krum}$ 0.880, $\text{nc} \rightarrow \text{mkrum}$ 0.943, $\text{rfa} \rightarrow \text{mkrum}$ 0.875,
 1323 $\text{krum} \rightarrow \text{rfa}$ and $\text{mkrum} \rightarrow \text{rfa}$ both 0.870. Neither miss is a mechanism error: $\text{fg} \rightarrow \text{krum}$ reads 0.078
 1324 at 9.9% clean accuracy (a collapsed model, the same hollow-suppression artifact as FoolsGold alone)
 1325 and $\text{krum} \rightarrow \text{cm}$ reads 0.450 against a 0.5 threshold as a degenerate pair reducing to CoordMedian
 1326 alone, whose own baseline is 0.519: threshold noise on a borderline constituent. This arm contains
 1327 *zero* predicted-LOW pairs, so like Tier 2 it bears on specificity only.

1328 **The suggestive contrasts, and why they are still confounded.** Three matched pairs hold d_1 fixed
 1329 while varying only d_2 ’s invariance class: $\text{rep} \rightarrow \text{fltrust}$ 0.046 vs. $\text{rep} \rightarrow \text{krum}$ 0.880; $\text{nc} \rightarrow \text{fltrust}$ 0.046
 1330 vs. $\text{nc} \rightarrow \text{mkrum}$ 0.943; $\text{rfa} \rightarrow \text{fltrust}$ 0.047 vs. $\text{rfa} \rightarrow \text{mkrum}$ 0.875: a 19–20 $\times$ gap turning on exactly
 1331 the property Proposition 1 identifies. They are nonetheless confounded in the way Lemma 2 pre-
 1332 dict: the invariant-signal d_2 is also the only d_2 in this menu that suppresses alone, so effectiveness
 1333 and mechanism co-vary across the contrast. This is why breaking the confound requires holding
 1334 standalone effectiveness fixed rather than choosing a different defense.

1335 R Adversary-correlated dispersion Δ_c , measured

1336 R.1 Heterogeneity is not the mechanism: ρ against Δ_c

1337 **The finding, stated first.** What costs a downstream defense its suppression is not how heteroge-
 1338 neous the upstream transform is but **whether its reweighting is correlated with adversary iden-**
 1339 **tity**, and not even the size of that correlation, which has **two opposite regimes**. This is why our own
 1340 headline predictor failed.

1341 **Two different things a coefficient vector does, and the theory sees only one.** (Symbols: Ap-
 1342 pendix G.) Proposition 1 turns on **dispersion**, $\rho := \max_i c_i / \min_i c_i$ (how far apart the coefficients
 1343 are, which decides whether an ordering survives) and says nothing about *which* clients receive which
 1344 coefficient. That second property is **adversary-correlated dispersion**,

$$\Delta_c := \mathbb{E}[c_i \mid i \in A] - \mathbb{E}[c_i \mid i \notin A], \quad A \text{ the adversarial participants,}$$

1345 and the two are **independently controllable**, not merely uncorrelated in our sample: a transform
 1346 can spread the coefficients arbitrarily at $\Delta_c = 0$ exactly, or leave them nearly equal while shifting
 1347 the adversary’s share, and Proposition 3 exhibits a family of each. **Proposition 1 and Theorem 1**
 1348 **constrain ρ ; Lemma 1’s attenuation channel is about Δ_c** , and §5 finds the real transforms differ
 1349 on the second. Δ_c is *not* a screening quantity: computing it requires knowing A (§6).

1350 **Two mechanism-preserving regimes, at opposite ends of ρ .** Lemma 1 is a second, inde-
 1351 pendent route to suppression, and the two together explain all three PASS compositions. Let
 1352 $r := w_a / \max_b w_b$ and S be the attack’s coordinate extremeness. **(A) Ordering preservation**
 1353 keeps the adversary in the tail where a rank aggregator discards it ($S \cdot r > 1$). **(B) Attenuation**
 1354 instead drives the adversarial transformed value *inside* the benign spread, where it cannot be an ex-
 1355 treme order statistic and its payload is scaled by r ($S \cdot r \ll 1$; $r = 0$ is exact annihilation). The
 1356 hazard is $S \cdot r \approx 1$, and our two weighting defenses sit at opposite ends: FoolsGold’s weights are
 1357 bounded ($\rho \leq 2.1$, $S \cdot r \geq 4.8$), regime A; Reputation’s are wildly unbounded (ρ mean 1.4×10^5)
 1358 precisely *because* it crushes the adversary ($S \cdot r = 0.017$), regime B. So Reputation→CoordMedian
 1359 violates the ordering condition by five orders of magnitude and still suppresses (Figure 3, App. N,
 1360 exhibits both regimes). Hence ρ does not order Δ_c (Proposition 3), and Δ_c ’s *magnitude* does not
 1361 order suppression either, the same large $|\Delta_c|$ being protective at one end and catastrophic at the
 1362 other. **Neither quantity is a predictor**; they buy an account of *why* a heterogeneity-based reading
 1363 mispredicts.

1364 $\Delta_c := \mathbb{E}[c_i \mid i \in A] - \mathbb{E}[c_i \mid i \notin A]$ for every per-client transform in the paper,
 1365 over the same 5 seeds $\times$ 3 rounds and, within a round, the *same* raw update stack for every
 1366 transform, so the rows are comparable by construction (measure_coefficient_targeting.py,
 1367 results/coefficient_targeting.json). **No ASR is computed and no model is trained per**
 1368 **transform**, which is what lets this be added without rescoring anything; the figures are **post-hoc and**
 1369 **exploratory** and no frozen prediction depends on them. Coefficients are read back from transformed
 1370 norms, $c_i = \|T(u_i)\|/\|u_i\|$, so Table 8 measures the shipped code path rather than the builder’s
 1371 description of it. As a provenance check the synthetic families’ adversarial coefficient shares repro-
 1372 duce results/admission_measurement.json to 0.0 absolute on all 13 cells, so these rows may
 1373 be read on the same axis as the Mode-S and Mode-A figures of §4.

1374 Three readings, in decreasing confidence. **(1) ρ is not a proxy for Δ_c** : the two are anti-ordered
 1375 on the anchors, and NormClip on the pixel backdoor has $\rho = 1$ exactly because no client exceeds
 1376 $\tau=5$, so $\Delta_c = 0$ there for a reason the invariance classification cannot see. **(2) The two anchors**
 1377 **that miss the ladder are the two most extremely targeted**, which explains the residuals of §5
 1378 without appeal to a range mismatch. **(3) Δ_c is nonetheless not a predictor of suppression**: RFA
 1379 is the most negative of the four real transforms and rfa→krum fails at 0.779, while Reputation is
 1380 more negative still and rep→coord_median suppresses. That is the two-regime structure of §2 again:
 1381 strong negative Δ_c can mean either annihilation or mere dilution, and Δ_c ’s magnitude alone does
 1382 not say which. We report it as the axis on which the ladder’s anchors were unmatched, not as the
 1383 quantity that replaces the statistic.

Table 8: ρ says how much a transform spreads the coefficients; Δ_c says who receives the spread. They are not ordered together: RFA and NormClip target the adversary 2–3.6 $\times$ harder than the ladder’s most dispersed rung at ρ that is 3.7–13 $\times$ *smaller*. Means over non-degenerate rounds; the synthetic rows are attack-pooled since their coefficients do not depend on the attack.

transform	attack	ρ	$\mathbb{E}[c \mid A]$	$\mathbb{E}[c \mid \neg A]$	Δ_c
NormClip	model-scaling	4.19	0.268	0.989	−0.721
NormClip	pixel	1.00	1.000	1.000	0.000
Reputation	model-scaling	8.3×10^5	0.003	1.387	−1.383
Reputation	pixel	1.68	0.915	1.019	−0.105
FoolsGold	model-scaling	1.9×10^{12}	1.497	0.834	+0.663
FoolsGold	pixel	1.8×10^{12}	1.365	0.986	+0.378
RFA	model-scaling	14.8	0.125	1.332	−1.208
RFA	pixel	1.69	0.923	1.012	−0.089
<i>Synthetic families of §5–§4, same code path:</i>					
ladder, $\kappa=0.5/1.0/2.0$	—	2.72/7.39/54.6	—	—	−0.140/−0.240/−0.336
Mode S, all rungs	—	1.00–54.6	—	—	0.000
Mode A, $\nu=-2 \dots +2$	—	7.39–7.39	—	—	−1.138 $\dots$ + 2.469

S Confusion matrix in full

Exactly 5 of the 42 pairs have max-committed ASR < 0.5 : the 3 PASS pairs plus the 2 false negatives. Pooling both in-distribution tiers (each pair scored against the label committed for it at the time) gives TP = 3, FP = 2, TN = 35, FN = 2 (38/42; specificity 94.6%, which we report but do not lead with, since 37 of 42 pairs are negatives and specificity is inflated by that imbalance; precision 1.0 at base rate 11.9% is the 8.4 $\times$ enrichment that matters, and recall is 40% under the tightened C1, since FG $\rightarrow$ RFA now counts as a miss). We report this pooled figure with its decomposition rather than on its own, because the tiers are not equivalent evidence: it is 16/18 development *fit* plus 22/24 held-out *prediction*. The four sub-tiers the body refuses to pool are development *fit*, held-out *specificity*, out-of-sample *positives with no informative negatives*, and the 13/20 *prospective* suite. Scoring instead against the criterion’s mechanism labels, which classify both failed LOW calls as HIGH, gives TP = 3, FP = 0, TN = 37, FN = 2 (40/42); we take the more conservative pooled figure as primary. The two false positives, fg $\rightarrow$ NC (0.927) and fedavg $\rightarrow$ CM (0.519), were hedged as speculative at pre-registration time and neither satisfies C1; every pair that actually satisfies C1 $\wedge$ C2 $\wedge$ C3 came out LOW. Category tally after the Proposition 1 correction: degenerate 18, C1-fail 11, C2-fail 3, C3-fail 7, PASS 3, and all 3 surviving C2 failures are $\cdot \rightarrow$ Reputation, exactly the non-invariant class of Proposition 1(c). With a positive set of only 2 certified in-distribution pairs and recall 40%, the 94.6% specificity figure is high partly because the positive class is rare, so we report it as a supporting statistic rather than a headline. The conditions are sufficient, not necessary, and they do reject working compositions. Extension to 3-way compositions: 9/10 triples correct (too small for confidence).

T Diagnosing the two false negatives in full

NormClip $\rightarrow$ CM (0.433) and RFA $\rightarrow$ CM (0.485) were labelled C3-FAIL on the reading that *any* heterogeneous upstream rescaling compresses a rank aggregator’s margin. Theorem 1(1) says something weaker and quantitative: ordering survives when the attack’s coordinate extremeness exceeds the realized weight ratio, $S > \rho$. Measured, NormClip’s ratio is mild. Over 9 live rounds it clips only 0–2 of 5 clients (11 of 45 client-rounds), so $\approx 76\%$ of clients receive $c_i = 1$ *exactly*, and the largest raw norm we observe is 24.5 against $\tau=5$, bounding $\rho \leq 4.9$; against $S=10$ the margin $S/\rho \geq 2.0$ still clears 1. RFA-reweighting is likewise near-uniform on benign clients. Both pairs therefore sit close to the *uniform* rescaling case, where ordering is preserved exactly, which is why their ASR lands in an intermediate band just under threshold rather than in the high band the binary label predicts. **The refinement does not recover them.** Both also fail C1: CoordMedian alone is already at 0.519, and neither NormClip (0.935) nor RFA (0.885) suppresses, so no reading of C3 makes either pair certifiable. What the diagnosis buys is a sharper condition and an account of the residual: a d_2 whose own baseline sits at 0.519 needs only a mild order-preserving upstream trans-

form to cross 0.5, so these two are threshold noise around a marginal aggregator, not suppression the framework failed to recognise. The operational refinement is to evaluate C3 for rank aggregators against the *realized* ρ from a 5–10 round pilot rather than the binary question of whether d_1 is heterogeneous: the same quantity Theorem 1(1) already requires, which the labelling had not been using.

U Condition-by-condition scoring, and the run-count saving

The results below are §3’s, moved here for space. None changes a reported number.

The out-of-sample bookkeeping, in full. All 4 out-of-sample PASS predictions the tightened criterion licenses held, as did both in-distribution $C1 \wedge C2 \wedge C3$ pairs. **But all 7 are *positive* predictions whose FAIL controls are uninformative** (DBA fails even undefended), which is why §3 declines to read them as an out-of-sample classification result: five positive predictions against an 11.9% base rate (5/42, Appendix S), with the positive direction separately unvalidated (§4).

Does every condition earn its place? C3 does not, reported as a finding. Scoring all 42 pairs on each condition *independently*: C1 alone selects 4 at 50% precision, C2 alone 14 at 36%, $C1 \wedge C2$ selects 2 at 100%, and **C3 alone selects 35 at 9%**, below the 12% base rate, never binding on a pair $C1 \wedge C2$ already admits. So **we demote C3 to a margin diagnostic and recommend $C1 \wedge C2$** , retaining it on theoretical grounds only: Theorem 1 makes margin compression a real failure mode and the $\alpha=0.1$ failure of §3 is exactly a margin collapse. It changes no number: a finding about the criterion, not an improvement to it.

What the saving is. Counted in runs rather than pairs (C1’s standalone evaluations are runs too) the saving is 420 **unscreened** against 80 **screened**, 81% **fewer composition-evaluation runs**, not the 95.2% of *pairs* we previously reported; **nothing is eliminated before any run**, and the structural claim is the scaling $|A| S D(D-1)$ (*quadratic* in the defense count) to $|A| S ((D-1) + |\text{admitted}|)$. GPU-hour figures are *estimates*; no timing telemetry was recorded.

V What the boundary permits, claim by claim

The inventory below separates the three kinds of *no* this paper reports: *structural* (no design satisfying the gate can answer the question, Prop. 4), *open* (answerable in principle, but not by an oracle instrument), and *scope* (outside what the committed-attack setting measures). Nothing here is a new claim; each row points at the section that establishes it.

Table 9: **What this design can and cannot conclude.** The middle column is about *identifiability*, not statistical power: a “no” marked *structural* cannot be repaired by more seeds, defenses or attacks (Proposition 4). Every entry is scoped to the measured setting of §5.

question	identifiable?	answer, and its evidence
Does disturbing d_2 ’s statistic change its <i>decision</i> ?	measured	Yes: Krum’s selection flips in up to 0.80 of rounds
Does it change the <i>admitted adversarial mass</i> ?	by construction	No: unchanged in every one of the 60 measured rounds
Does it change <i>suppression</i> ?	instrumented	No: $ \Delta =0.026$, inside margin (Prop. 2(a))
Does <i>adversary-correlated</i> reweighting change suppression?	instrumented	Yes, non-monotonically: single-peaked (Mode A)
Does admitted mass <i>order</i> outcomes across defenses?	pre-registered	No: refuted: +0.098 against a frozen +0.150
Is (P4) <i>sufficient</i> for (P5)?	control arm	No: <code>cos_krum</code> admits 0.000 while ASR falls
Can (P4) be screened <i>without</i> adversary identity?	no: open	Both instruments are oracles: the stated open problem
Does C1 identify C2’s contribution?	no: structural	Empty cell by the gate, for every menu (Prop. 4)
Does any of this certify <i>deployed</i> security?	no: scope	Committed attack only; one adaptive test breaks it

W Anticipated objections

Five objections we expect, answered against text that is already in the paper rather than by new argument. Nothing here is a new claim; each answer is a pointer plus the scope it turns on.

1451 **1. This is small-scale federated learning, not decentralized foundation-model training.**
 1452 Agreed, and we claim **methodological** relevance only (§1): no transformer arm, no foundation-
 1453 model measurement, and no sentence implying either. What transfers is a statement about *causal*
 1454 *structure*: any criterion whose eligibility gate is defined on the same outcome it predicts conditions
 1455 on a descendant of both mechanism and effectiveness (Prop. 4), and that argument does not depend
 1456 on model scale or participant count. What does *not* transfer is every empirical magnitude, which
 1457 is why each is reported per arm, tiered at **very unequal weight and never pooled** (Table 2), and
 1458 scoped in Table 1. The result forbids an inference; it does not measure a system.

1459 **2. The interventions are oracles: no deployed system can pin an adversary’s coefficient or**
 1460 **read adversary identity.** Correct, and that is why Mode S is an **instrument, not a screen**. The
 1461 Conclusion says so in those terms: Mode S measures level (P4) only by reading adversary identity,
 1462 which no deployable defense can, so our experiments **identify the causal channel and are not a**
 1463 **(P4) screen**. Pinning $c=1.0$ has one purpose, to cut the confounded path of Figure 1(a) by construc-
 1464 tion, exactly as an oracle-labelled control isolates a channel in an ablation. The deployable question
 1465 is left open as the paper’s closing sentence. The one control that *removes* an oracle component
 1466 (score-only, which scores the transformed stack but aggregates the untransformed selected update)
 1467 does not change the verdict (App. E, D).

1468 **3. $n=5$ seeds is small.** It is: $n=5$ per arm, $n=8$ for `cos_krum`, with Student- t 95% intervals
 1469 throughout and no claim that a small effect is zero. Where suppression does not move we report **no**
 1470 **evidence of a practically meaningful change under the pre-specified ± 0.15 equivalence margin**,
 1471 a margin frozen before the runs (Table 1), never a claim that suppression is unchanged. Two design
 1472 choices do more for precision than seeds would: the identity rungs are imported **bit-identical per**
 1473 **seed** (largest deviation 0.00e+00), so the contrast is within-seed rather than across noise, and the
 1474 central claim is a **sign reversal** (-0.272 against $+0.098$ on the identical cell, Table 3), which is not
 1475 a null result. Finally, Prop. 4 says the cell that would validate the criterion is *empty by design*: more
 1476 seeds cannot repair it (Table 9).

1477 **4. If ASR does not move, why does a change of decision matter?** Because the change of de-
 1478 cision is what a preservation criterion *measures*, and the gap between the two is the finding. An
 1479 invariance argument certifies the top three levels of Def. 1 while the security content sits at level
 1480 (P4), the adversarial contribution admitted. A criterion reading the Krum arm would call that com-
 1481 position broken (the decision flips in up to 80% of rounds) for a composition whose suppression
 1482 shows no measurable change: a false alarm in the direction that costs coverage (Prop. 5 caps such a
 1483 screen’s recall at 80%). The converse fails too: `cos_krum`’s selection is bit-identical at every rung
 1484 yet it never suppresses (0.584 at the identity) and its ASR *falls* 0.173 at the first accuracy-gate-
 1485 passing rung, the full-ladder endpoint of 0.425 lying below the accuracy floor, so exact invariance
 1486 of the statistic neither implies suppression nor holds it fixed (App. J.5).

1487 **5. The theory covers only positive per-client rescalings.** Yes, and we say so wherever it is used:
 1488 **a laboratory for studying composition, not a general model of it** (§1; Prop. 1, Thm. 1, App. I).
 1489 The class is not arbitrary (it is what norm clipping, contribution weighting and reputation weighting
 1490 do to a client’s update) but it excludes coordinate-wise filters and any transform that changes an
 1491 update’s direction, and the theorem’s own margin closes under heterogeneity ($m=3.24 \rightarrow 1.01$ at
 1492 $\alpha=0.1$), so **it certifies nothing where heterogeneity is high**. The negative does not rest on the
 1493 theory: it is a *measured* counterexample, and a narrower theory makes the counterexample sharper
 1494 rather than weaker, since **the theorem certifies mechanism preservation, not security**, it bounds
 1495 no ASR.

1496 X Experimental Details

1497 Figure 7 places the evaluated settings in the (γ, p_{eff}) plane. At the persistence we measure ($\hat{\gamma} \approx$
 1498 0.66) the one temporal mix we ran reaches ASR 0.956, which is the reason the paper composes
 1499 rather than switches; the figure records those points and draws no boundary between them.

1500 **FL configuration.** $N=10$ clients, $K=5$ per round, 50 rounds, SGD with momentum 0.9, learning
 1501 rate 0.01, local epochs 3, batch size 32, Dirichlet $\alpha=0.5$ (non-IID). Scale experiments: $N=100$,
 1502 $K=10$, $f=0.05$, ResNet18 with GroupNorm.

1503 **Attacks.** Model-scaling: scale factor 10, 50% poison fraction, target class 0. Pixel backdoor:
 1504 5×5 bottom-right trigger, target class 0. DBA: 4 sub-triggers at image corners, $f=0.4/0.6$. Joint-
 1505 constraint: $\epsilon \in \{0.5, 1.0, 2.0, 4.0\}$, disjoint coordinate blocks per adversary, single-step ϵ -ball pro-
 1506 jection per round.

1507 **Defenses.** NormClip $\tau=5.0$, i.e. $u_i \mapsto u_i \min(1, \tau/\|u_i\|_2)$ over the whole flattened update;
 1508 TrimmedMean $\beta=0.2$ (trim 20% each tail); RFA: unweighted Weiszfeld, 50 iterations, tolerance
 1509 10^{-6} , initialised at the mean; FoolsGold: $w_i \propto \log((1 - s_i)/(s_i + \epsilon) + \epsilon)$ with s_i the client's
 1510 maximum absolute pairwise cosine similarity normalised by the round maximum, $\epsilon=10^{-5}$, then
 1511 shifted by $-\min_i$ and normalised, so the least-trusted client of each round receives weight exactly
 1512 0 (empirically 2 of 5 clients per round).

1513 **Reputation.** Despite the name this defense is *memoryless*, it keeps no cumulative per-client state
 1514 and is recomputed from scratch each round from that round's K participants. Consensus is the
 1515 coordinate-wise median m of the K updates; scores are ℓ_2 distances $d_i = \|u_i - m\|_2$; the bandwidth
 1516 is adaptive, $\sigma = \text{median}_i(d_i)$ (clamped $\geq 10^{-6}$); weights are $w_i = \exp(-d_i/\sigma)$, normalised to sum
 1517 to 1. No clipping, no initialisation. Its discriminative property is therefore *consensus distance*, not
 1518 similarity spread, which is why Proposition 1(c) predicts that $\cdot \rightarrow \text{Reputation}$ compositions fail C2.
 1519 Weights are strictly positive, so Reputation never produces the zero vectors of Lemma 1.

1520 **Adversary count per round.** Adversaries are the fixed client identities $\{0, \dots, fN - 1\}$; each
 1521 round samples K participants uniformly without replacement from all N , with no stratifica-
 1522 tion. For $N=10$, $f=0.2$, $K=5$ the per-round adversary count is hypergeometric: $P(n_a=0) =$
 1523 $P(n_a=2) = 56/252 = 0.222$, $P(n_a=1) = 140/252 = 0.556$; measured over 100 rounds we ob-
 1524 serve 0.19/0.54/0.27. Two consequences, both matching the data. CoordMedian's count majority
 1525 $n_a < K/2 = 2.5$ holds for every attainable $n_a \in \{0, 1, 2\}$, i.e. in 100% of rounds. TrimmedMean's
 1526 budget is $\lfloor \beta K \rfloor = \lfloor 1.0 \rfloor = 1$, so it is exceeded whenever $n_a=2$: 22.2% of rounds, ≈ 11 of 50.
 1527 This is why FG $\rightarrow$ TrimmedMean fails (0.508) while FG $\rightarrow$ CoordMedian succeeds (0.107) under an




Figure 7: **The evaluated settings, plotted. No regimes, and one line, which is a measurement.** X-axis: persistence γ . Y-axis: best single-defense ASR (p_{eff}). Each marker is a setting we ran and carries its own measured value; the dotted line is the measured FL persistence $\hat{\gamma} \approx 0.66$. **An earlier version shaded four regime bands and drew threshold guides at $\gamma=0.15$, $p_{\text{eff}}=0.3$ and $p_{\text{eff}}=0.6$; none of those numbers was measured, so all six are gone.** Five points cannot locate a boundary, and nothing here is evidence of a phase transition in (γ, p_{eff}) space.

otherwise identical upstream transform. We retain the pre-registered C3-FAIL label for FG→TM; the count-budget violation is an additional mechanism, not a relabelling.

Theorem quantities. FoolsGold weight ratio across 5 seeds $\times$ 20 rounds ($N=10$, $f=0.2$, model-scaling): max $\rho = 2.1$, giving extremeness ratio $S/\rho \geq 4.8$ with $S=10$ (regime A). Reputation, same protocol (5 seeds $\times$ 20 rounds, 81 rounds containing both adversarial and benign clients): ρ mean 1.44×10^5 , max 2.84×10^6 , p95 6.1×10^5 ; mean adversarial weight 4.83×10^{-4} against mean benign weight 0.278, and $w_a < 10^{-3}$ in 85% of rounds. Reputation therefore fails the ordering condition by ~ 5 orders of magnitude and suppresses via regime B instead. RFA, measured on the transformed points $u'_i = w_i u_i$ that RFA actually receives (25 adversary-present rounds over 3 seeds): FG zeroes every adversary in 32% of rounds, where Lemma 1 applies and the measured contraction $1 - \Lambda_a$ is 0.29 on average (range 0.001–0.52). In the other 68% ($w_a > 0$) the separation condition of Theorem 1(2) is satisfied in *every* round, with margin $\Delta'_{\text{sep}}/2(R'_B + \delta') \geq 2.8$ (mean 3.2) and δ'/R'_B mean 0.19, max 0.72. The two cases are exhaustive, so the analysis covers all measured rounds. Proposition 1(a) verified in the shipped implementation: FoolsGold’s weight vector changes by at most 3×10^{-8} (float32 rounding) under NormClip across 1000 synthetic trials and live rounds in which 1–2 of 5 clients were genuinely clipped (raw norms up to 24.5 against $\tau=5$). Exactly 2 of the 5 clients per round receive zero FoolsGold weight. **Corollary 1 in full:** with $D_b := R'_B + \delta'$ bounding every benign residual and $D_a := \Delta'_{\text{sep}} - R'_B - \delta' = (2m - 1)D_b$ lower-bounding every adversarial one, the Weiszfeld coefficients satisfy $\sum_j 1/d_j \geq n_b/D_b$ and $1/d_a \leq 1/D_a$, so $\Lambda_a \leq (n_a/D_a)/(n_b/D_b + n_a/D_a) = n_a/(n_b(2m - 1) + n_a)$. At $m \rightarrow 1$ this is $n_a/(n_b + n_a)$, i.e. 40% for $n_a=2$, $n_b=3$; at the measured $m \geq 2.82$ it is $\leq 12.6\%$, and at the all-round mean $m = 3.24$, $\leq 10.8\%$.

Pre-registration provenance. We distinguish the tiers explicitly. The 18-pair development set is where the criterion was *developed*; its predictions were not committed in advance and we report it as fit. The 24-pair held-out set was pre-registered in `pre_registration_wave2.md` and committed before those experiments ran; it contains no predicted-PASS pair, so it bears only on specificity. The out-of-sample PASS predictions (DBA transfer, $N=100$ /ResNet18 scale transfer, composition dominance) were pre-registered separately in `pre_registration_oos_pass.md`, also committed in advance. The 10 three-way predictions were pre-registered in `pre_registration_three_way.md`. Two held-out LOW calls, hedged as speculative at pre-registration time, failed; we score them as misses rather than dropping them. The prospective suite of §5 was pre-registered in `pre_registration_prospective.md` at commit 262cf35, which precedes the first write to `results/prospective_suite/`; single-defense baselines for the three new defenses were measured *before* the freeze, because C1 and C0 take them as inputs (FLTrust 0.048, Krum 0.583, Multi-Krum 0.783 max-committed; FedAvg 0.667 scaling / 0.768 pixel, clearing the sanity gate). The predicted labels in that file are reproduced verbatim in the PAIRS list of `run_prospective_suite.py` and the confusion matrix is recomputed from per-seed ASR in `results/prospective_suite/summary.json`, not transcribed.

Freeze commits and run tallies for the interventional suites. Every design, grid and threshold below was git-committed before the corresponding results directory existed, and each hash is verifiable in the repository history. *Dose-response ladder* (§5): `pre_registration_dose_response.md` at commit e711a95, before `results/dose_response/`; 80 runs. *Metric-swap mirrors* (§5): 8 frozen labels at commit 356897b. *Prospective suite* (§5): commit 262cf35, as above. *Targeted intervention, Modes S and A* (§4, §5): design and thresholds at commit 5130cec before `results/targeted_dose/` existed; 97 runs plus 25 identity runs verified bit-identical per seed (largest deviation 0.00e+0); the adversarial coefficient share was verified with no ASR involved, before any run. *Rank-discordant replication* (§4): `pre_registration_dose_replication.md` at commit 06d9bce before `results/dose_replication/` existed; 15 new runs, identity rung imported bit-identical from the ladder and verified by `--harness-check`. *Second dataset and architecture* (§4): `pre_registration_dose_femnist.md` at commit 478b555 before `results/dose_femnist/` existed; 12 new runs, with $\kappa=0$ run in-suite because no EMNIST-byclass ladder exists to import from. *Score-only control* (§4, Appendix E): `pre_registration_score_only.md` at commit 35788d9 before `results/score_only/` existed; 15 new runs, identity rung imported

1582 bit-identical and verified by `--harness-check`, and the `cos_krum` corollary checked at one seed
1583 as pre-registered.